

# VCCP: A Transparent, Coordinated Checkpointing System for Virtualization-based Cluster Computing

***N. Saragol*** \*      ***Hong Ong*** #  
***Box Leangsuksun*** +

***K. Chanchio*** \*

**\* Thammasat University, Patumtani, Thailand**

**# Oak Ridge National Laboratory, Oak Ridge, TN, USA**

**+ Louisiana Tech University, Ruston, LA, USA**

# Outline

---

- Introduction
- Goals and Motivations
- Related Works
- VCCP Mechanisms
  - Protocol and Analysis
- Experimental Results
- Progress and Future Works

# Introduction

---

- Fault tolerance is necessary for HPC
  - The more hw/sw components, the higher the chances some of them will fail
  - Bad for long-running parallel applications
- Checkpoint/Restart is a common technique used to provide fault tolerance. Common approaches include:
  - Modifying App. source code
  - Linking App with User-level library
  - Modifying OS kernel

# Problems

---

Existing checkpoint/restart tools for parallel processing are:


- **Complex:** require runtime systems and/or compilation tools, thus adding more components system-wide
- **Not user-friendly:** require additional works such as software installation, code modification, kernel modification, recompilation/re-linking.
- **Vendor-specific:** each software tool for parallel processing has a different checkpoint/restart implementation

# Goals

---

- High transparency
  - Checkpoint/restart mechanisms should be **transparent** to applications, OS, and runtime environments; no modification required
- Efficiency
  - Checkpoint/restart mechanisms should **not** generate unacceptable overheads
    - Normal Execution
    - Communication
    - Checkpointing Delay

# Goals (cont.)



network of VMs as **an abstract system unit** that all  
checkpointing to be conducted **transparently** and  
**efficiently** without modifying guest OS or applications.

# Motivations

---

- Leverage **Virtualization Technology** to provide highly transparent checkpointing mechanisms for HPC
  - Hypervisor can *transparently* **save** and **restore** VM state
- **A Cluster of Virtual Machines** can be built to run parallel applications
- Virtual Machine Technology **keeps getting better**
  - E.g. HW supports (Intel VT), Virtual I/O

# Related Works

---

- Most coordinated checkpointing approach are implemented at library level
  - E.g. LAM/MPI, MPICHV, CoCheck
- Some is hypervisor-based, but require software component on guest OS
  - Scarpuzza et al.: a Xen-based system on Infiniband
- Highly-transparent, hypervisor-based, but not coordinated checkpointing approach
  - Kangarlou et al. implement a non-blocking snapshot on Xen and VIOLIN virtual network



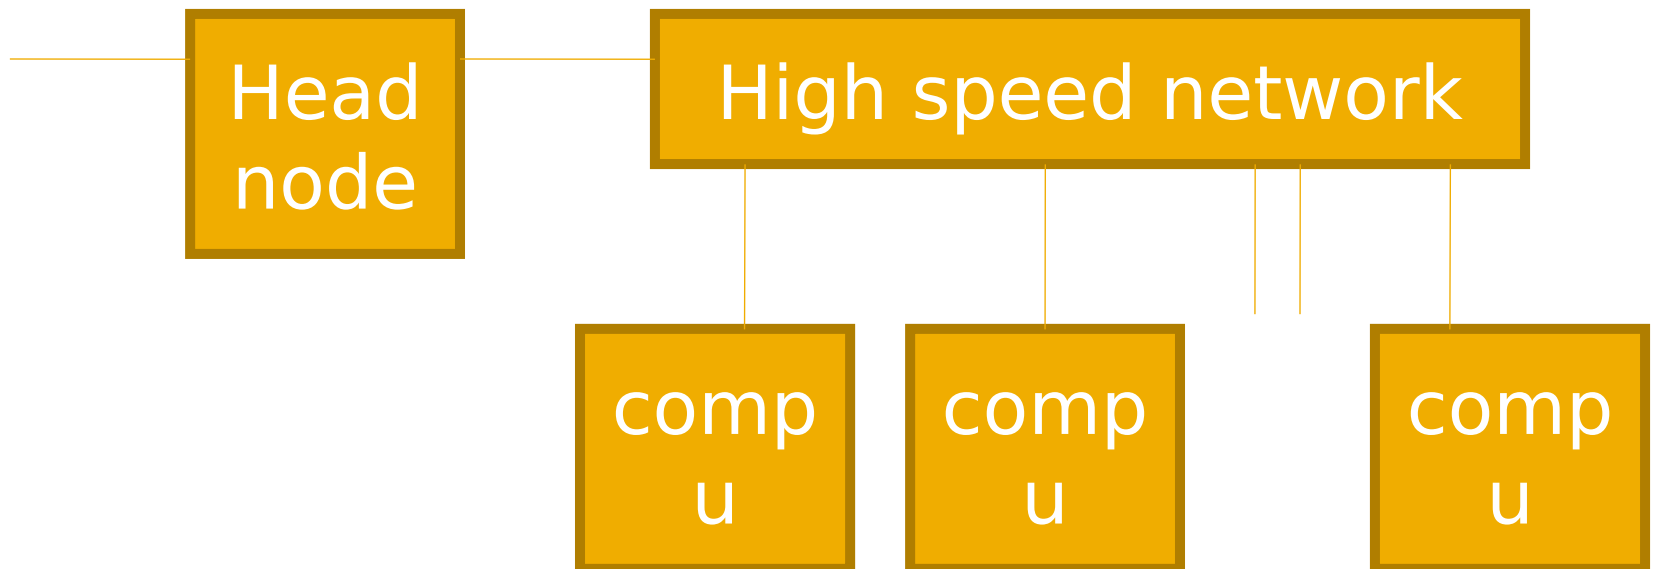
# Contributions

---


- Layered Virtual Cluster Architecture
- Novel Hypervisor-based Coordinated Checkpointing Protocols
  - Transparent C/R mechanisms
- Correctness Analysis
- Preliminary Implementation & Experiments
  - A proof of concepts

# Cluster Architecture

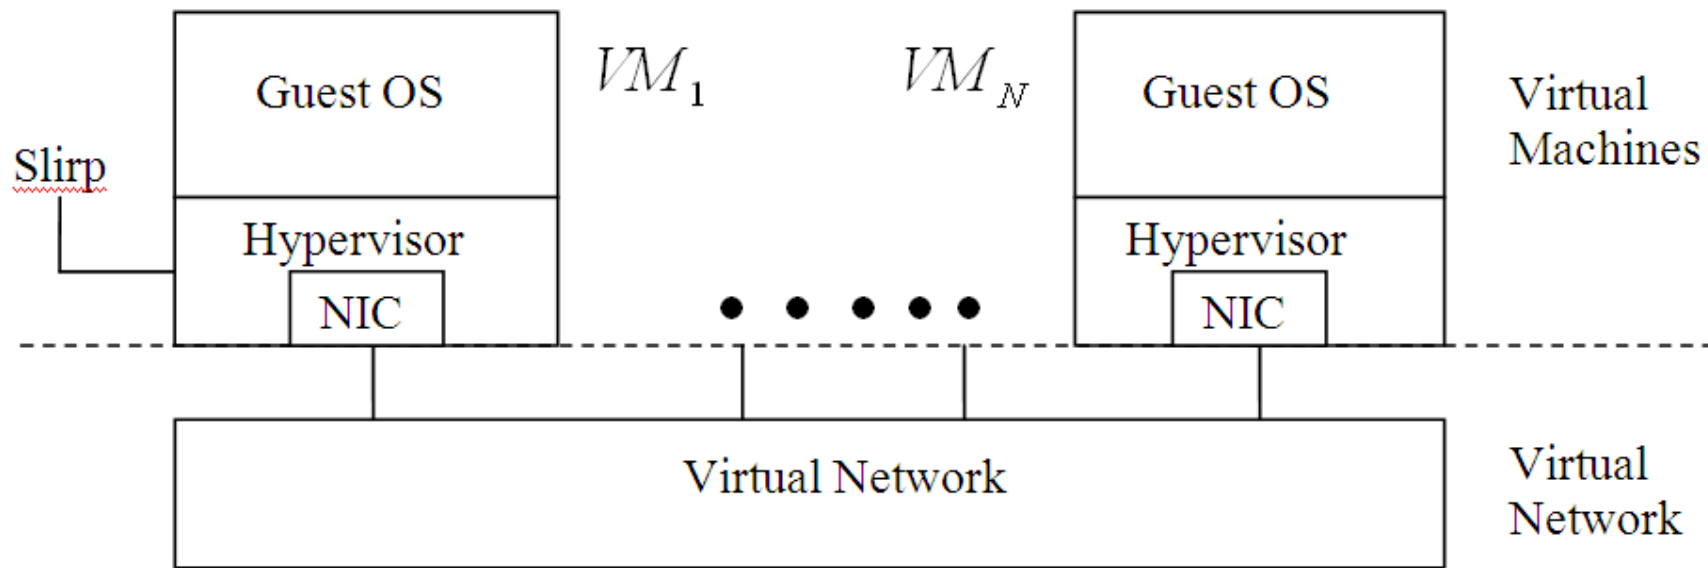
---



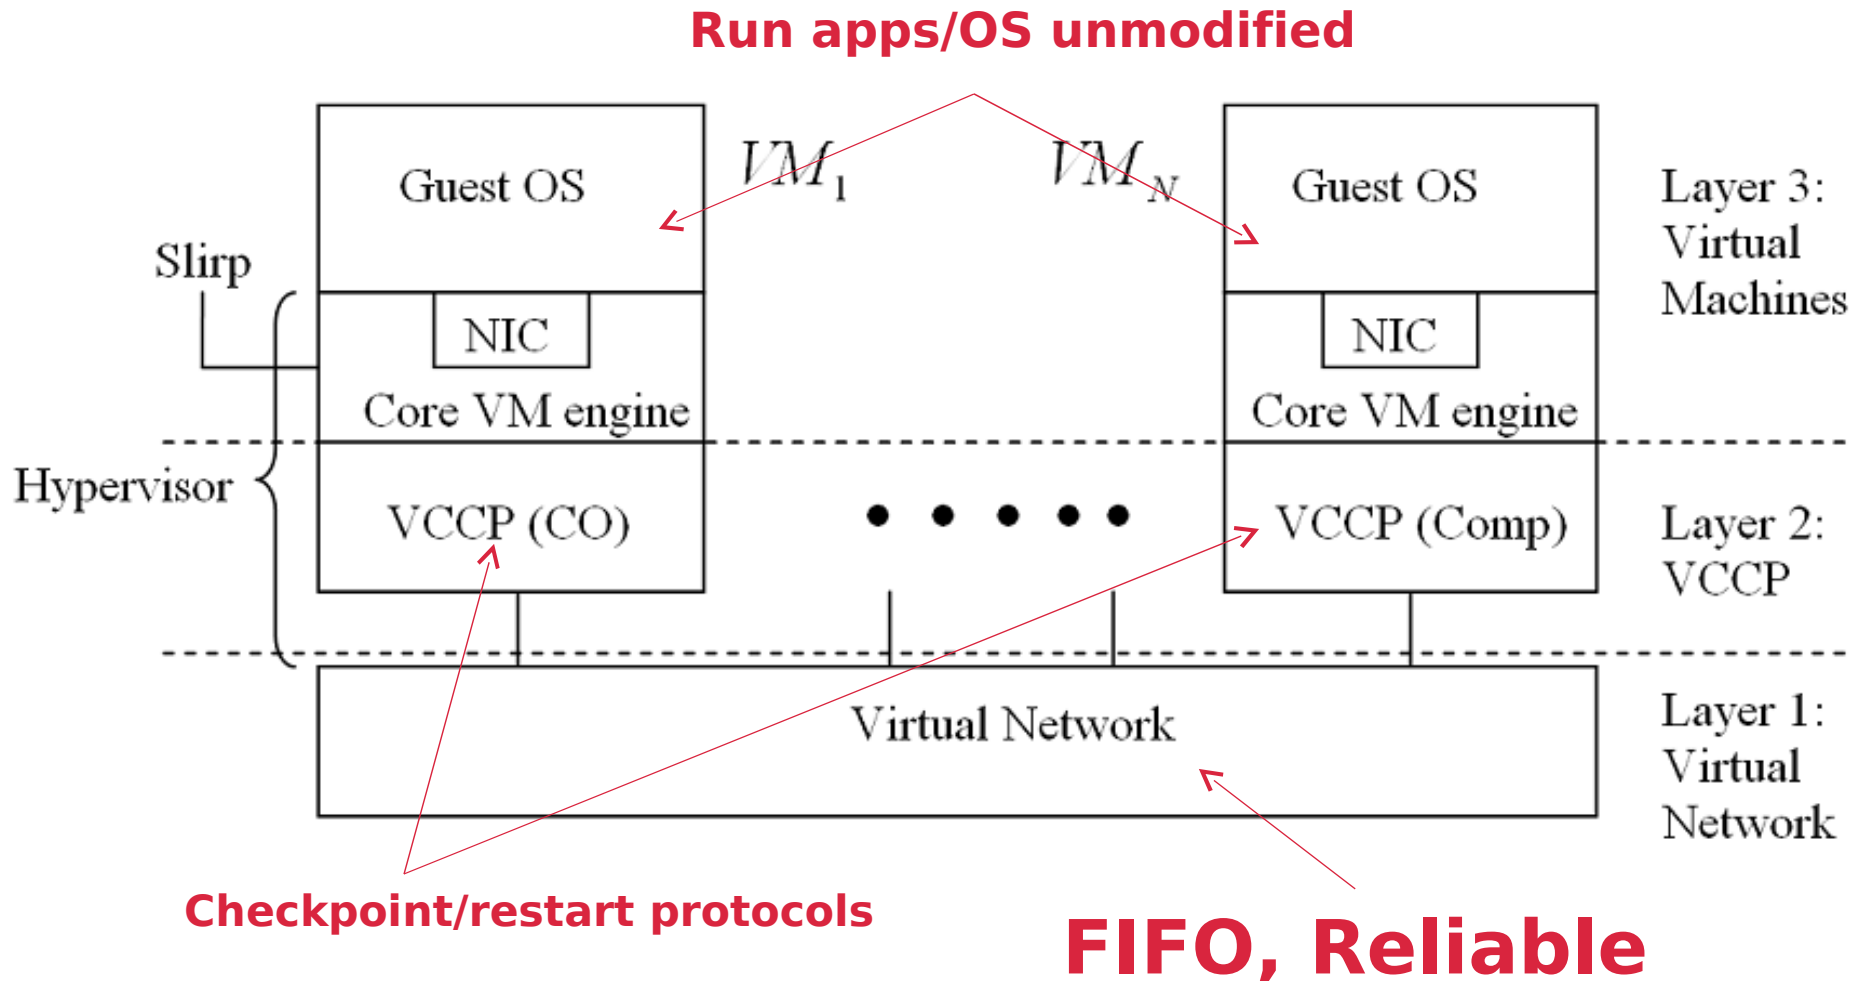
# Virtual Cluster Architecture



# Typical Virtual Cluster Architecture



# Virtual Cluster Architecture





# Virtual Cluster CheckPointing (VCCP) Protocol

1. Stop VM computation
2. Flush messages out of the network
3. Locally Save State of every VM
4. Continue computation

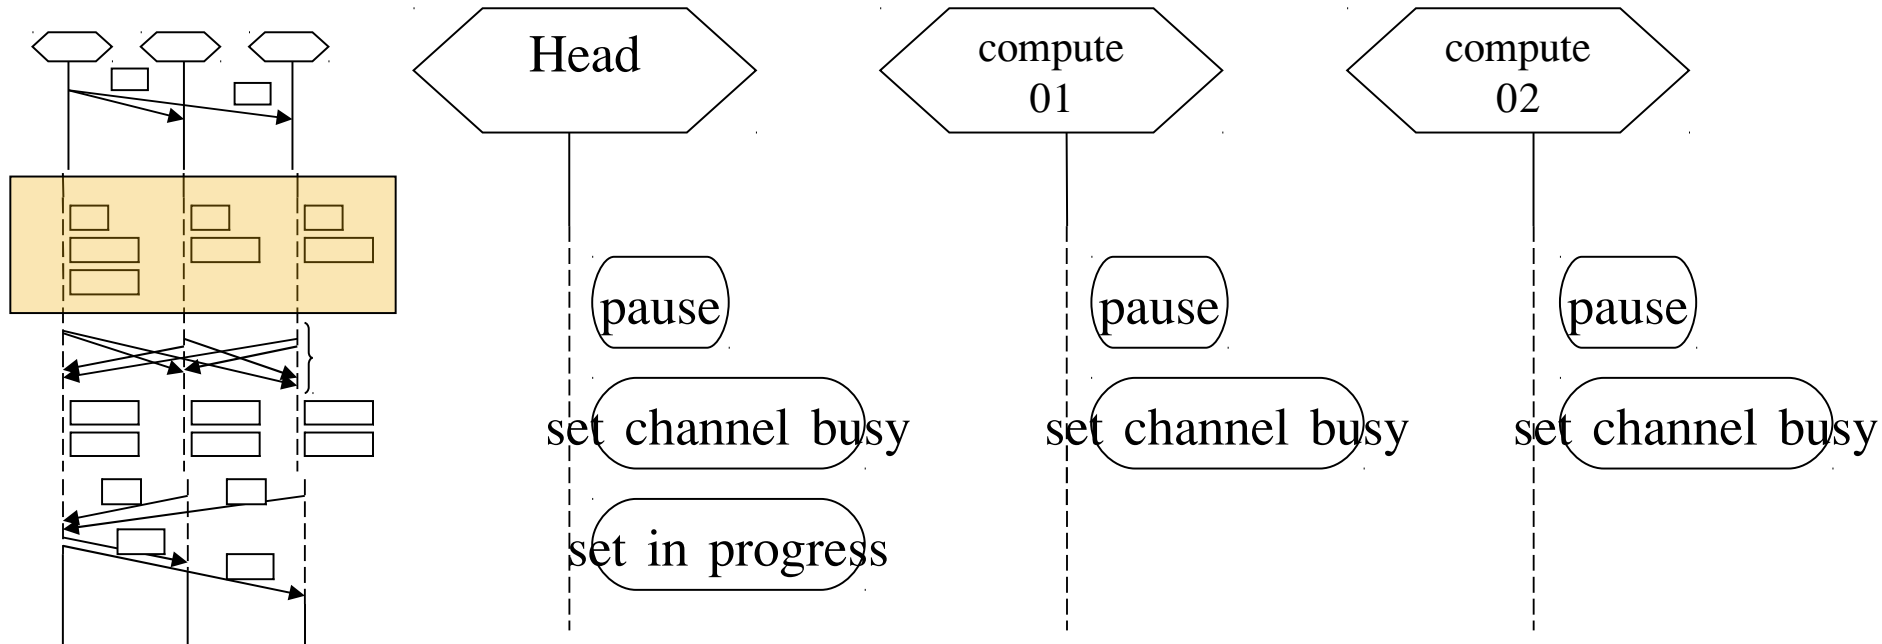
# VCCP

## VCCP checkpoint protocol



# VCCP

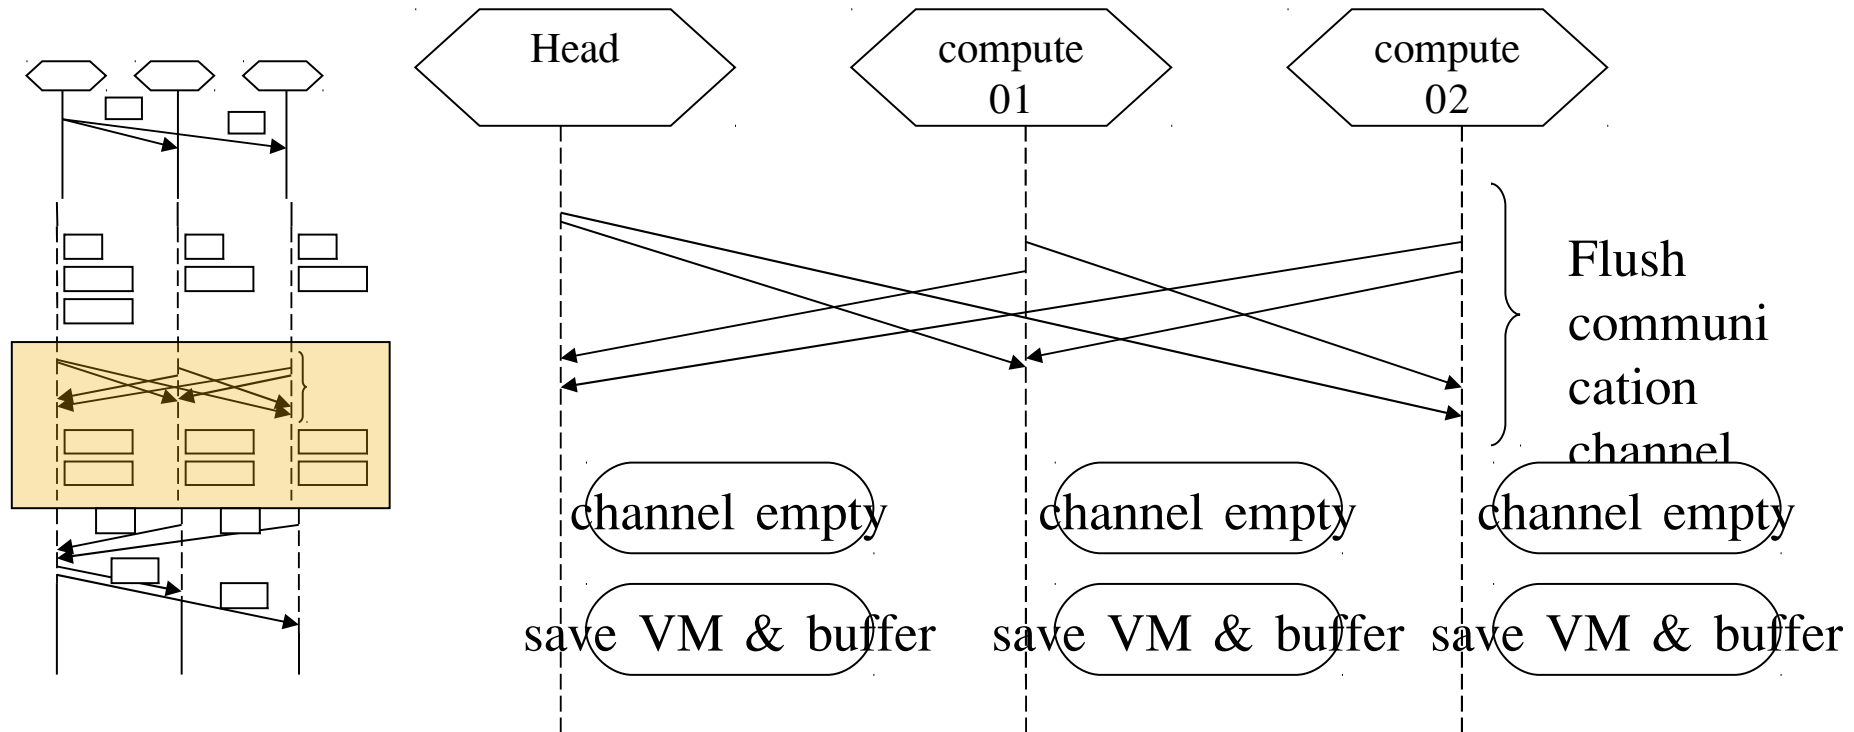
## VCCP checkpoint protocol





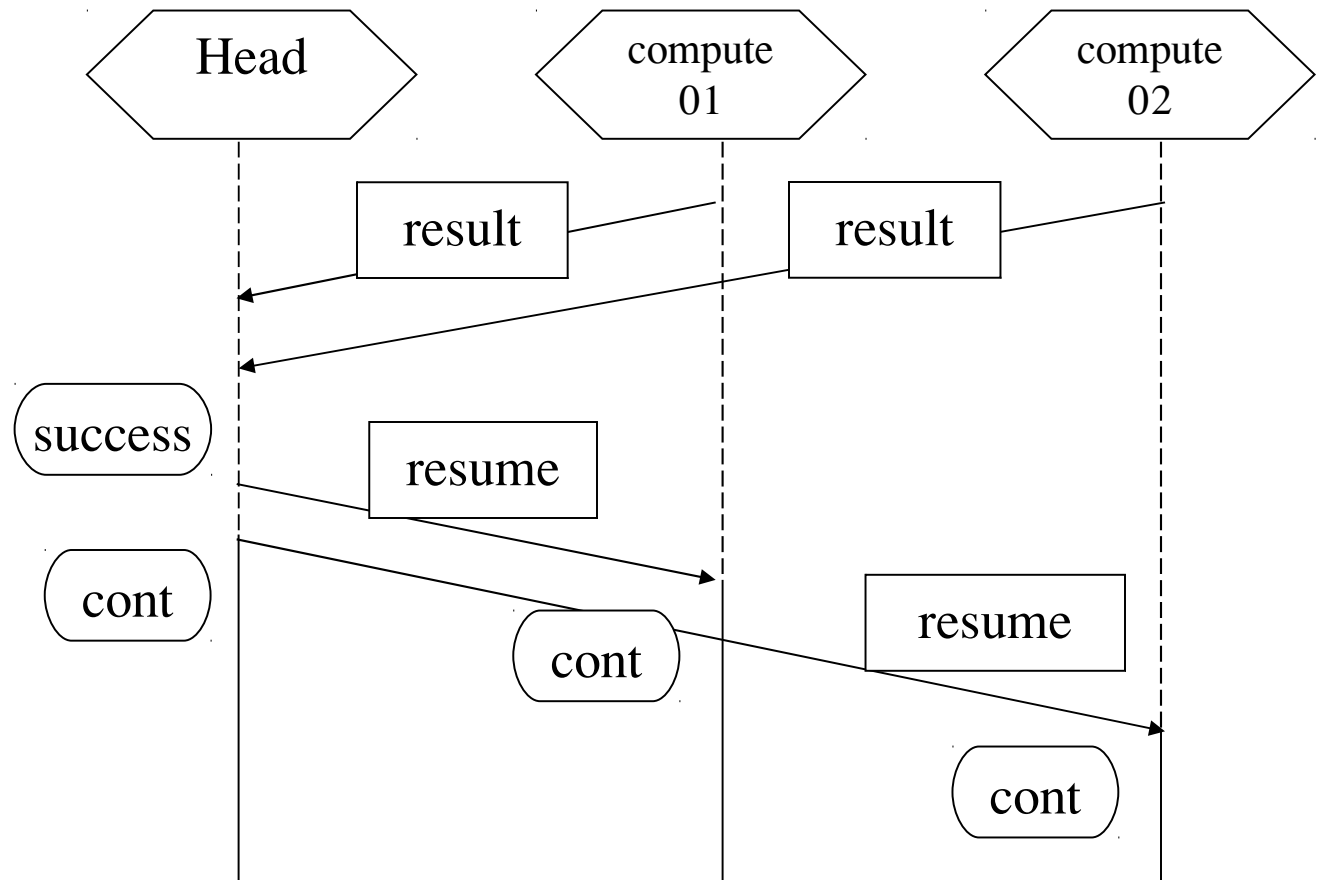

# VCCP

## VCCP checkpoint protocol



# VCCP

## VCCP checkpoint protocol



# VCCP recovery protocol

---

Coordinator:

1. Broadcast *loading* requests
2. Load VM state and Recv Queue
3. Repeat
4.     Receive frames
5. Until (*restore done* frames are received  
          from all other VMs)
6. Broadcast *resume VM* request
7. Resume local VM

Compute:

Upon receiving a *loading* request,

1. Load VM and Recv Queue
2. Send *restore done* frame to Coordinator
3. Wait until receiving a *resume VM* frame
4. Resume VM

# Correctness Analysis

- **The checkpoint/recovery protocols do not cause problems to normal execution**
  - TAB/VDE/TCP maintain FIFO ordering and reliable
- The protocols can cause clock skew
  - Cannot guarantee pausing and resuming at the exact same time on every VM
- Checkpointing and recovery events can affect RTT (See Next slide)

# RTT analysis


1. Checkpointing can affect message RTT; depending on the clock skew value
2. This problem also occurs on traditional checkpointing mechanisms
3. Solutions:
  - Make the timeout threshold large
  - Keep the clock skew value small

# Experiment

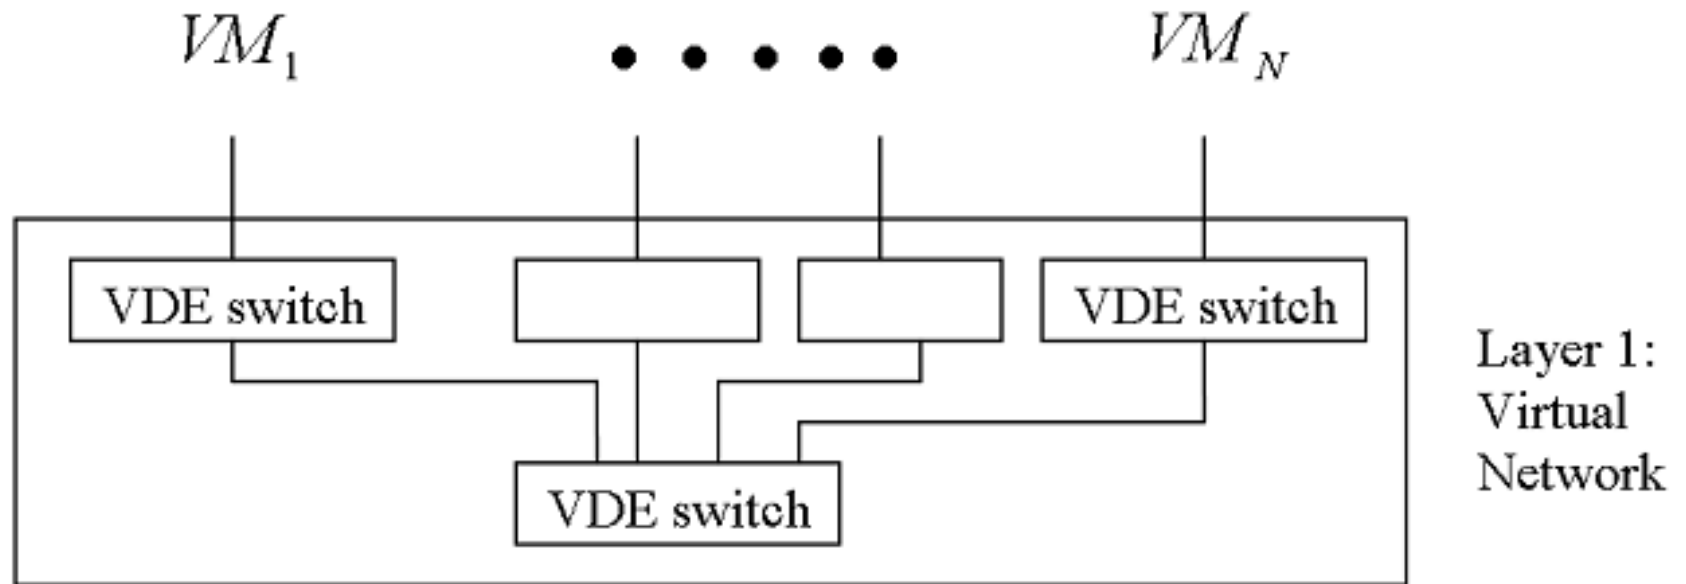
---

- Cluster Node x 9 (8 compute, 1 virtual switch)
  - Processor: Intel Xeon 2.6 GHz x 2
  - Memory: 2 GB
  - HD: 40 GB
  - Host OS: Rock 4.2.1
  - 1 Gigabit Ethernet
- Virtual Cluster x 8
  - Hypervisor: Modified QEMU
  - Guest OS: Damn Small Linux 3.4.1 + OpenMPI 1.2.3
  - Memory: 512 MB
  - HD : 300 MB (4 GB Maximum)

# Testbed Configuration



# Virtual Distributed Ethernet (VDE)





# Benchmarks

---

## **NAS Parallel Benchmark**

- EP : estimate floating point performance  
minimal interprocessor communication
- CG : estimate unstructured matrix vector  
multiplication performance with  
interprocessor communication
- IS : estimate integer sort performance  
with interprocessor communication
- MG: estimate data communication  
performance

# Overheads

<b><i>Kernel EP</i></b>	<i>2 nodes</i>	<i>4 nodes</i>	<i>8 nodes</i>
<i>VCCP cluster</i>	57.78	28.71	14.63
<i>QEMU cluster</i>	57.41	28.7	14.59
Overheads	0.37	0.01	0.04
Overheads %	0.64%	0.03%	0.27%

<b><i>Kernel IS</i></b>	<i>2 nodes</i>	<i>4 nodes</i>	<i>8 nodes</i>
<i>VCCP cluster</i>	2.48	1.73	1.46
<i>QEMU cluster</i>	2.45	1.69	1.42
Overheads	0.03	0.04	0.04
Overheads %	1.22%	2.37%	2.82%

<b><i>Kernel CG</i></b>	<i>2 nodes</i>	<i>4 nodes</i>	<i>8 nodes</i>
<i>VCCP cluster</i>	91.02	128.68	110.45
<i>QEMU cluster</i>	90.37	124.57	109.74
Overheads	0.65	4.11	0.71
Overheads %	0.72%	3.30%	0.65%

<b><i>Kernel MG</i></b>	<i>2 nodes</i>	<i>4 nodes</i>	<i>8 nodes</i>
<i>VCCP cluster</i>	7.88	17.05	21.29
<i>QEMU cluster</i>	7.7	16.87	21.17
Overheads	0.18	0.18	0.12
Overheads %	2.34%	1.07%	0.57%

# Overheads

<b><i>Kernel EP</i></b>	<i>2 nodes</i>	<i>4 nodes</i>	<i>8 nodes</i>
<i>VCCP cluster</i>	57.78	28.71	14.63
<i>QEMU cluster</i>	57.41	28.7	14.59

<b><i>Kernel CG</i></b>	<i>2 nodes</i>	<i>4 nodes</i>	<i>8 nodes</i>
<i>VCCP cluster</i>	91.02	128.68	110.45
<i>QEMU cluster</i>	90.37	124.57	109.74

Overheads			4.11	0.71
Overheads %	0.03		3.30%	0.65%

<b><i>Kernel IS</i></b>	<i>2 nodes</i>	<i>4 nodes</i>	<i>8 nodes</i>
<i>VCCP cluster</i>		17.05	21.29
<i>QEMU cluster</i>		16.87	21.17
Overheads		0.18	0.12
Overheads %	1.07	1.07%	0.57%

VCCP v.s. Qemu cluster  
Average 1.33 %  
Minimum 0.03% (EP 4 Nodes),  
Maximum 3.30% (CG 4 Nodes)

# Overheads (cont.)

<b><i>Kernel EP</i></b>	<i>2 nodes</i>	<i>4 nodes</i>	<i>8 nodes</i>
<i>VCCP cluster</i>	57.78	28.71	14.63
<i>REAL cluster</i>	51.39	25.73	12.92
Overheads	6.39	2.98	1.71
Overheads %	12.43%	11.58%	13.24%

- VCCP v.s. Qemu cluster
  - Average 1.33 %
  - Minimum 0.03% (EP 4 Nodes),  
Maximum  
3.30% (CG 4 Nodes)

- VCCP v.s. Real cluster
  - Average 12.41 % (EP)
  - Minimum 11.58% ,(EP)  
  
Maximum 13.24% (EP)

# Overheads (cont.)

<i><b>Kernel EP</b></i>	<i>2 nodes</i>	<i>4 nodes</i>	<i>8 nodes</i>
<i>VCCP cluster</i>	57.78	28.71	14.63
<i>REAL cluster</i>	51.39	25.73	12.92

## VCCP v.s. Real cluster

Average 12.41 %

Minimum 11.58% ,

Maximum 13.24%

(No hardware  
supports)

### ■ VCCP v.s. Q

■ Average 1.3

■ Minimum 0.

3.30% (CG

al cluster

1 % (EP)

58% ,(EP)

24% (EP)

# Overheads

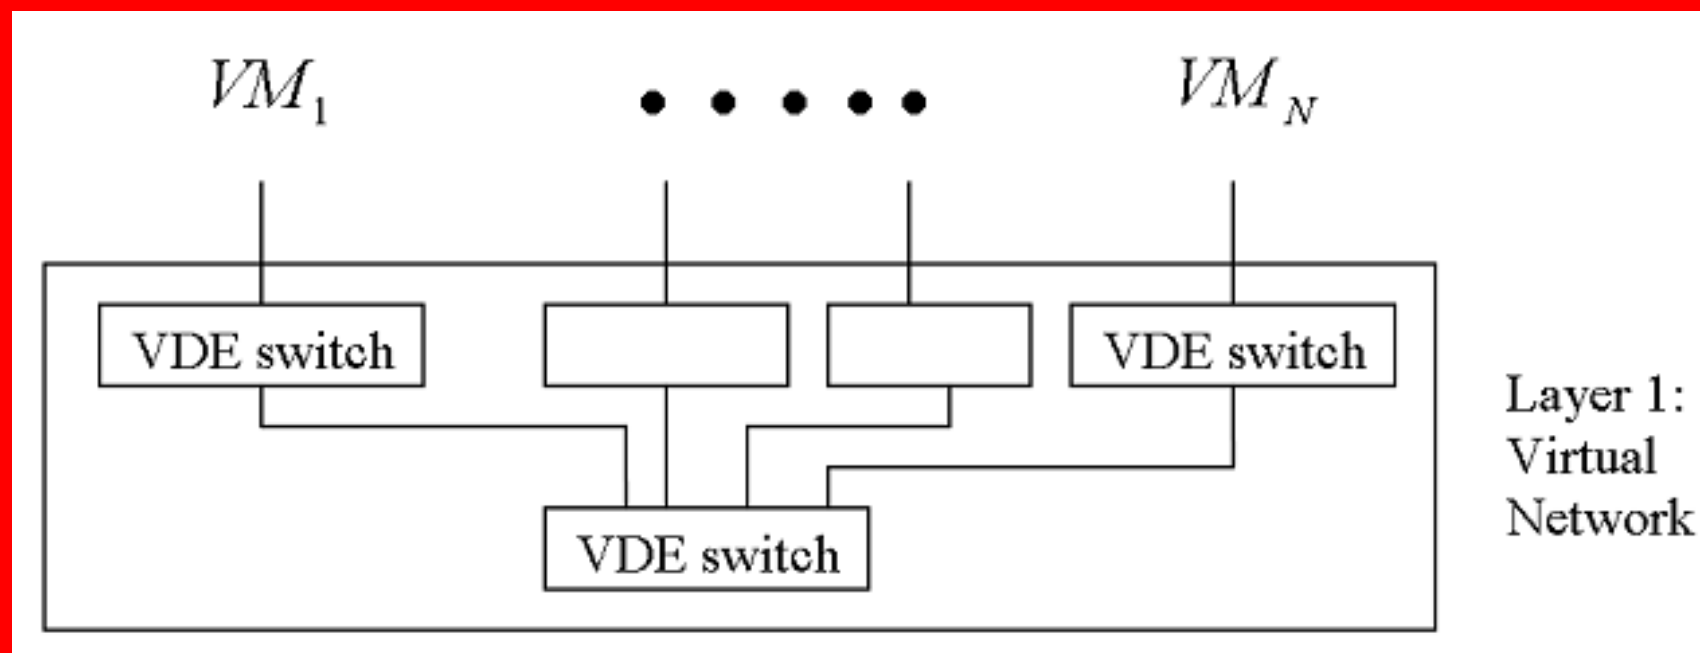
## Virtual Network Problems:

**Kernel CG**

2 nodes

4 nodes

8 nodes



Overheads	0.65	0.64	0.64	Overheads	0.70	0.70	0.72
Overheads %	1.22%	2.37%	2.82%	Overheads %	2.34%	1.07%	0.57%

# Overheads

## Virtual Network Problems:

1. VM and VDE compete for resources
2. Hot spot on central VDE switch

<b><i>Kernel IS</i></b>	<i>2 nodes</i>	<i>4 nodes</i>	<i>8 nodes</i>
<i>VCCP cluster</i>	2.48	1.73	1.46
<i>QEMU cluster</i>	2.45	1.69	1.42
Overheads	0.03	0.04	0.04
Overheads %	1.22%	2.37%	2.82%

<b><i>Kernel CG</i></b>	<i>2 nodes</i>	<i>4 nodes</i>	<i>8 nodes</i>
<i>VCCP cluster</i>	91.02	128.68	110.45
<i>QEMU cluster</i>	90.37	124.57	109.74
Overheads	0.65	4.11	0.71
Overheads %	0.72%	3.30%	0.65%

<b><i>Kernel MG</i></b>	<i>2 nodes</i>	<i>4 nodes</i>	<i>8 nodes</i>
<i>VCCP cluster</i>	7.88	17.05	21.29
<i>QEMU cluster</i>	7.7	16.87	21.17
Overheads	0.18	0.18	0.12
Overheads %	2.34%	1.07%	0.57%

# Checkpointing Performance

---


Checkpoint time = Flush time

- + Time to save VM State
- + Save messages to disks
- + Wait time

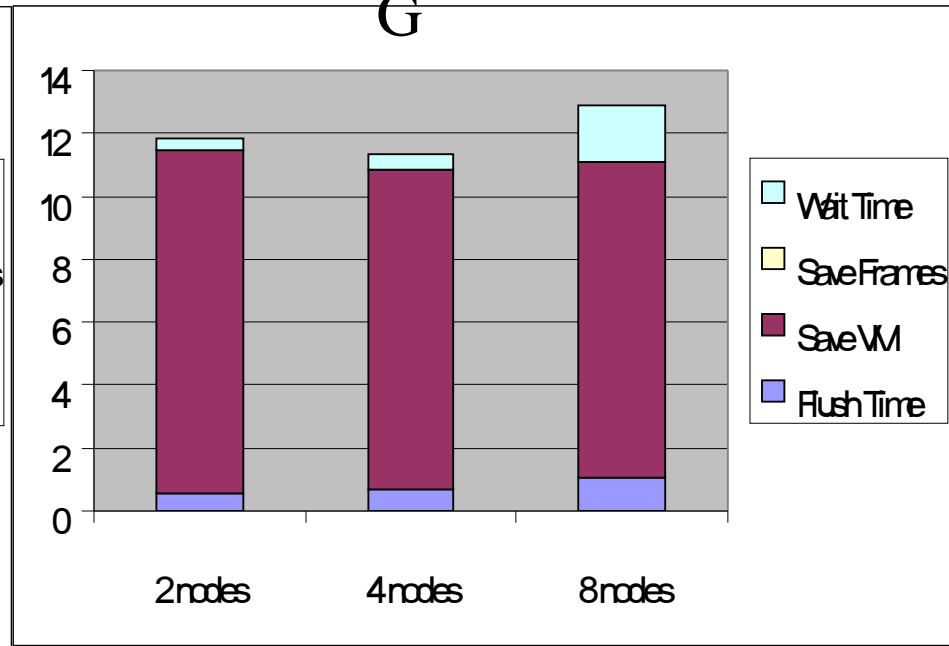


# Checkpointing Performance

E  
P




C  
G

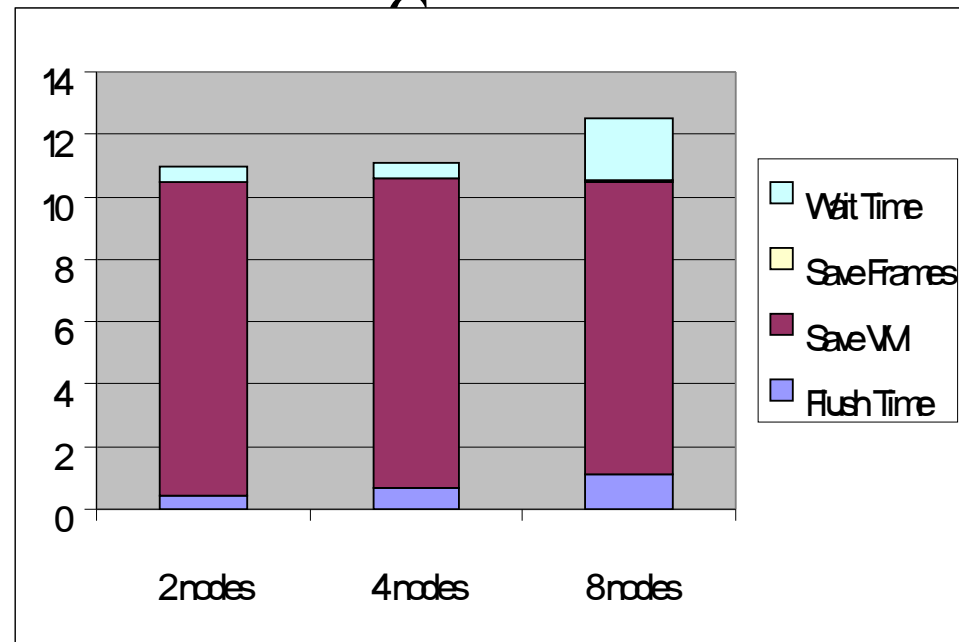


# Checkpointing Performance

IS



M



# Progress

---

- We have developed a transparent hypervisor-based coordinated checkpointing system for Virtual Clusters
- **Transparency** is achieved
- **Efficiency** requires further investigations
  - **Low** normal execution overhead
  - **High** communication overhead
  - **High** local checkpointing overhead

# Future Works

---

- Study and develop new virtual network
  - Direct TCP connections among VM
  - Virtio
- We are investigating low-latency local VM checkpointing mechanisms
  - Multi-threaded mechanism to save VM state
  - Diskless checkpointing

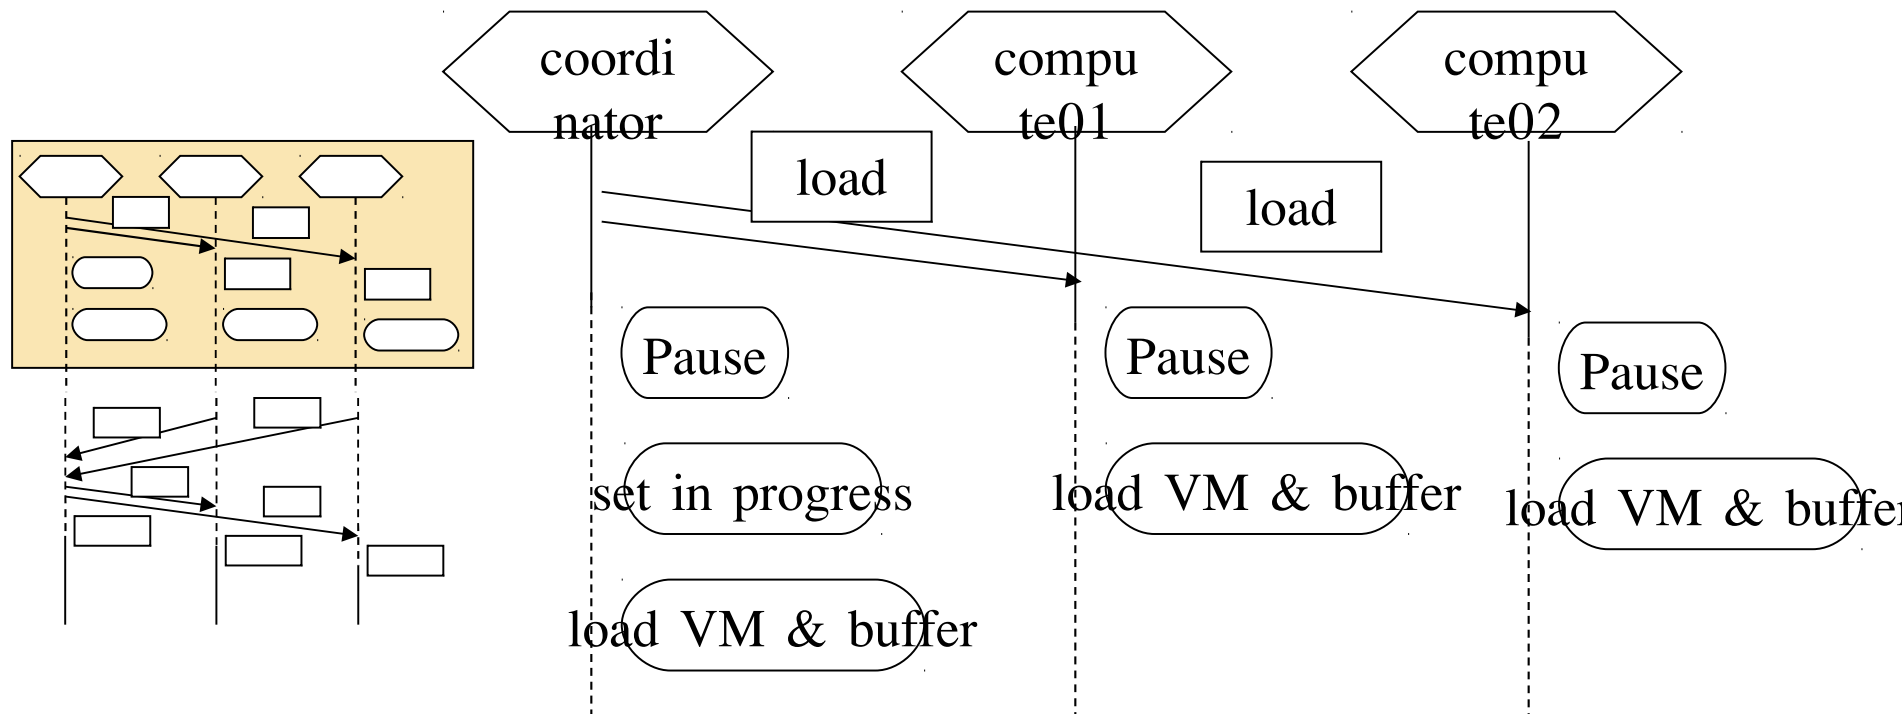
---

THANK YOU and QUESTIONS?

# Backup

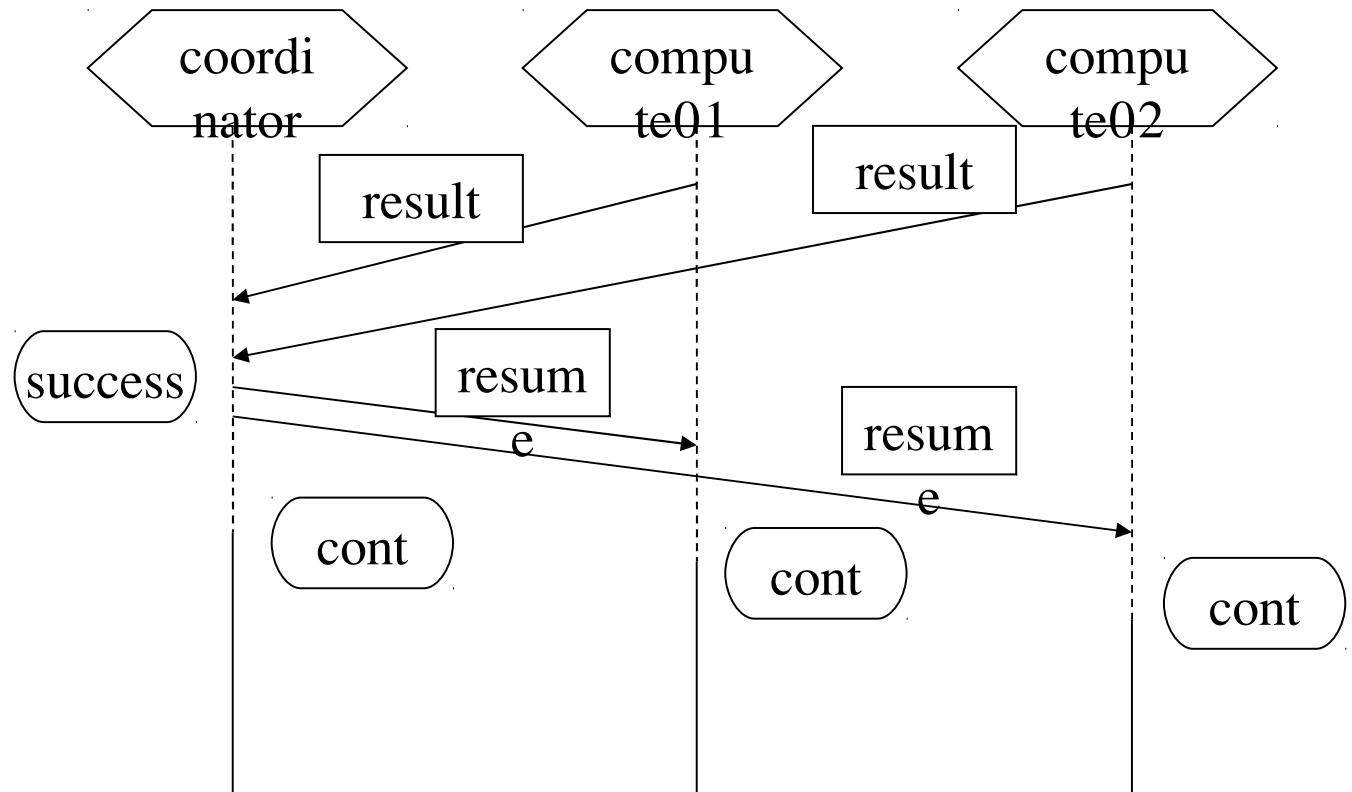

# VCCP

## VCCP recovery protocol



# VCCP

## VCCP recovery protocol





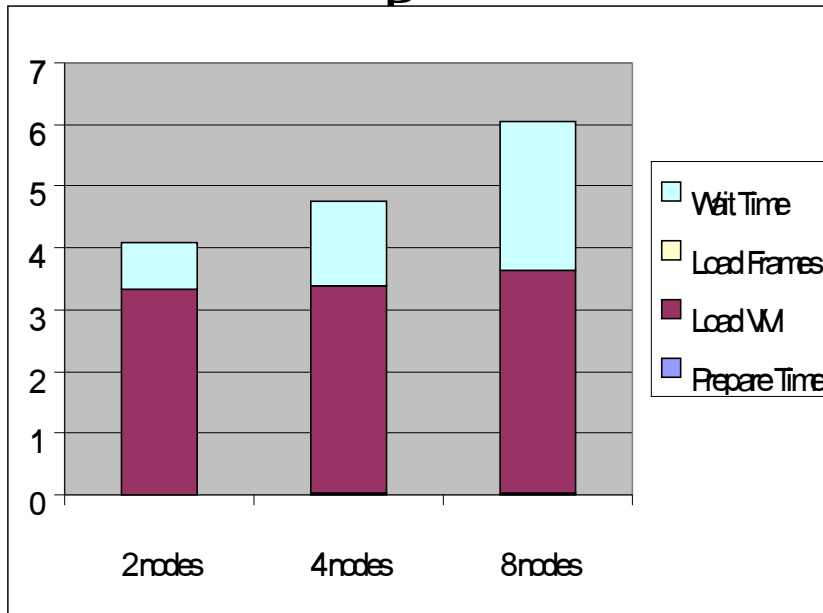
# Recovery Performance

---


Restoration time = coordinate time + load coordinator's  
vm time + load coordinator's frame time + wait  
time

# Recovery Performance

E  
D




C  
G



# Recovery Performance

IS



M  
G

