

OpenStack

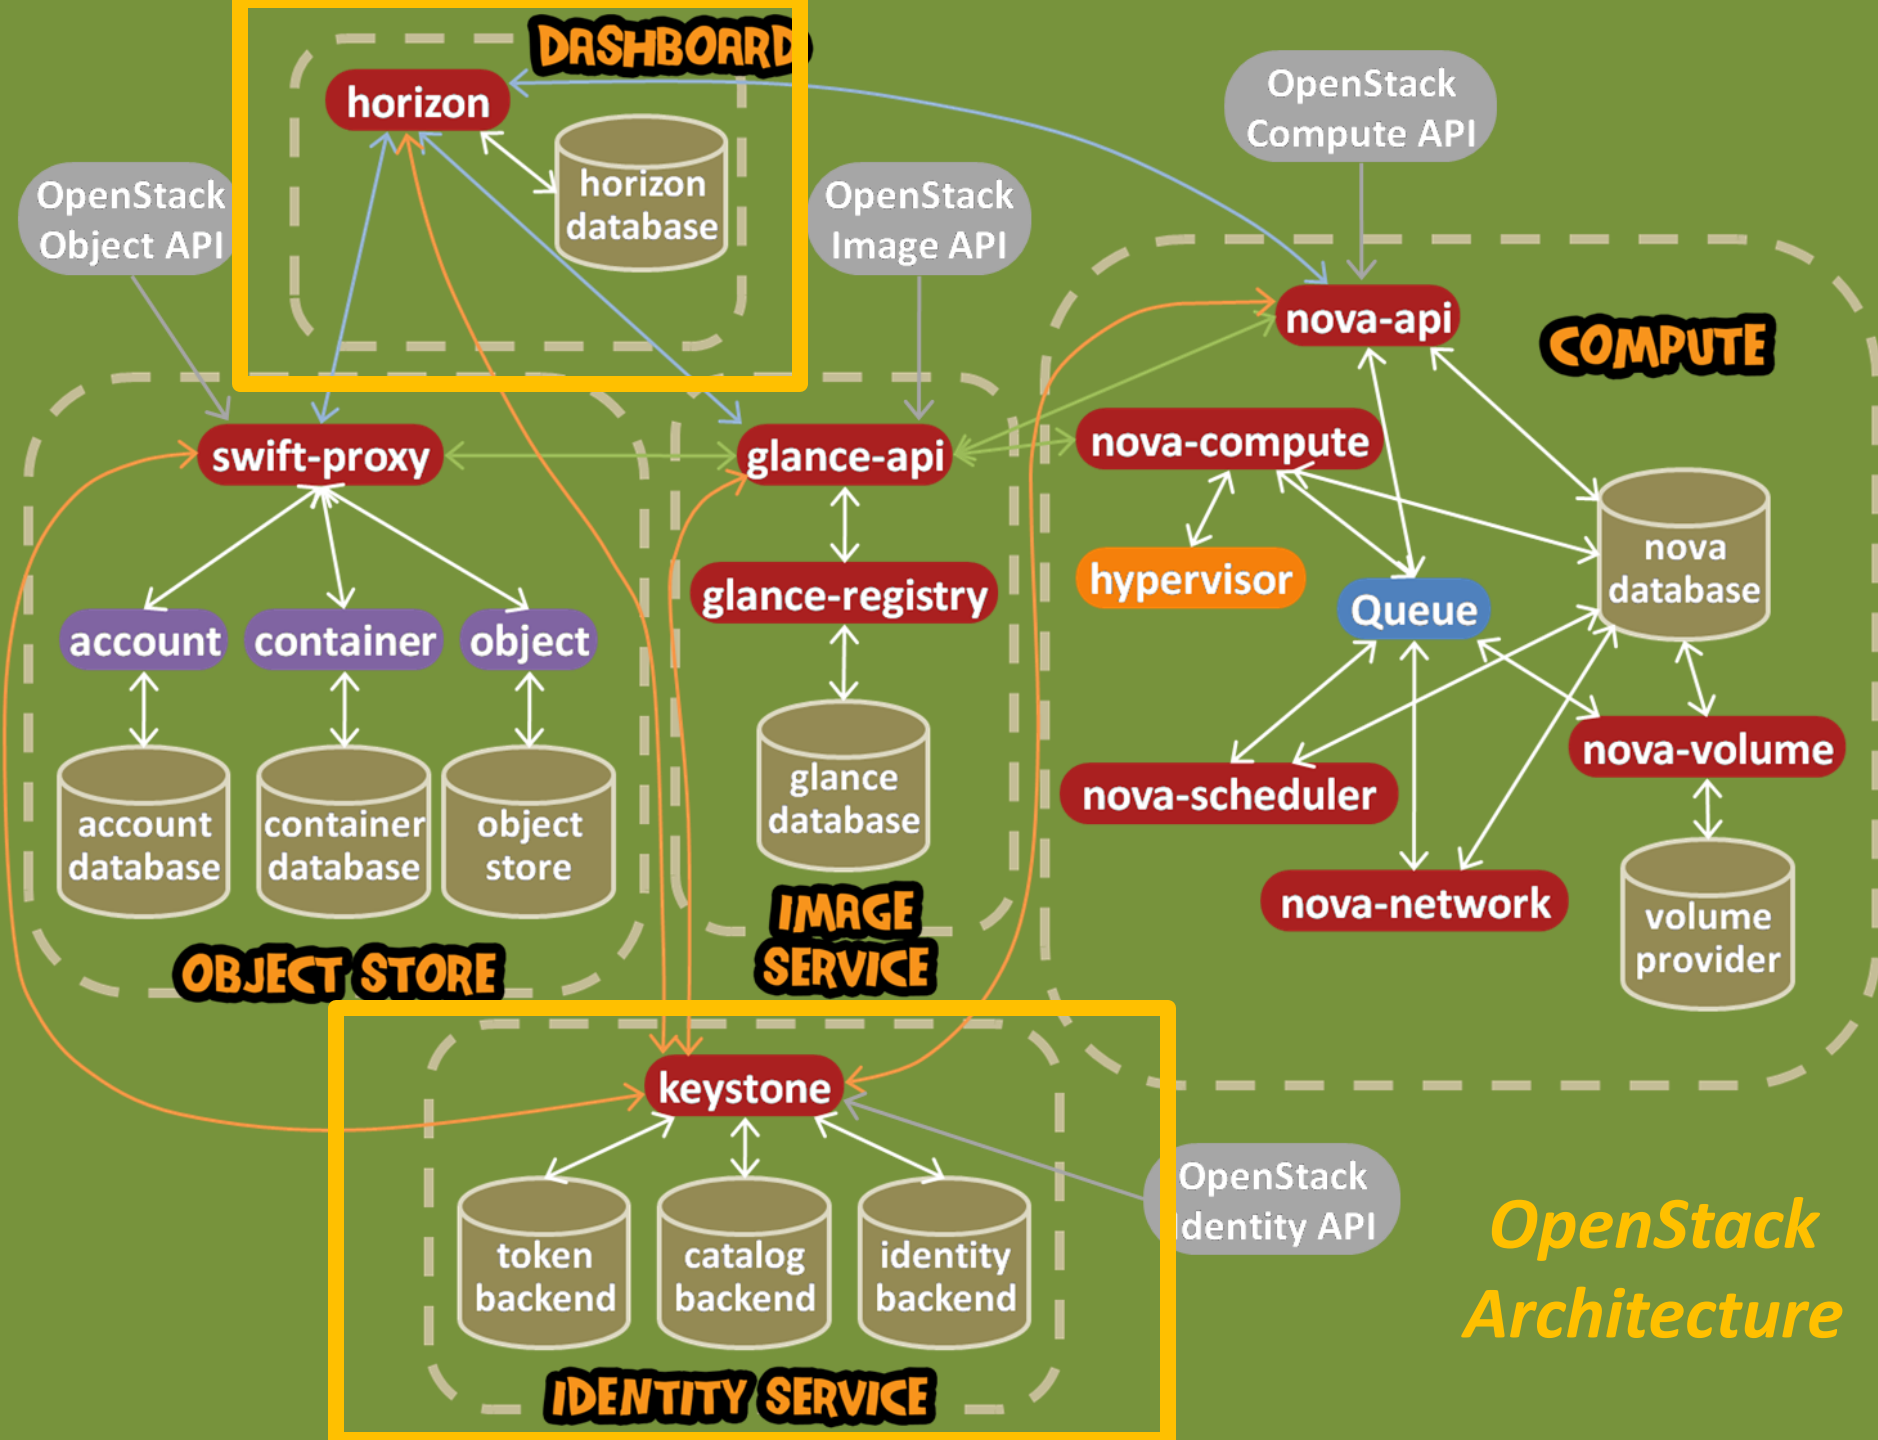
Architecture and Operation

OpenStack Tutorial Day 2


Kasidit Chanchio
Vasabilab,
Thammasat University

How OpenStack Components work

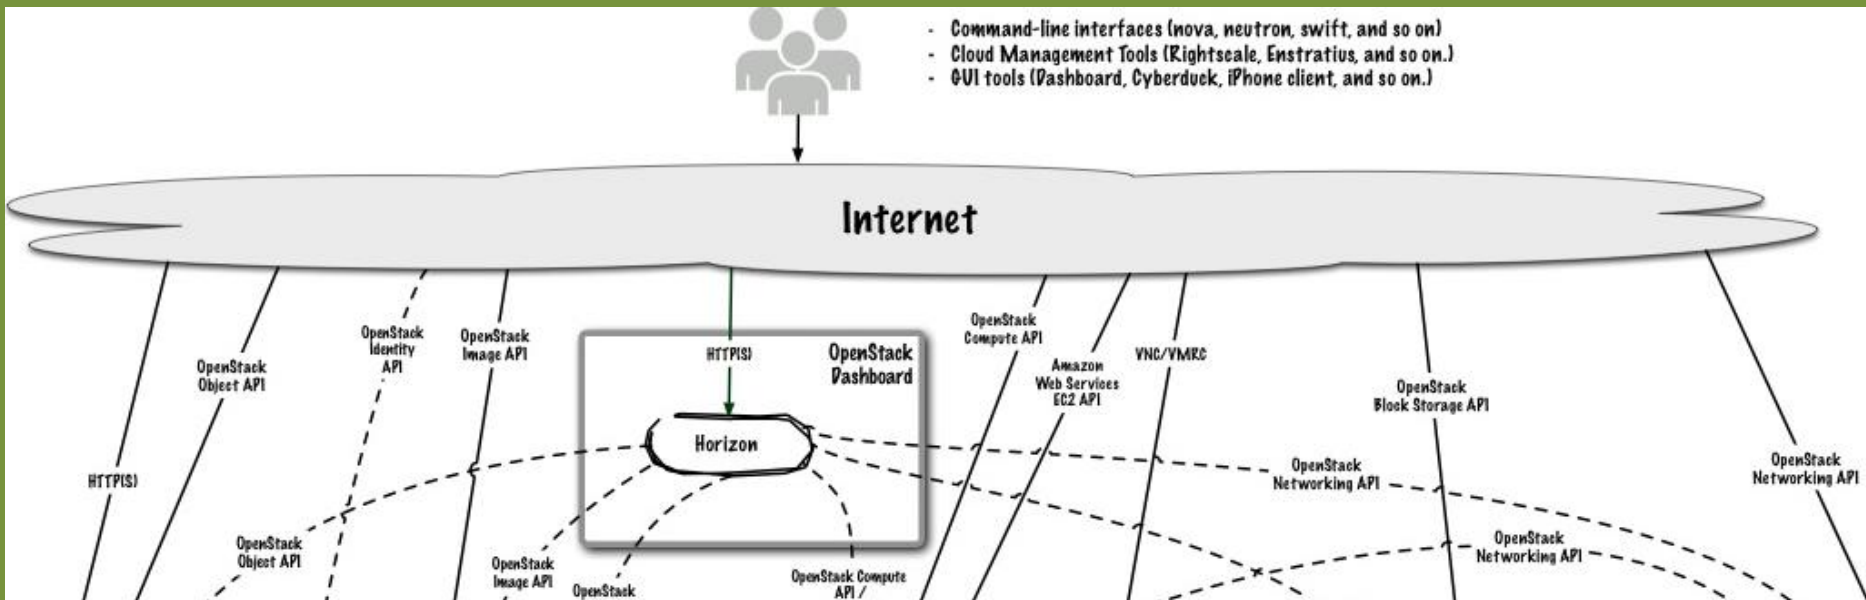
- Components in our focuses:
 - Keystone
 - Nova
 - Glance
 - Networking
 - Orchestration



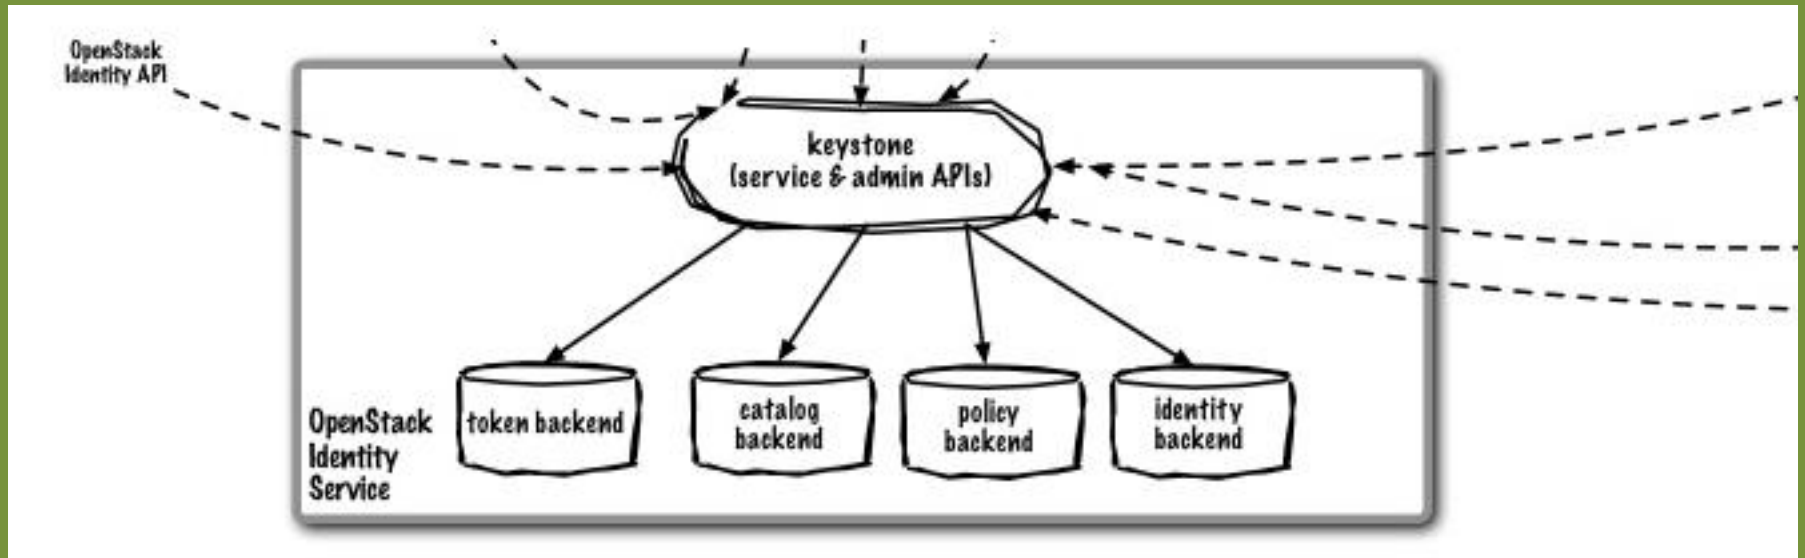
- ## Internet



OpenStack Architecture

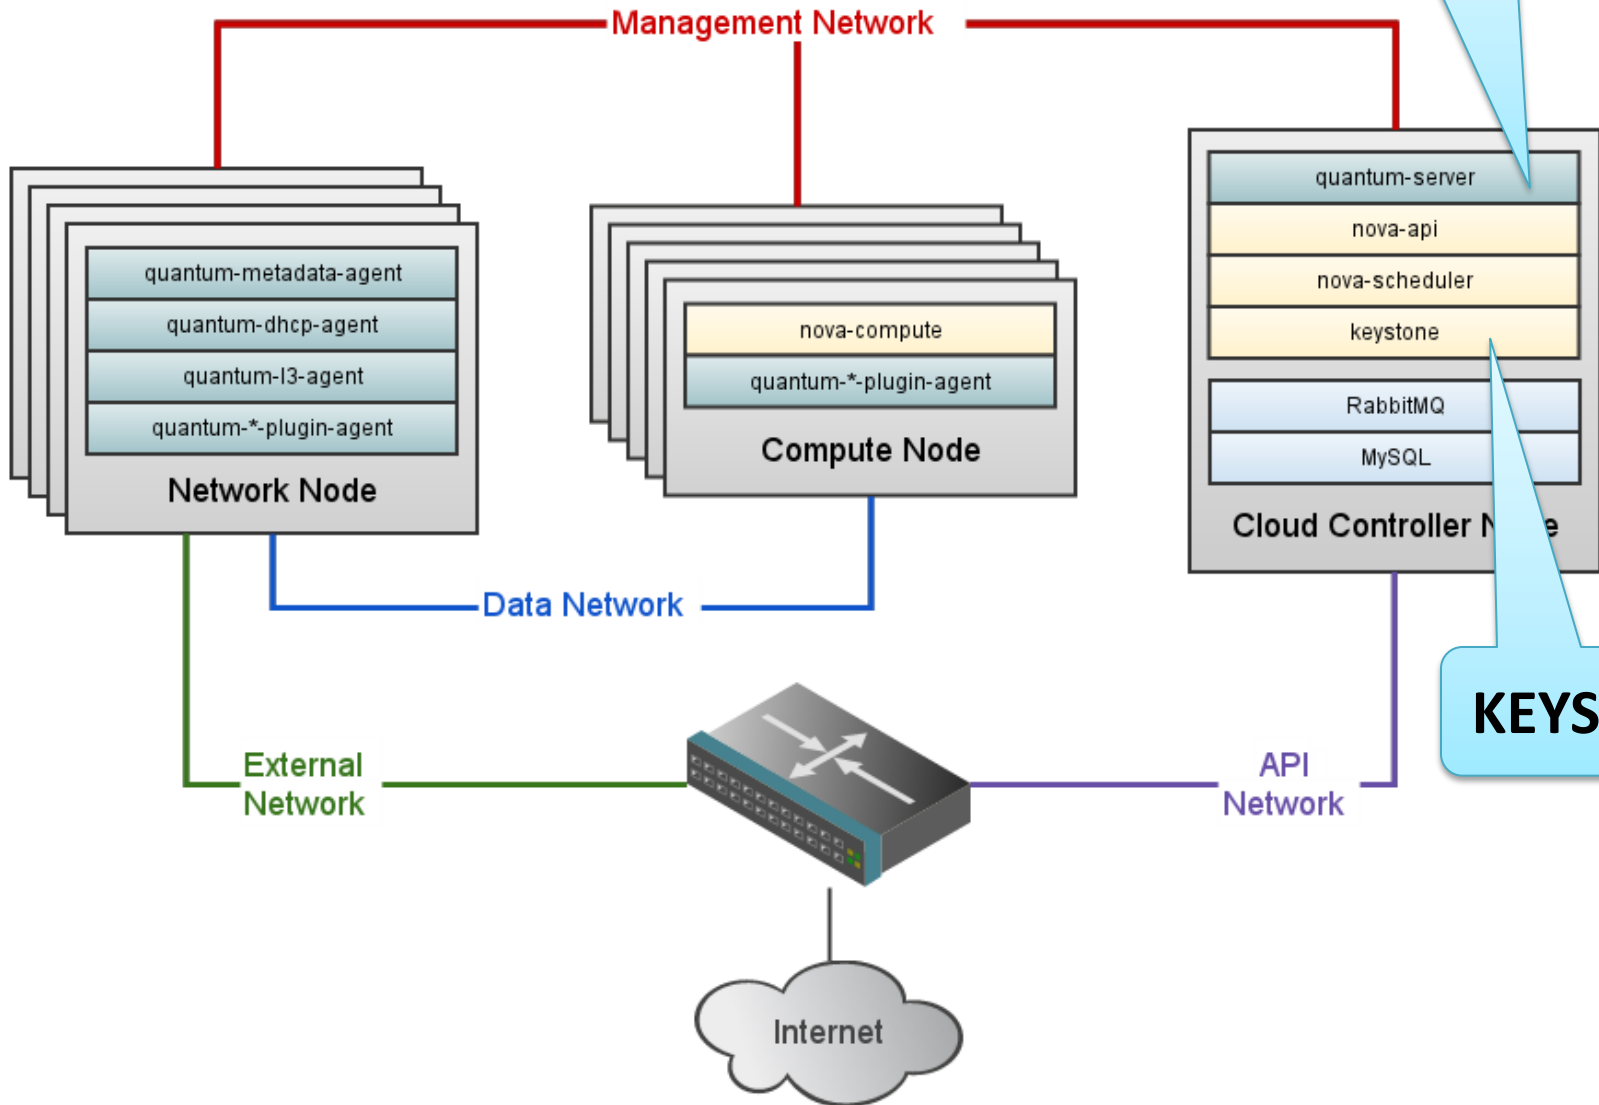


OpenStack Architecture



Component Layout

DASHBOARD



KEYSTONE

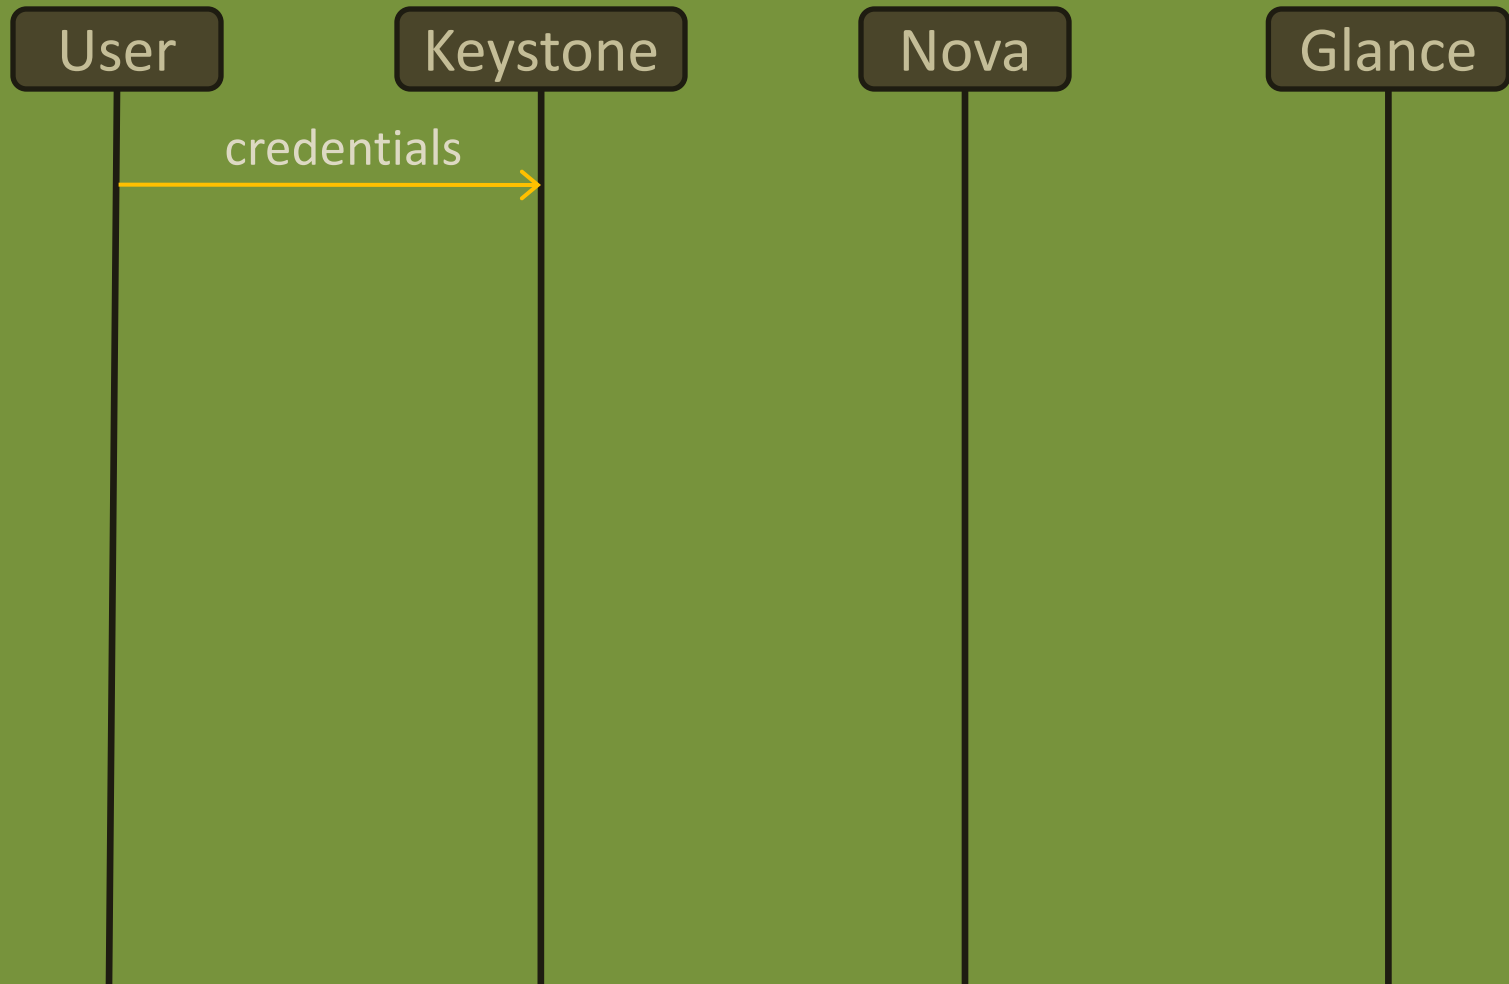
Keystone

- A central authentication and authorization
- **User** represents someone or something that can gain access through Keystone. Users come with credentials that can be checked like passwords or API keys.
- **Tenant** represents what is called **the project** in Nova. Users are bound to a tenant by assigning them a role on that tenant.
- **Role** represents a number of privileges or rights a user has or actions they are allowed to perform.
- To access a service, we have to know its endpoint. So there are endpoint templates in Keystone that provide information about all existing endpoints of all existing services.

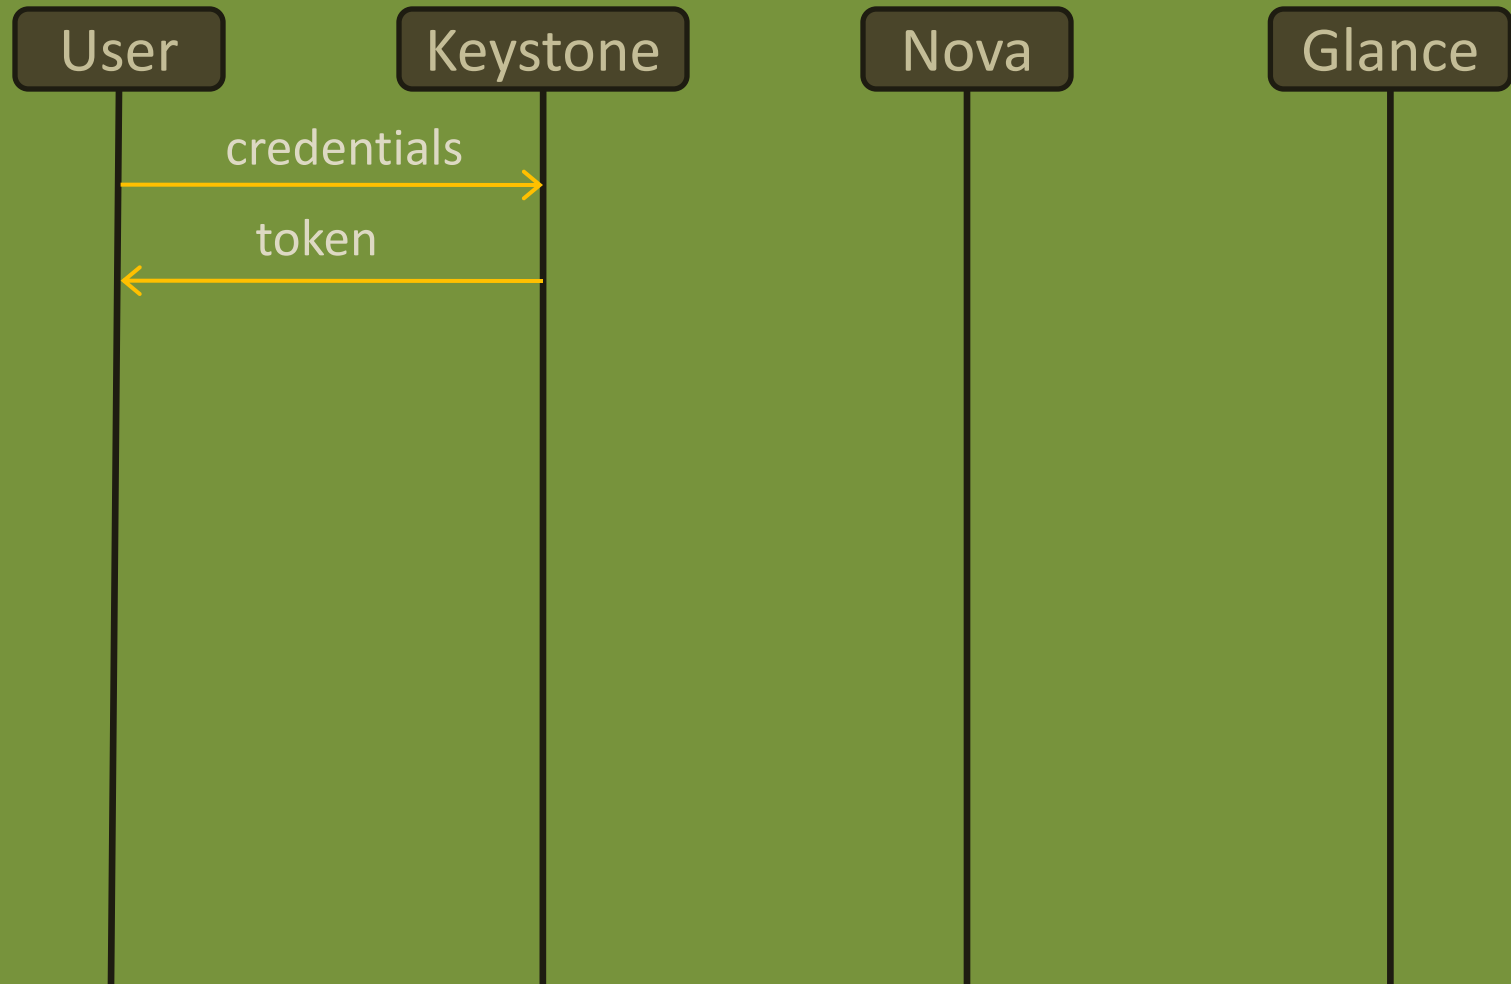
Keystone

- To access some service, users provide their credentials to Keystone and receive a token.
- If the user, for example, wants to spawn a new VM instance in Nova, one can find an URL to Nova in the list of endpoints provided by Keystone and send an appropriate request.
- After that, Nova verifies the validity of the token in Keystone and should create an instance from some image by the provided image ID and plug it into some network.
- All the way this token travels between services so that they can ask Keystone or each other for additional information or some actions.

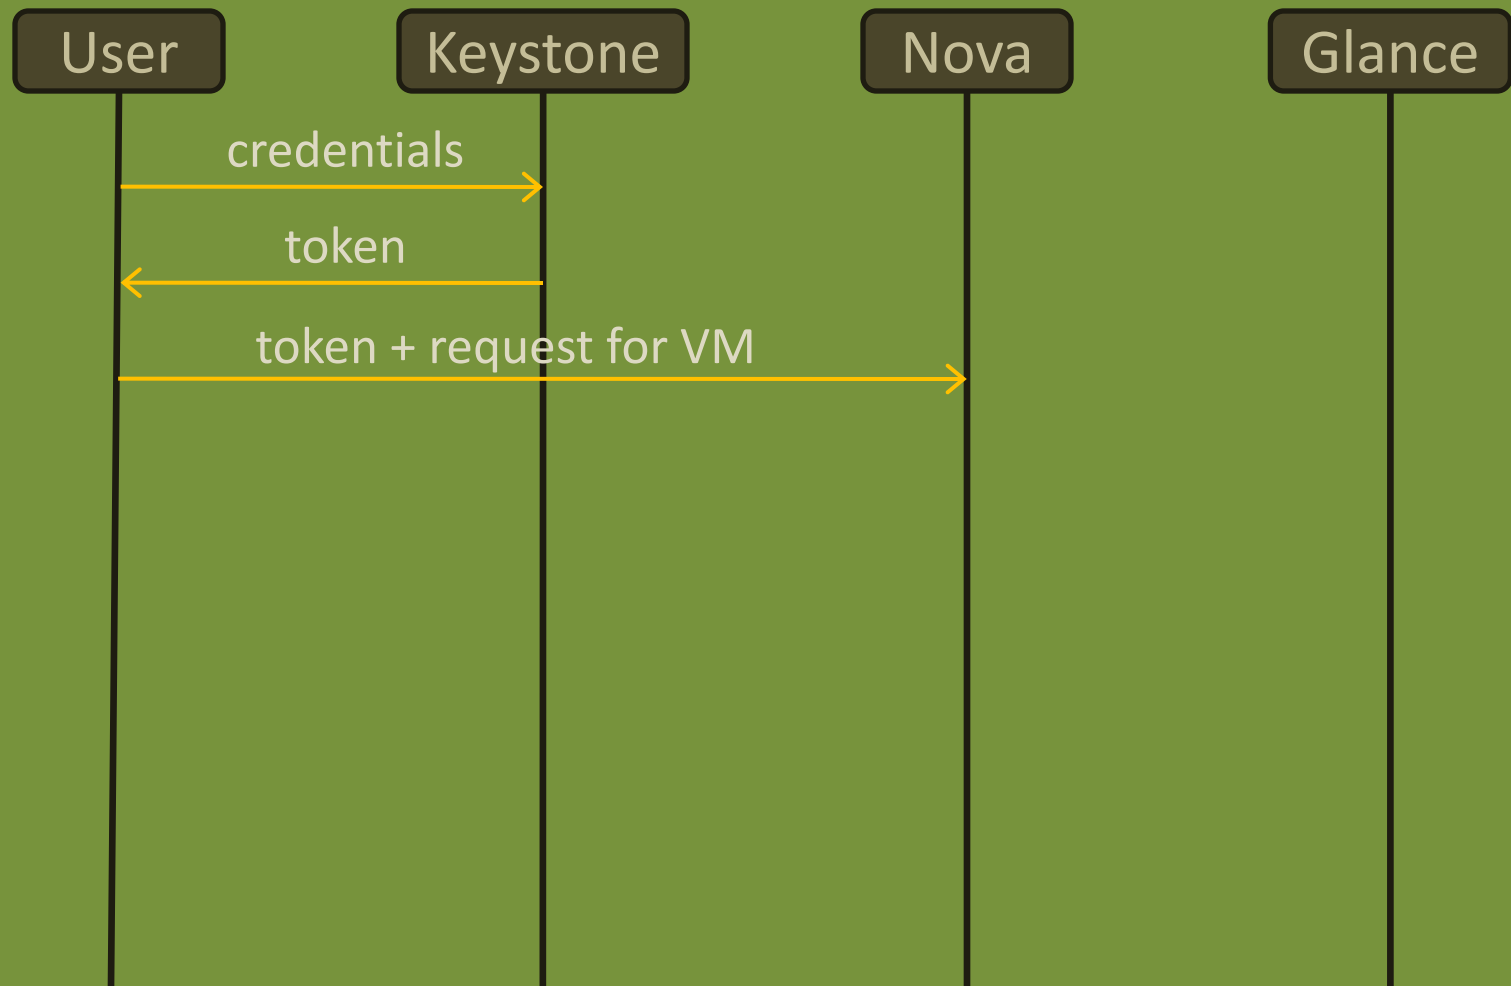
Keystone Control Flow



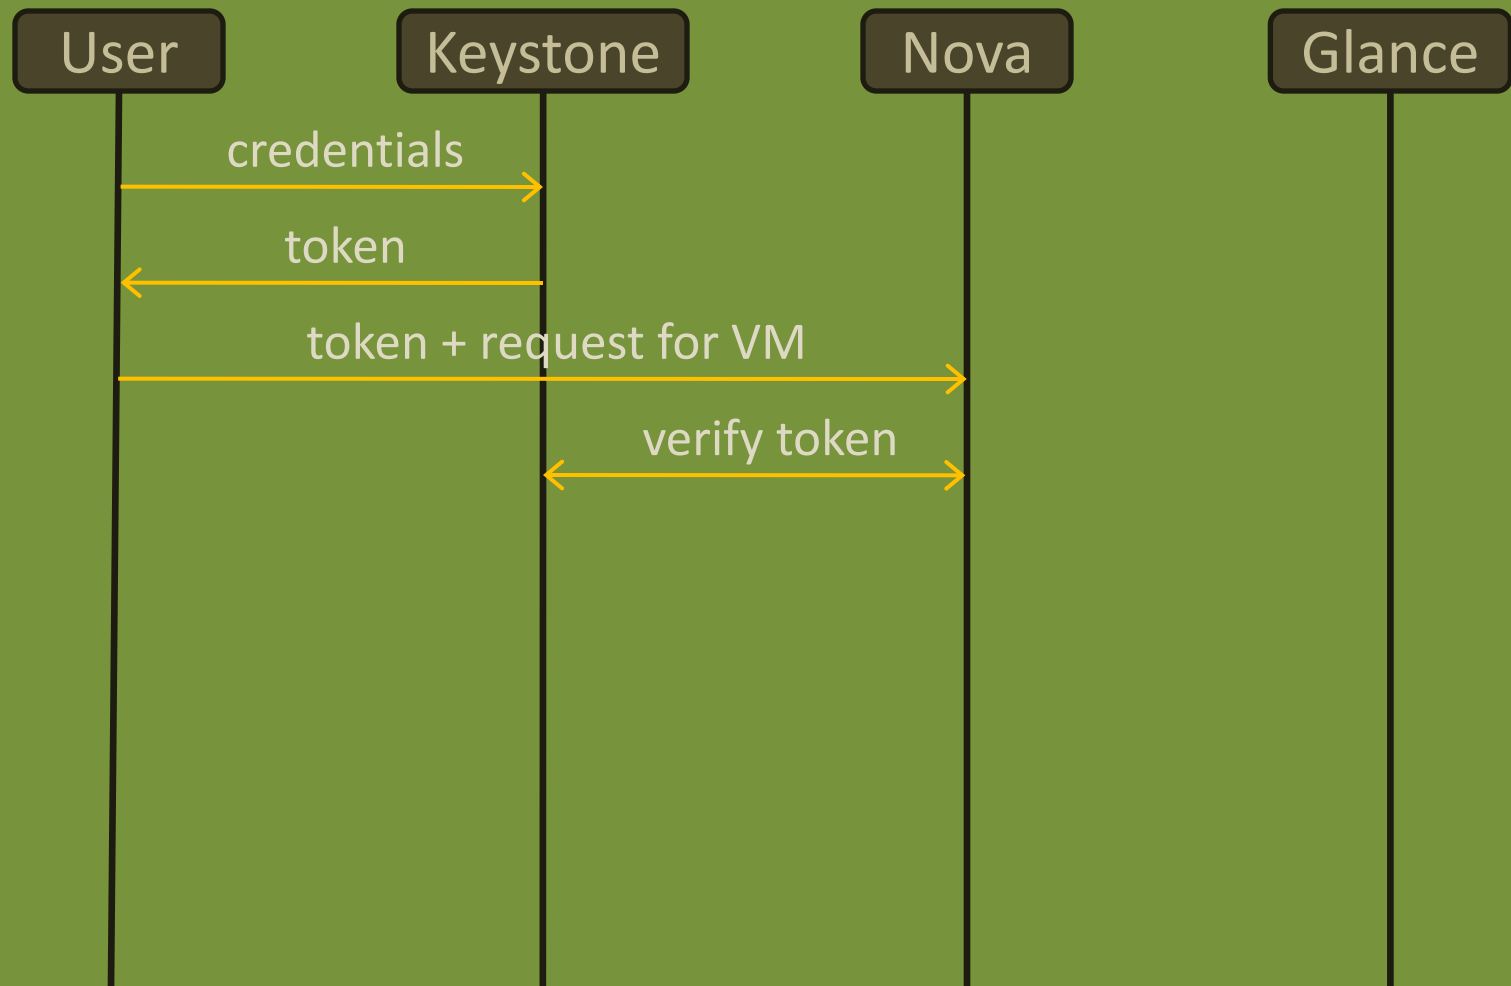
Keystone Control Flow



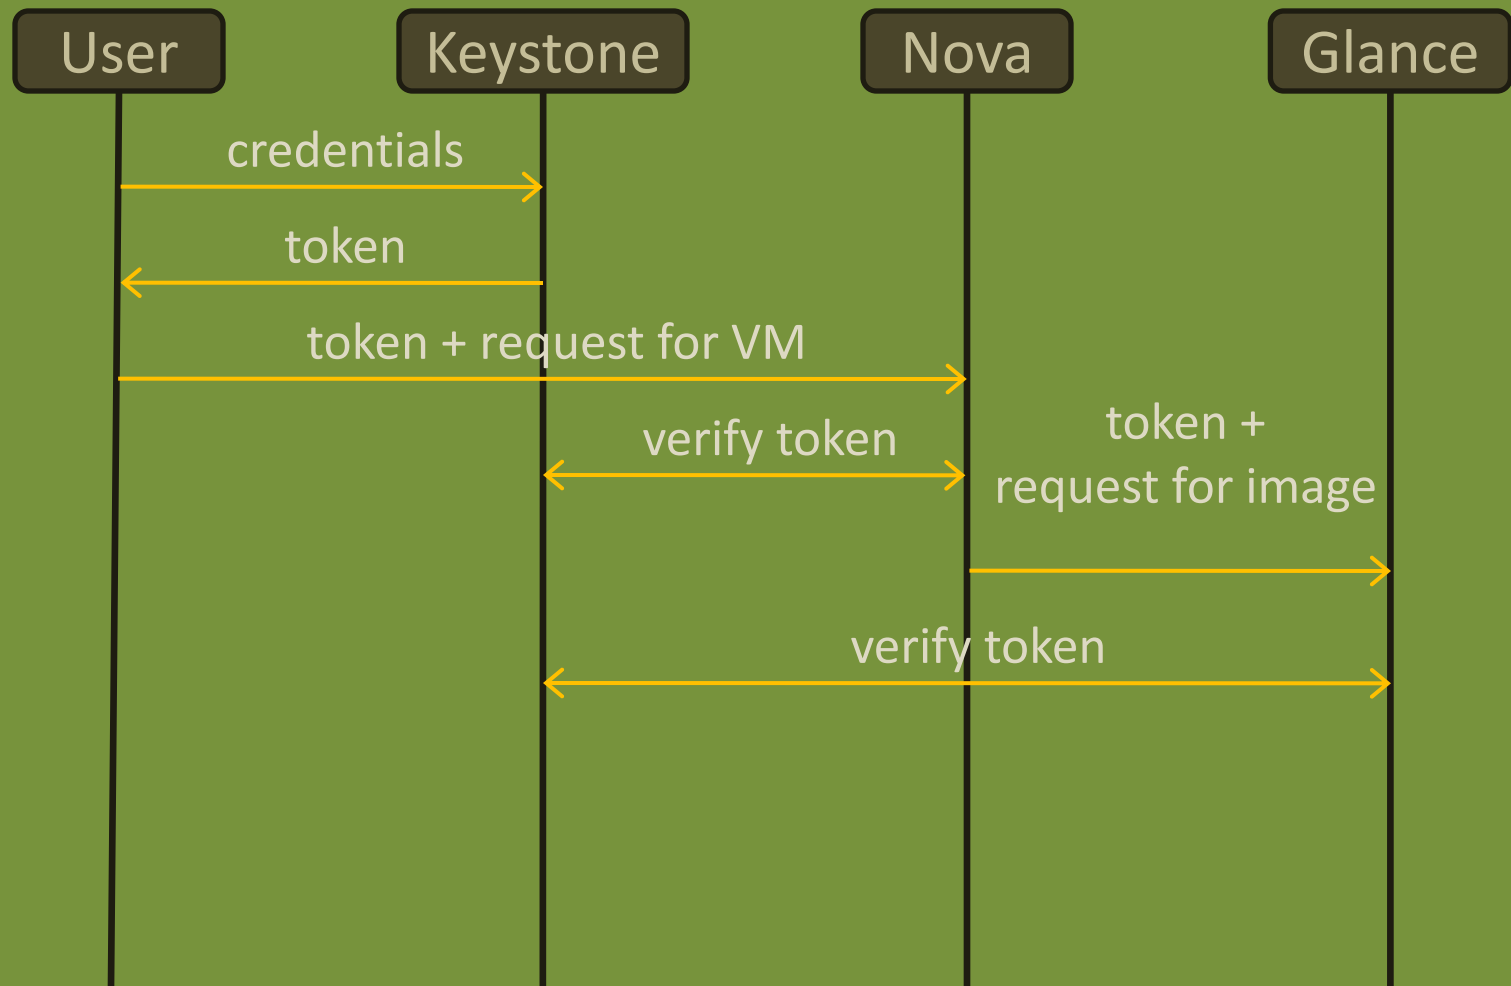
Keystone Control Flow



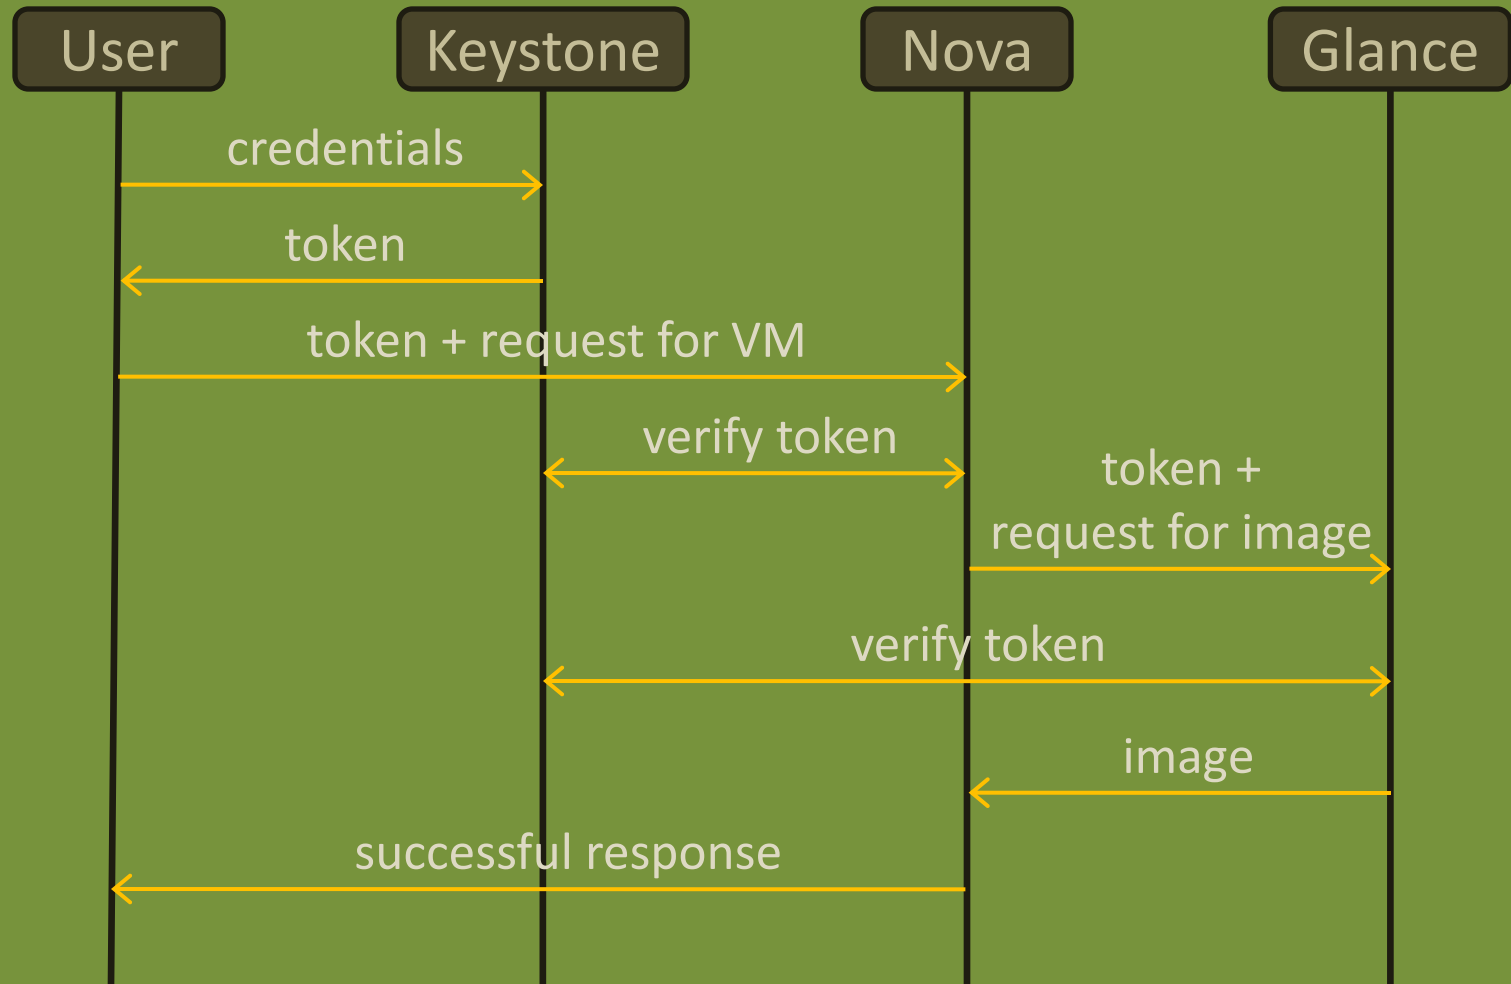
Keystone Control Flow

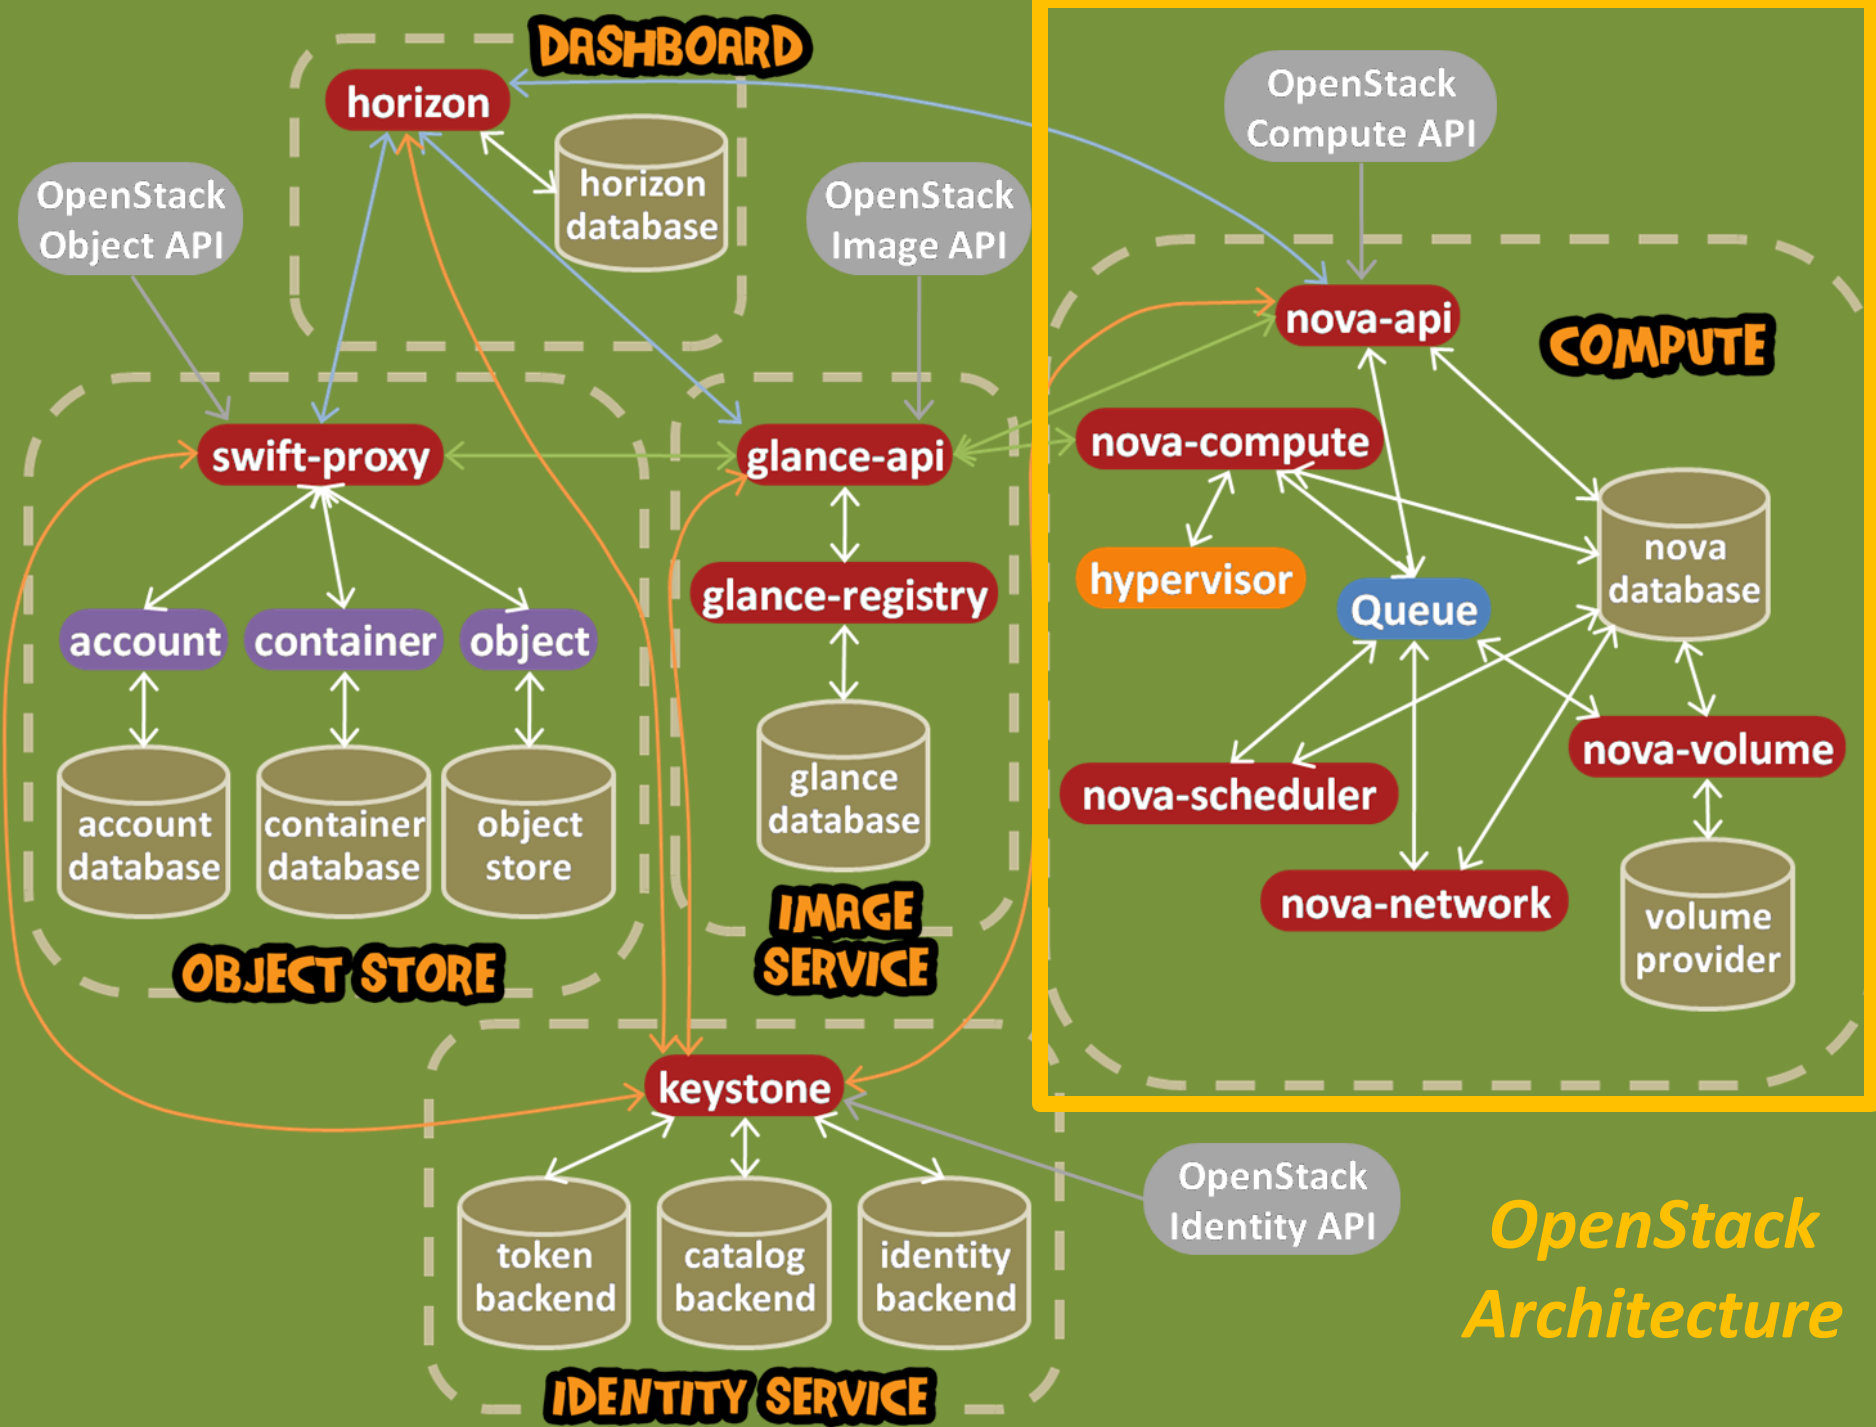


Keystone Control Flow




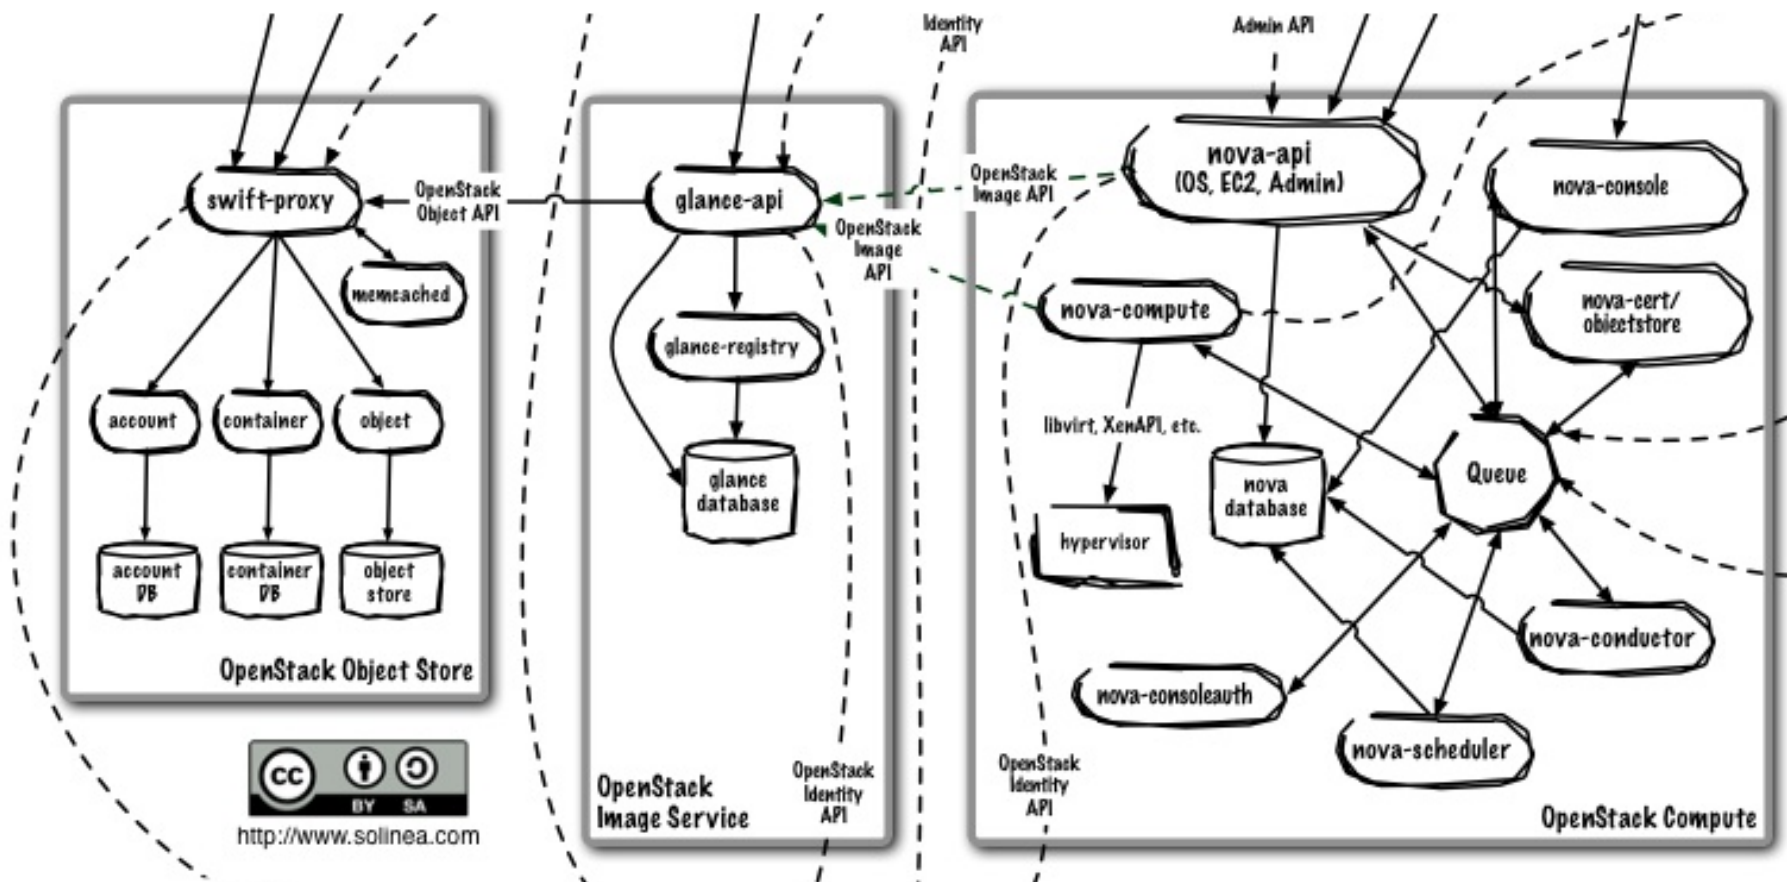
Keystone Control Flow





- ## Internet

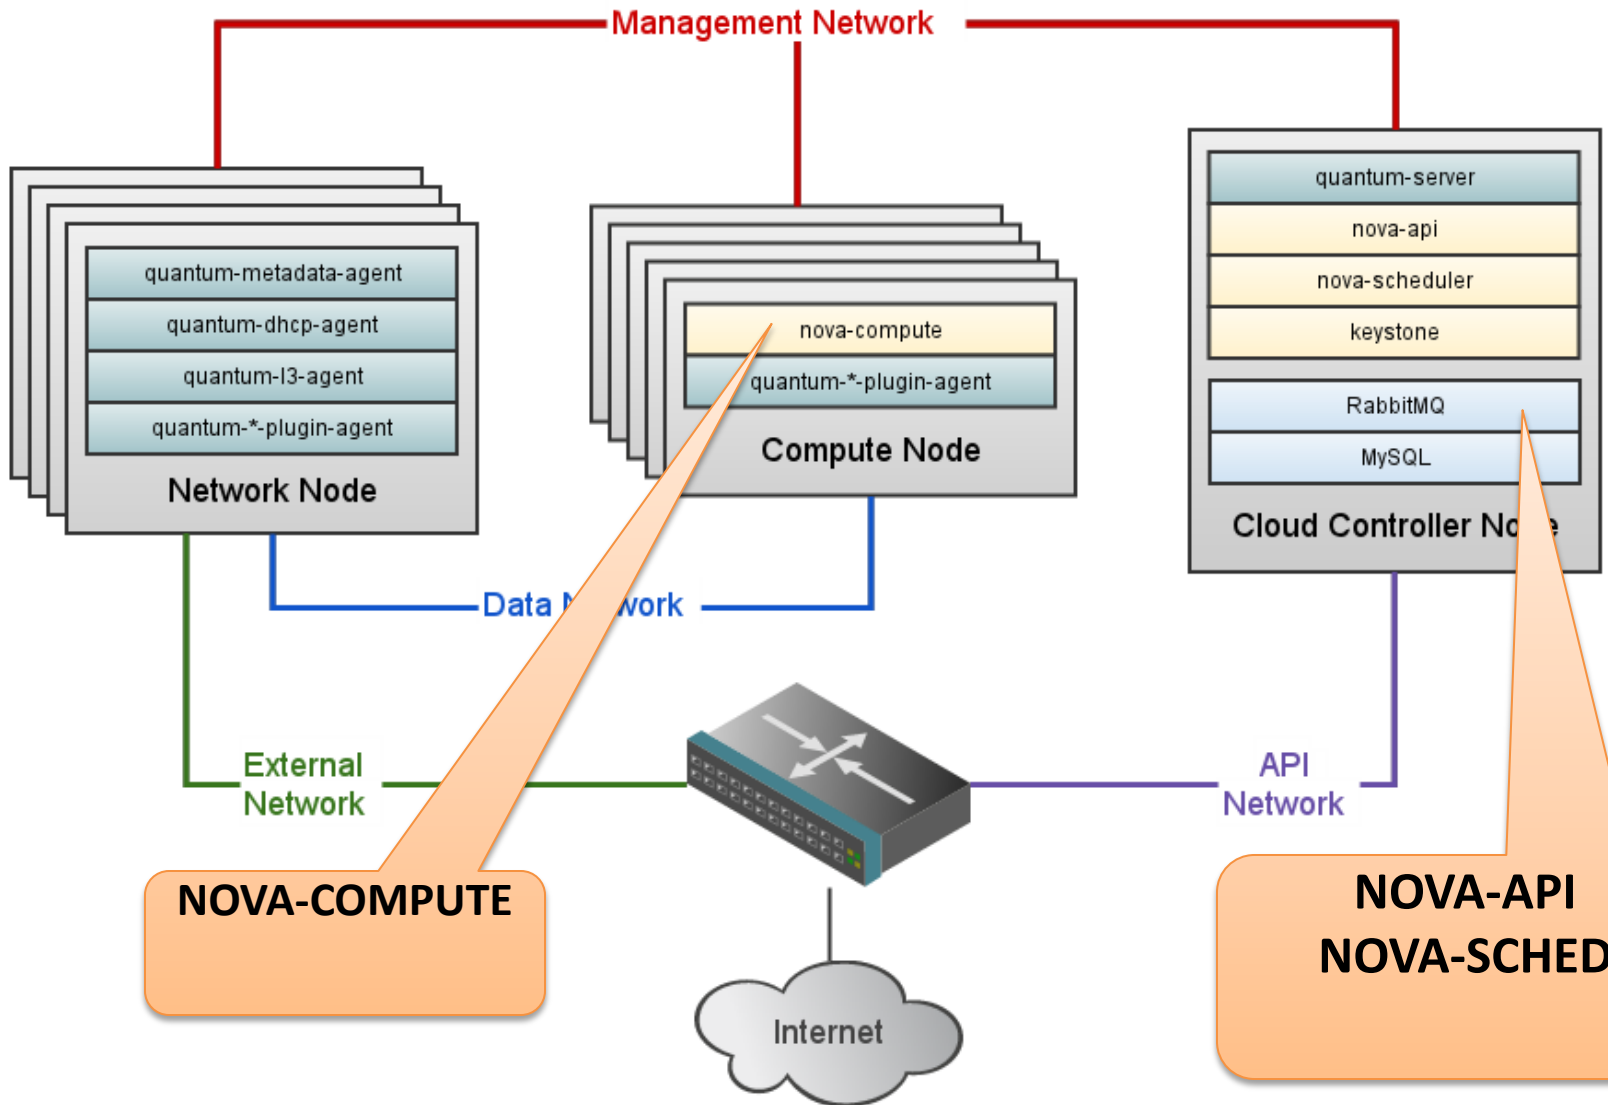




Nova

- Nova handles instances provisioning on compute resources.
- Nova-api initiates most activities
- Nova components communicate via queue and nova database
- Nova-scheduler decides where to launch instances
- Nova-compute launches instances
- Nova-compute periodically report host and network capabilities to Nova-scheduler

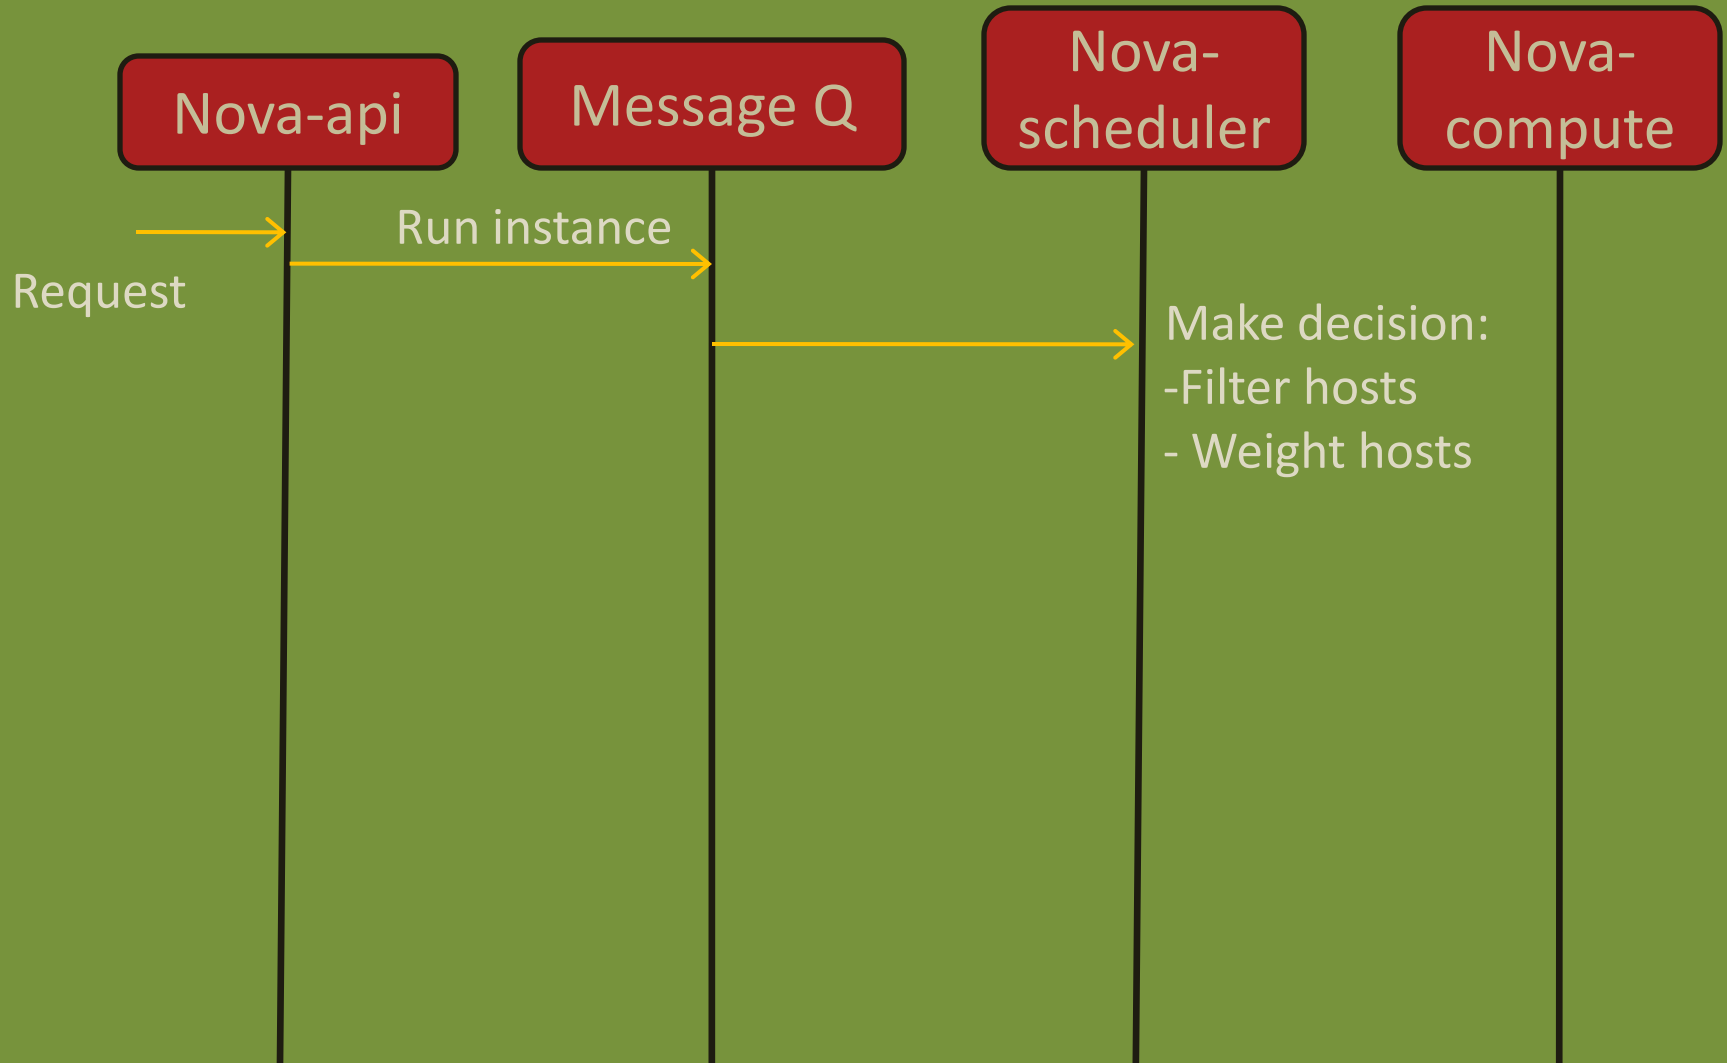
Component Layout



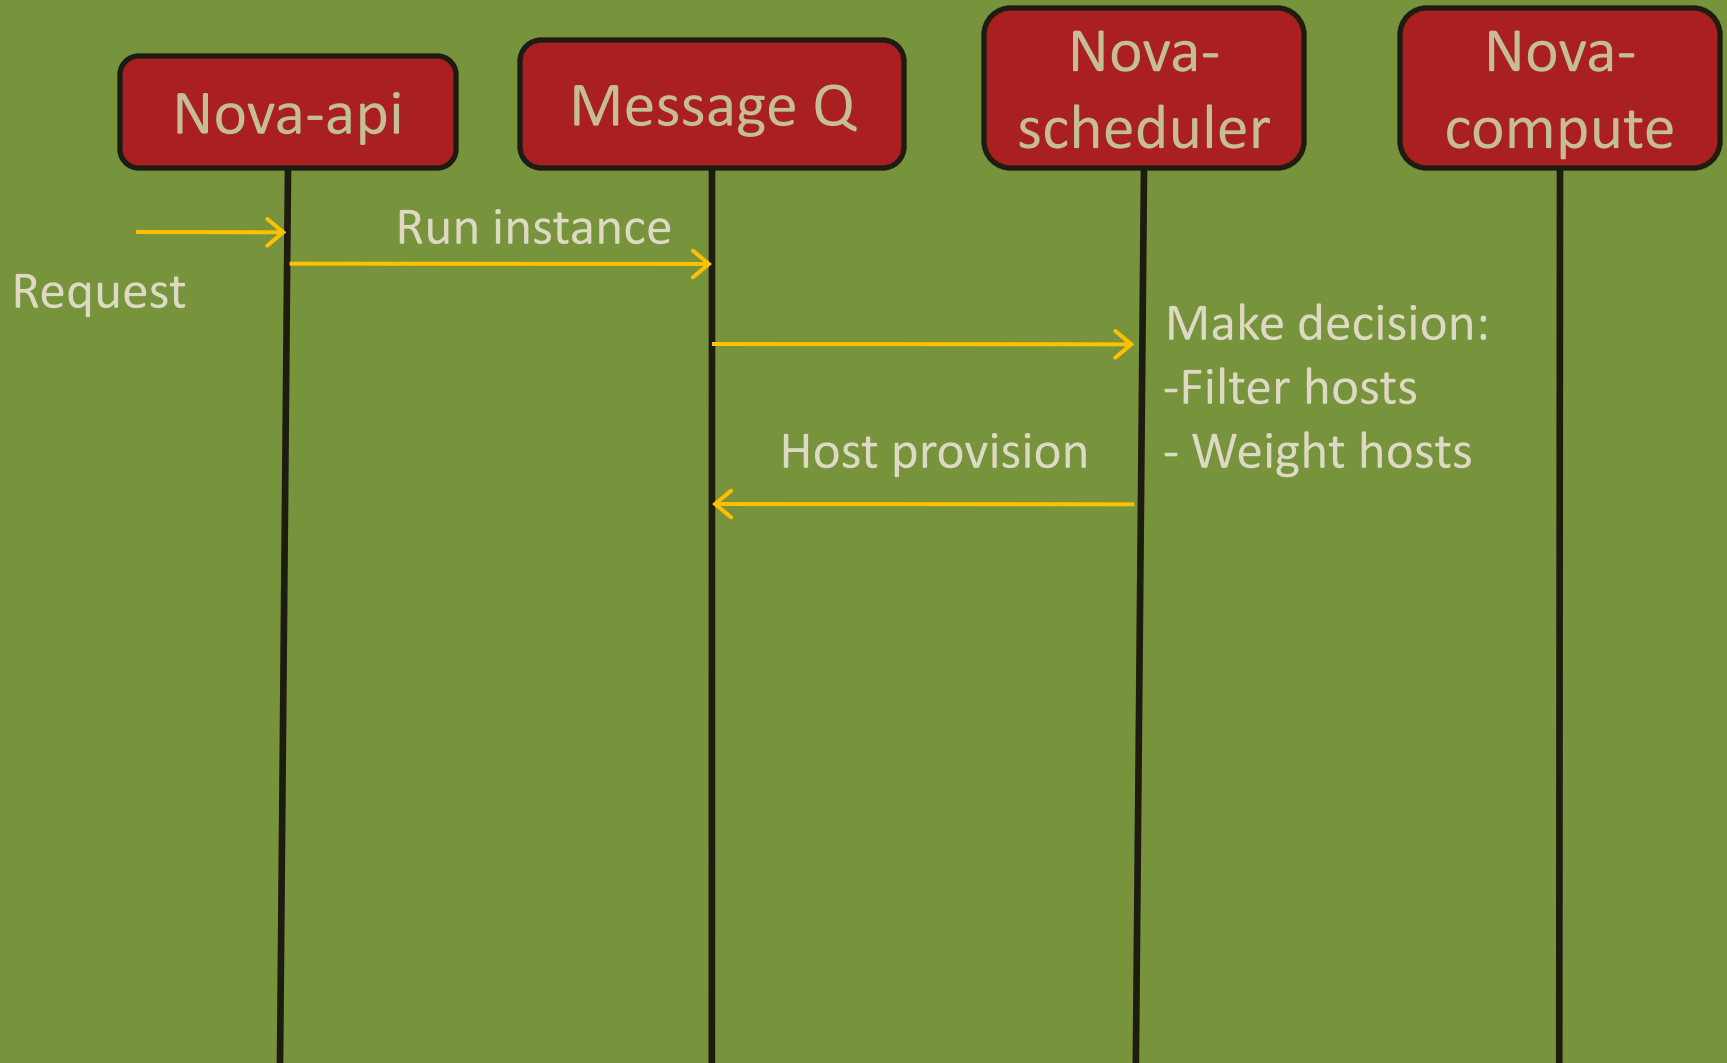
Nova Control Flow



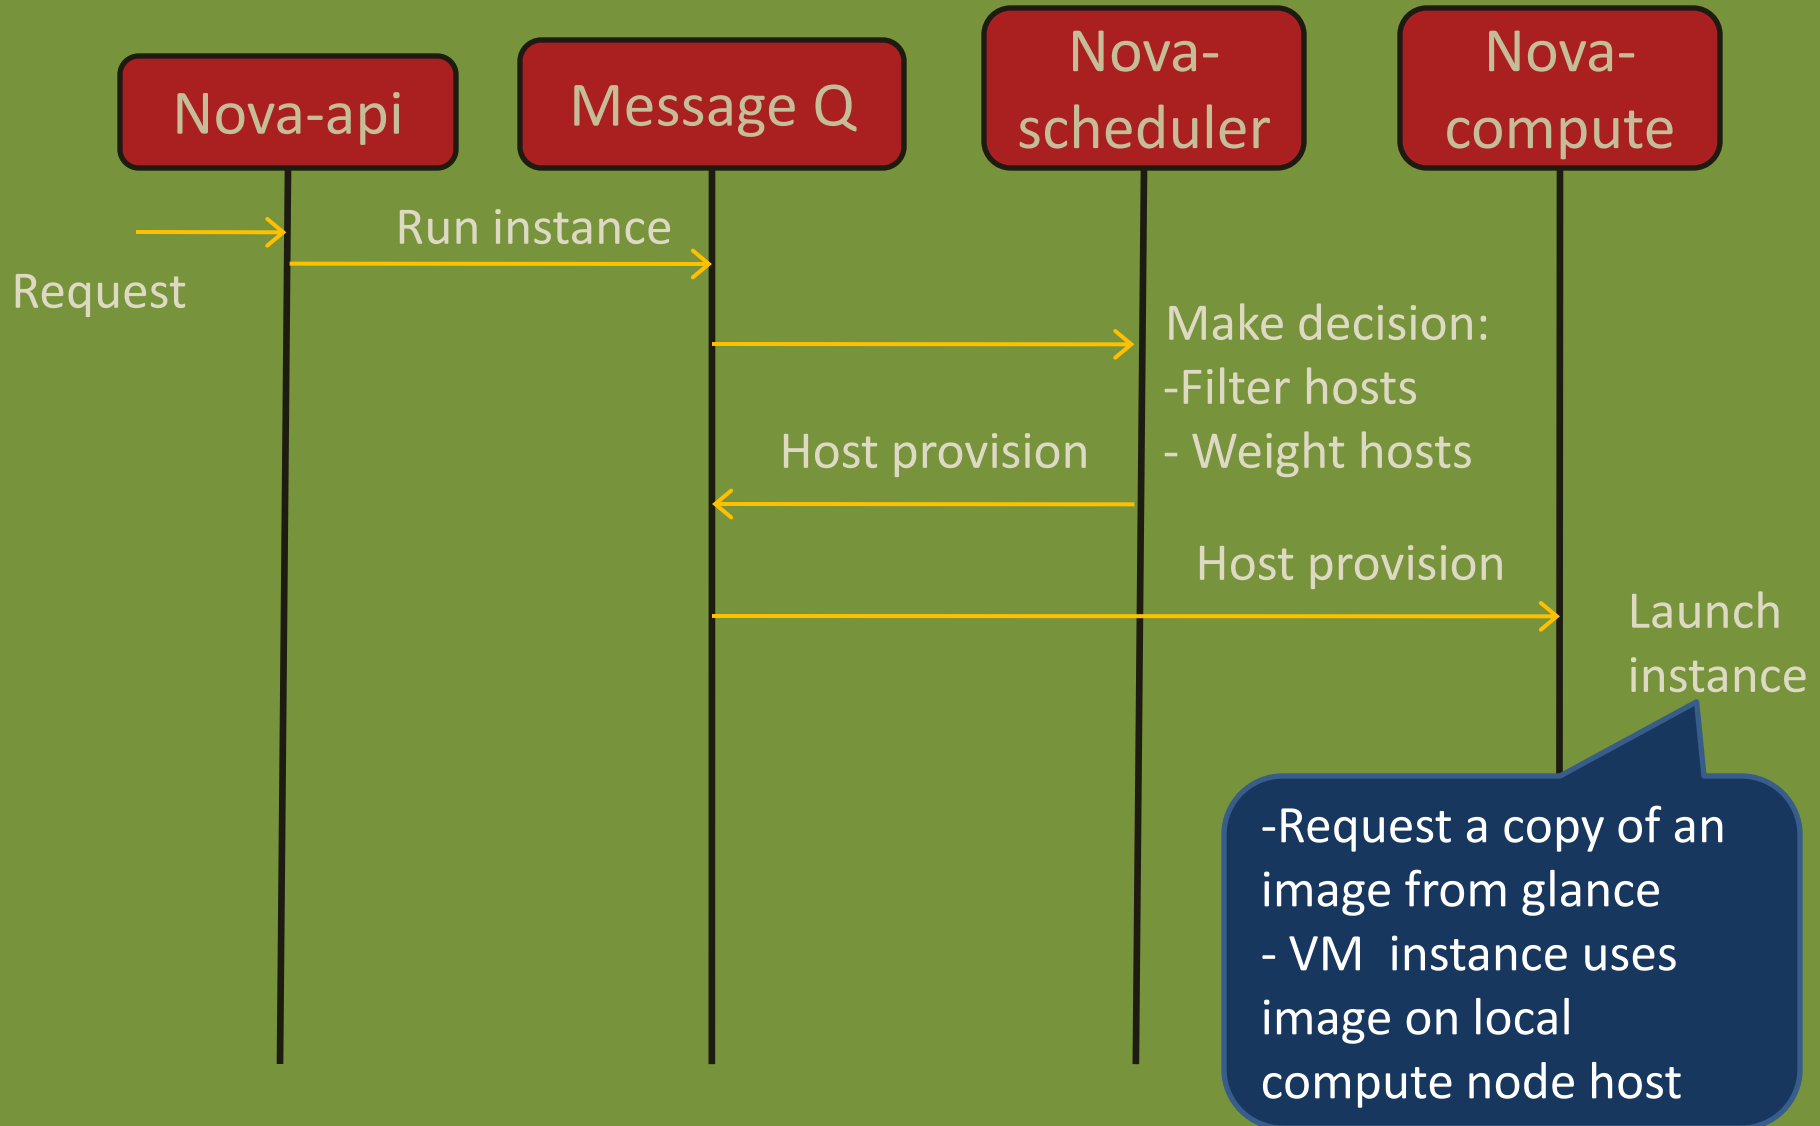
Nova Control Flow



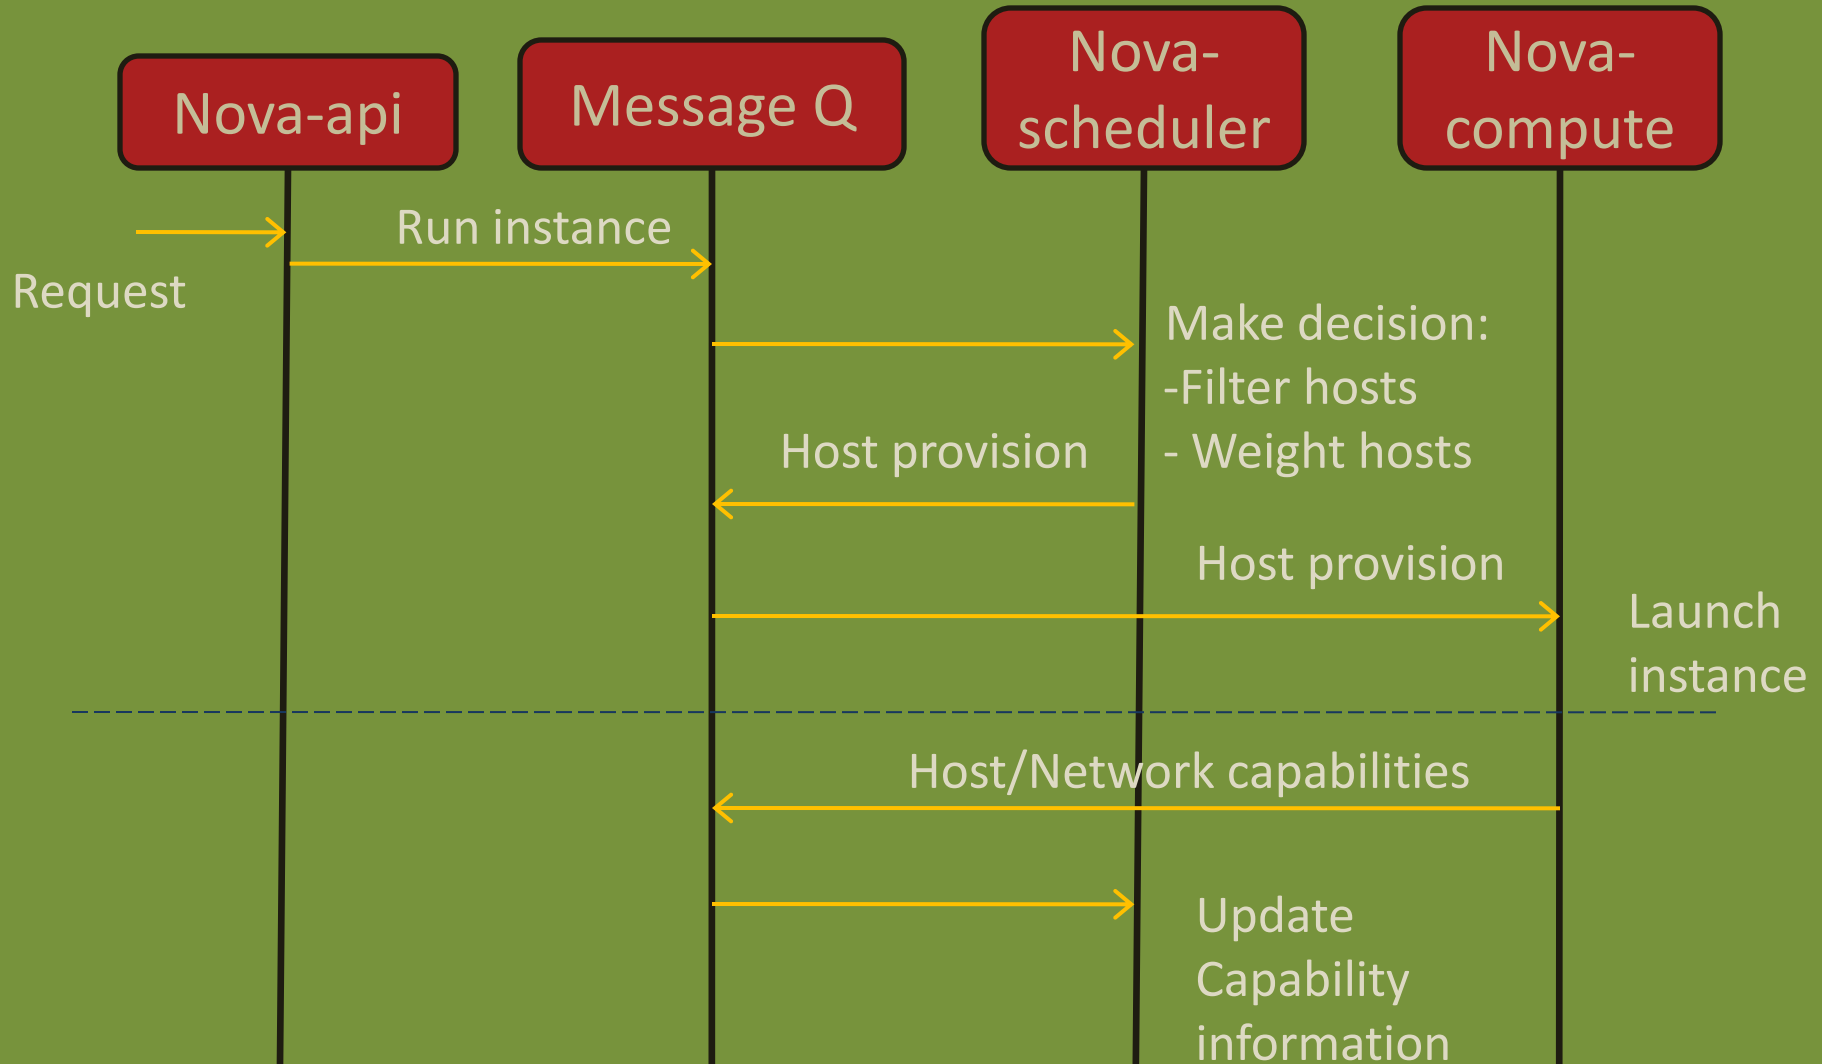
Nova Control Flow

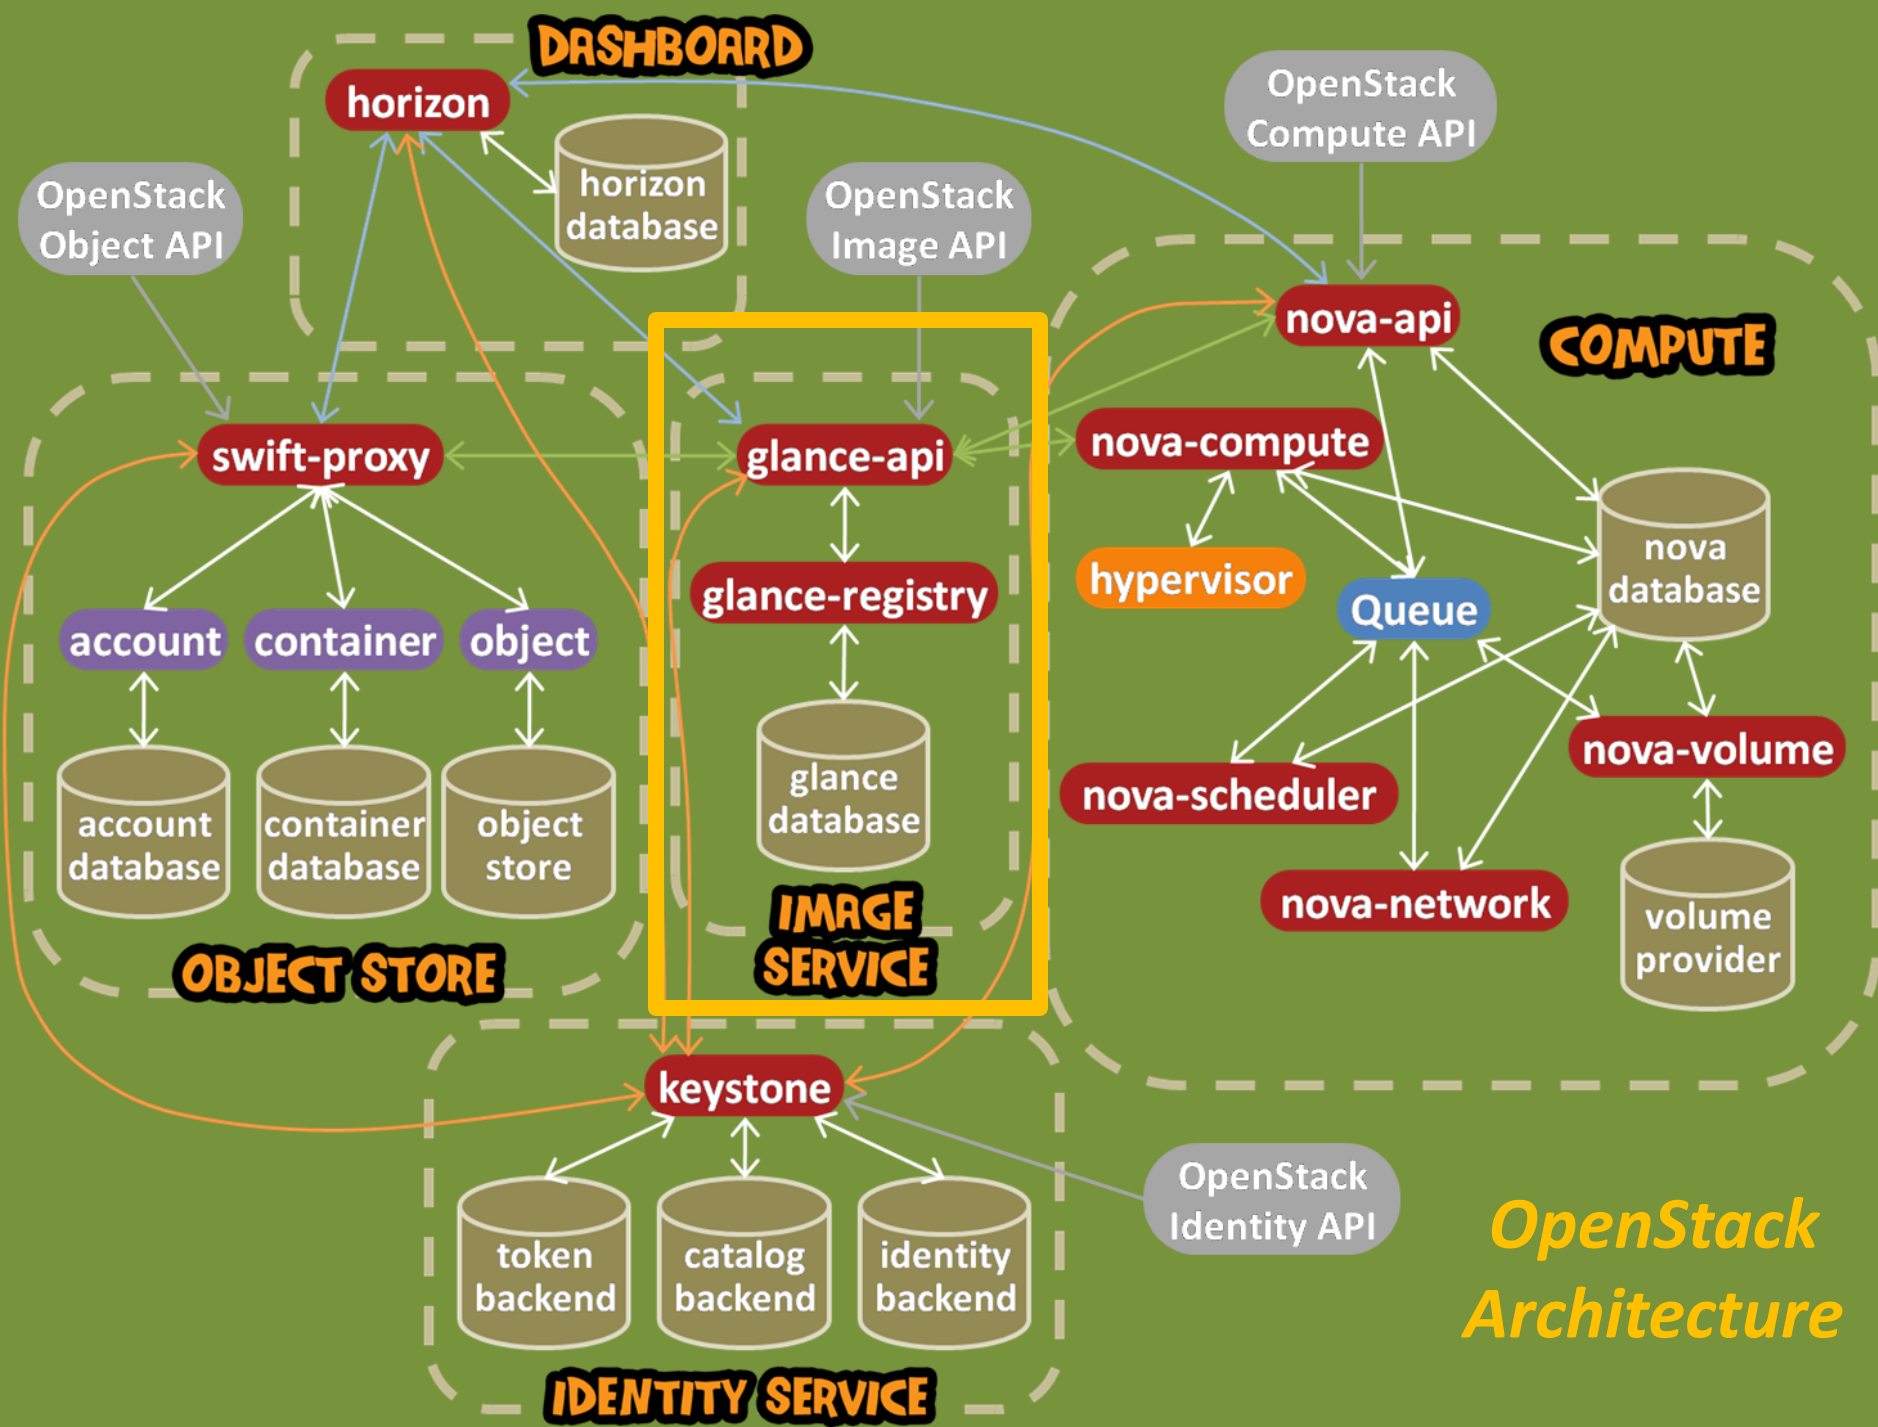


Nova Control Flow



Nova Control Flow

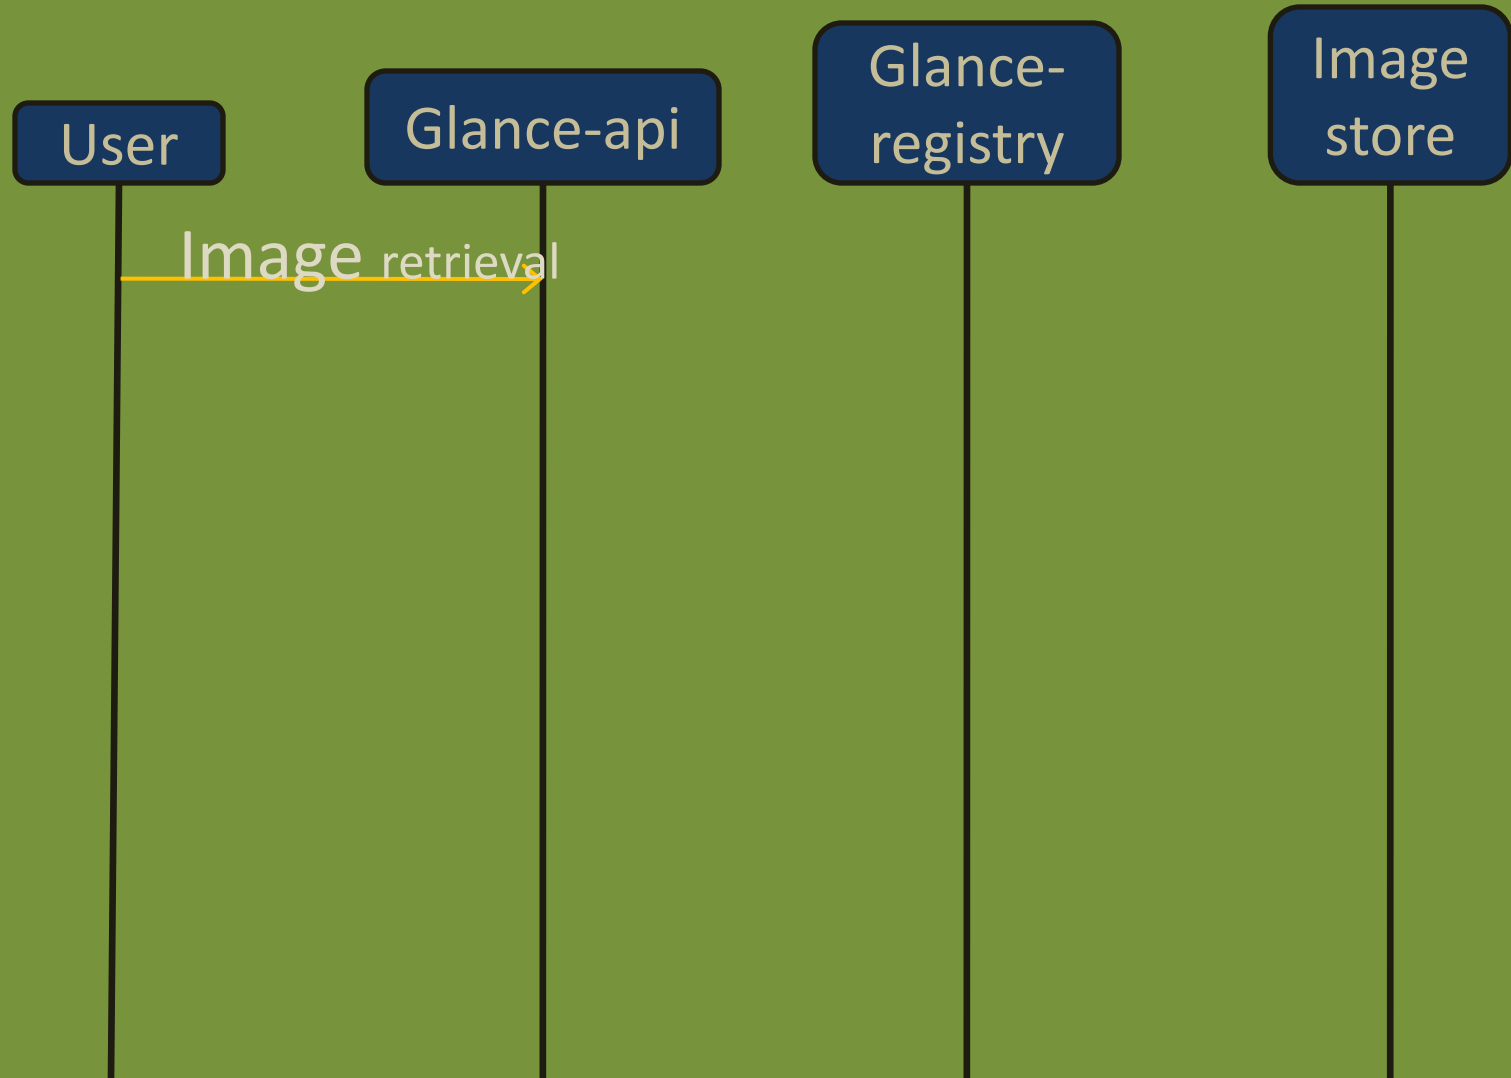




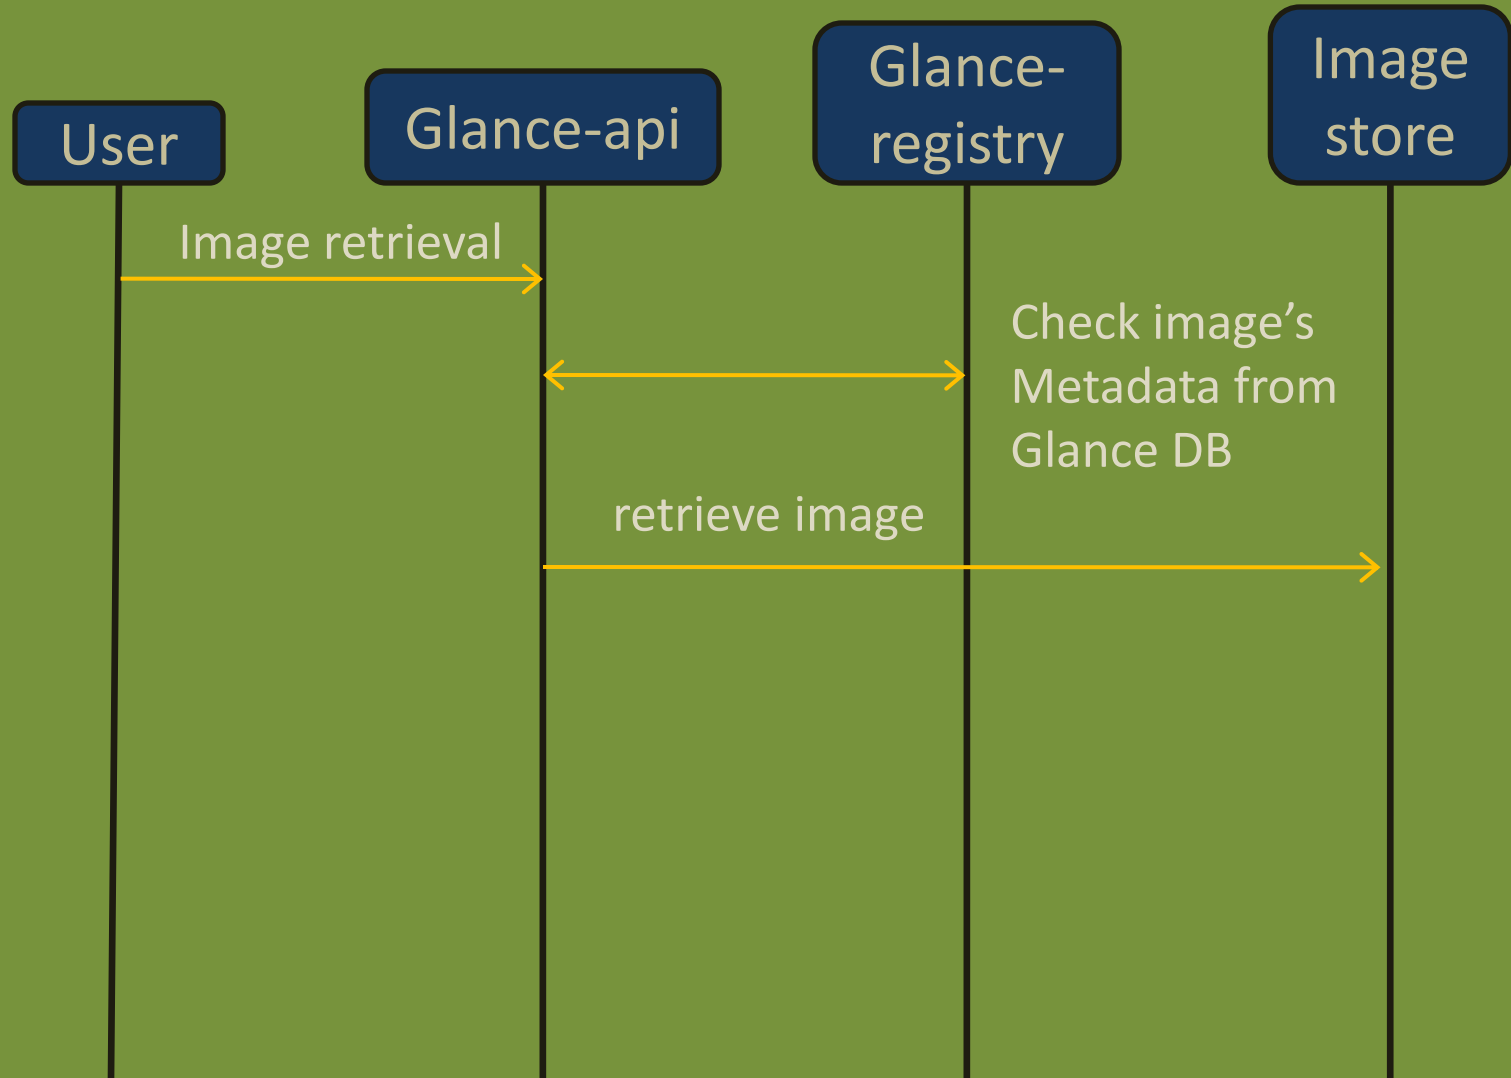
Glance

- Glance manage all kinds of images to instantiate VM instances
- **Glance-api** takes image retrieval requests from nova-compute and pass them to glance-registry
 - OpenStack create **a new copy** of the image on a host where the VM instance runs
- **Glance-registry** check image metadata from database
- Glance stores Image data in its **image store** (S3, HTTP, Local, Swift)

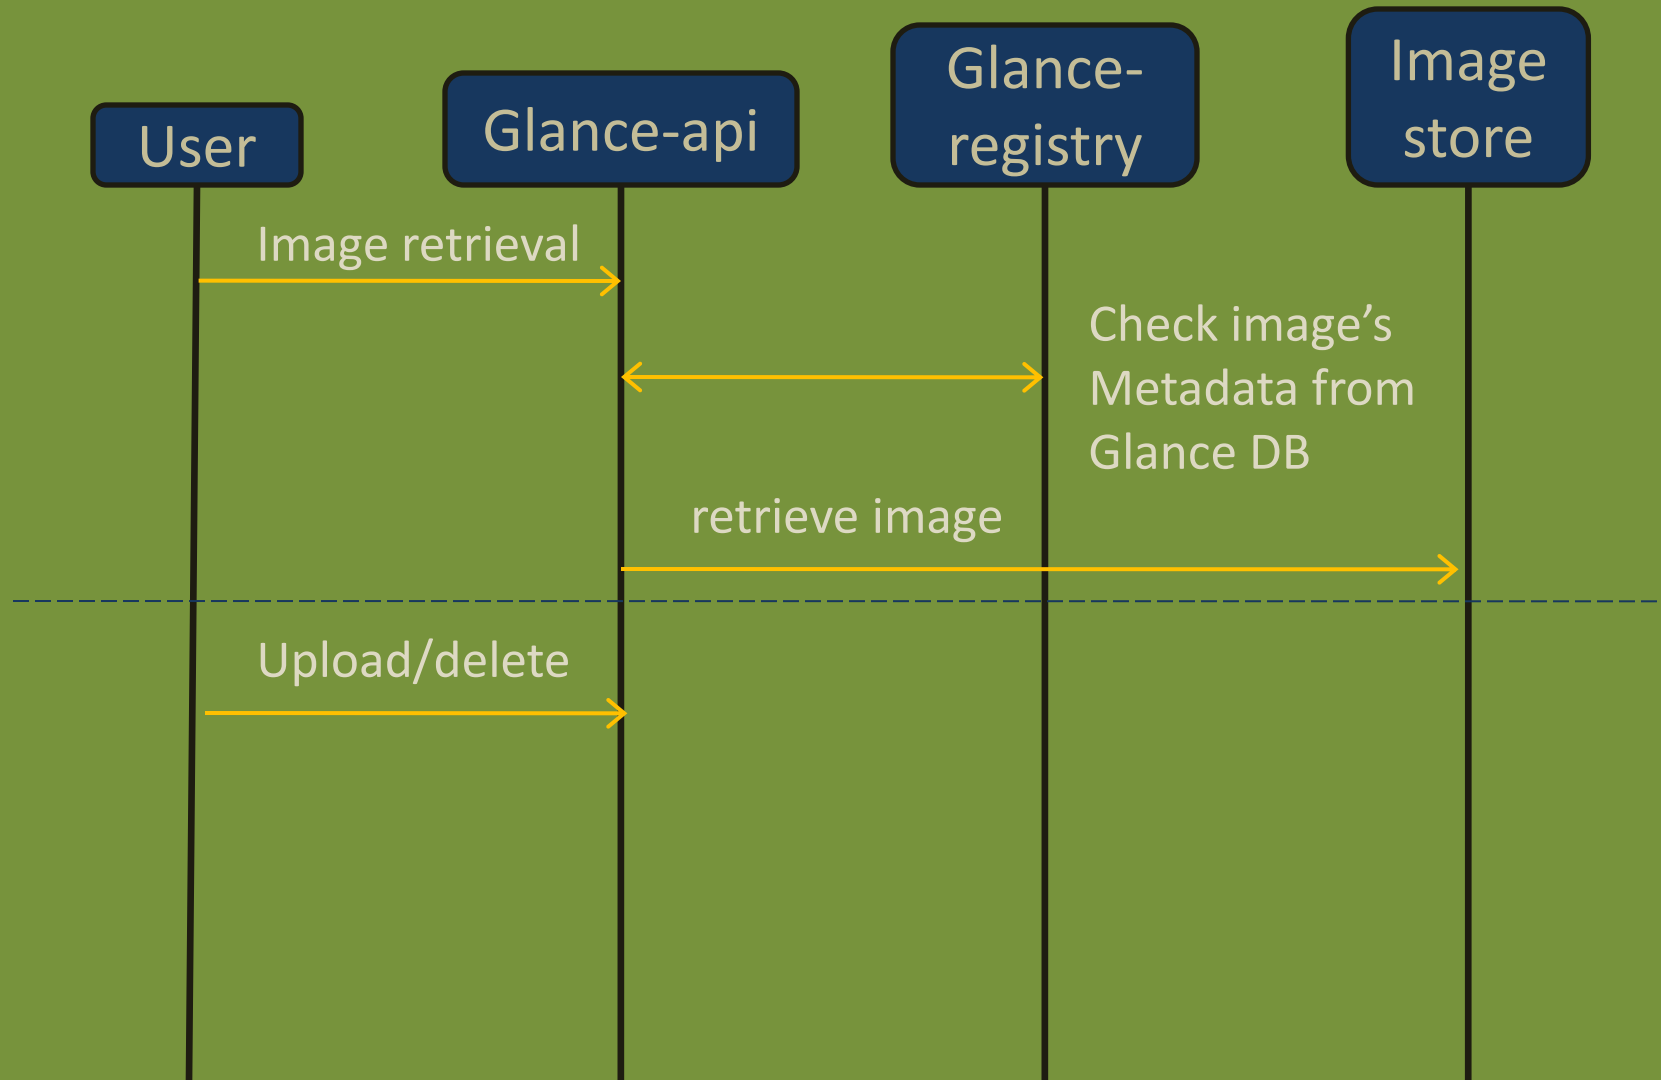
Glance Control Flow



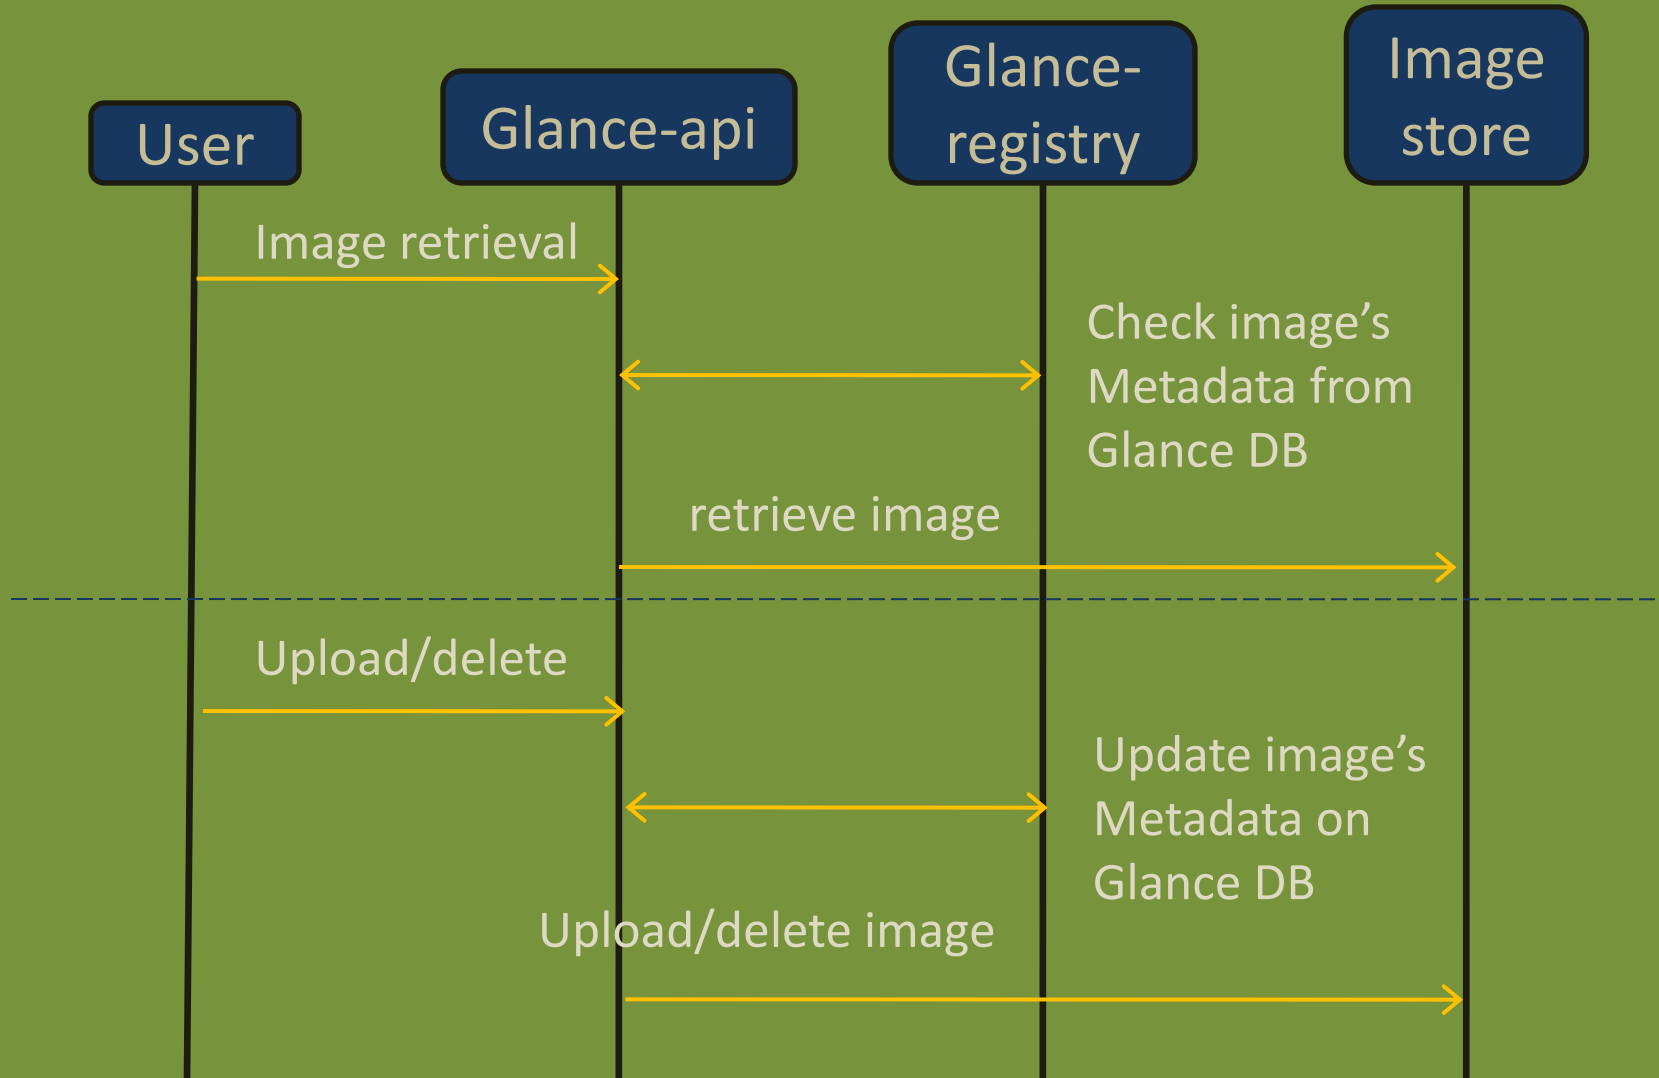
Glance Control Flow



Glance Control Flow



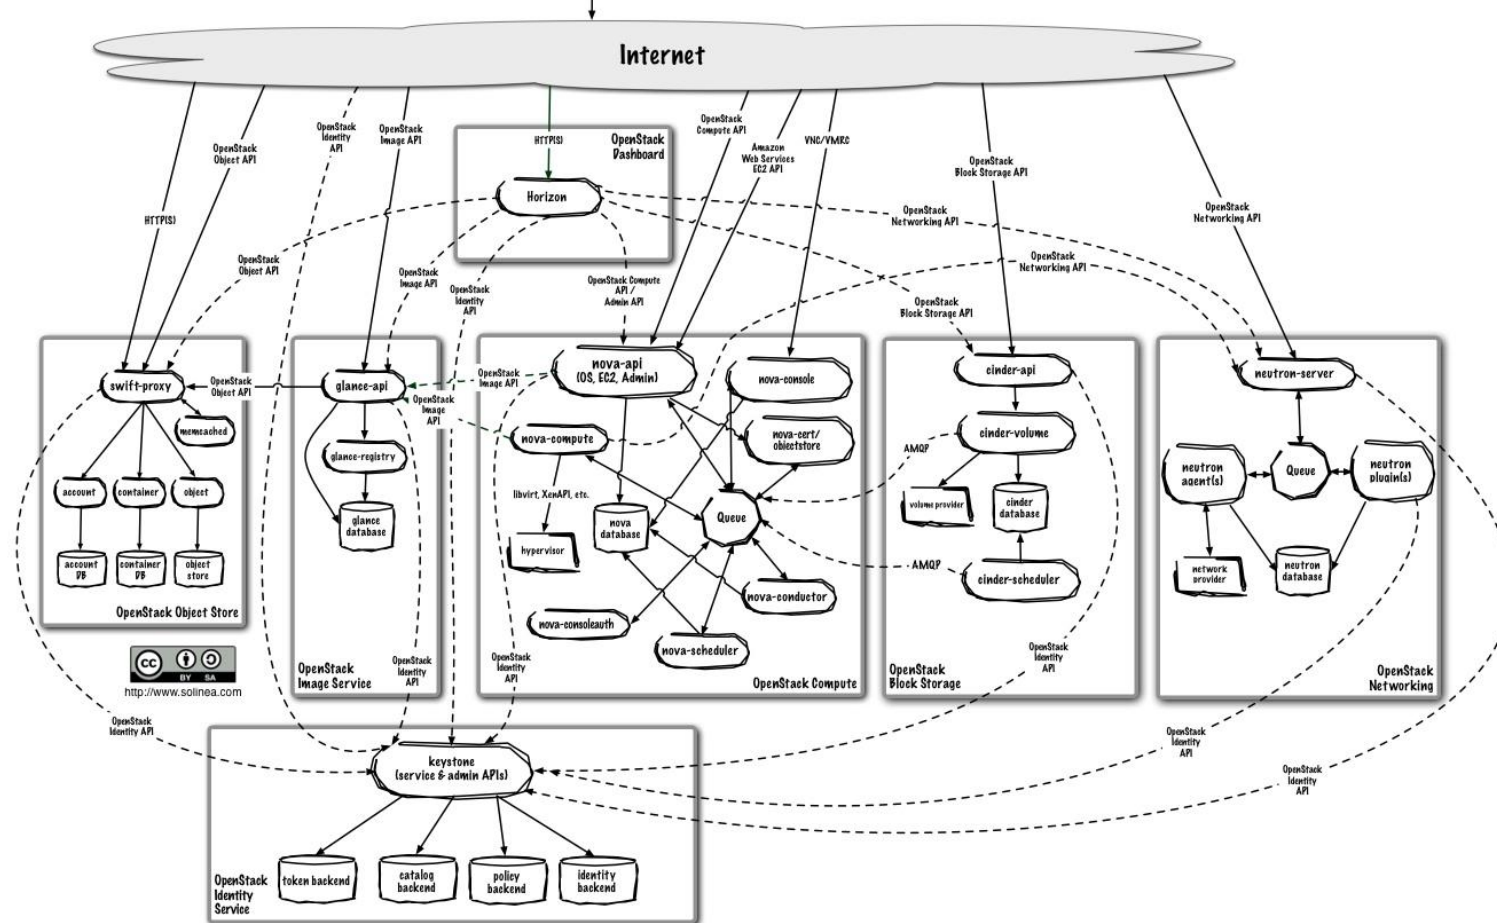
Glance Control Flow

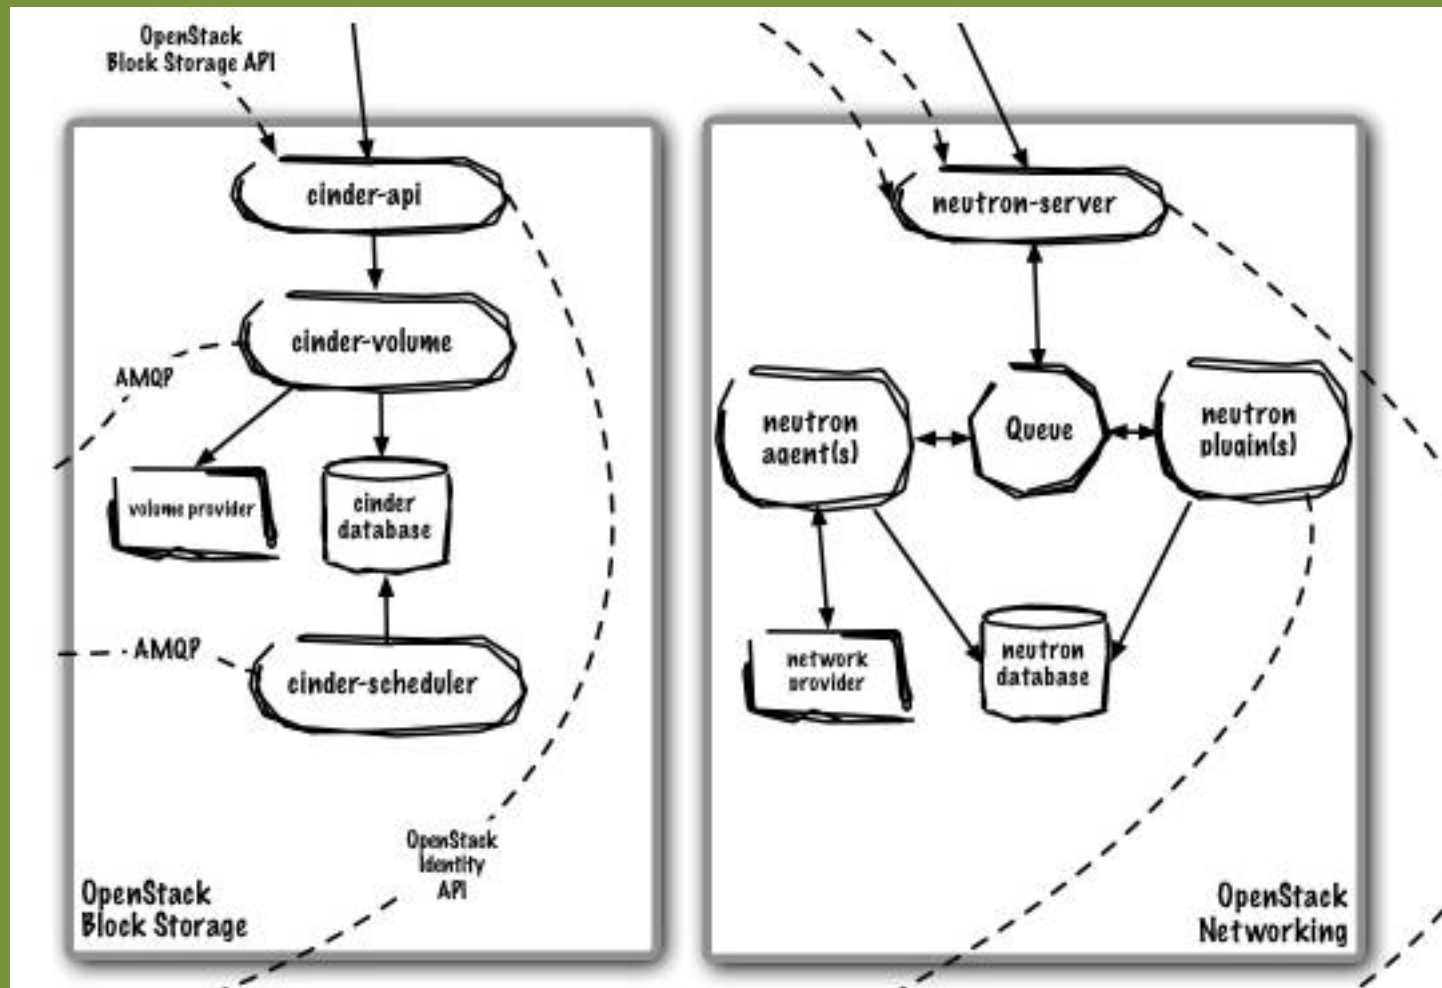


OpenStack Architecture



- Command-line interfaces (nova, neutron, swift, and so on)
- Cloud Management Tools (Rightscale, Enstratus, and so on.)
- GUI tools (Dashboard, Cyberduck, iPhone client, and so on.)







Neutron

- Neutron handle network virtualization
- It allows user/tenant to create logical networks to fit the application's need.
- Use REST API to manage network functions
- Multi-tenancy: Isolation, Abstarction
- Modular: Support Plug-in from multiple vendors

Traditional Host & Network

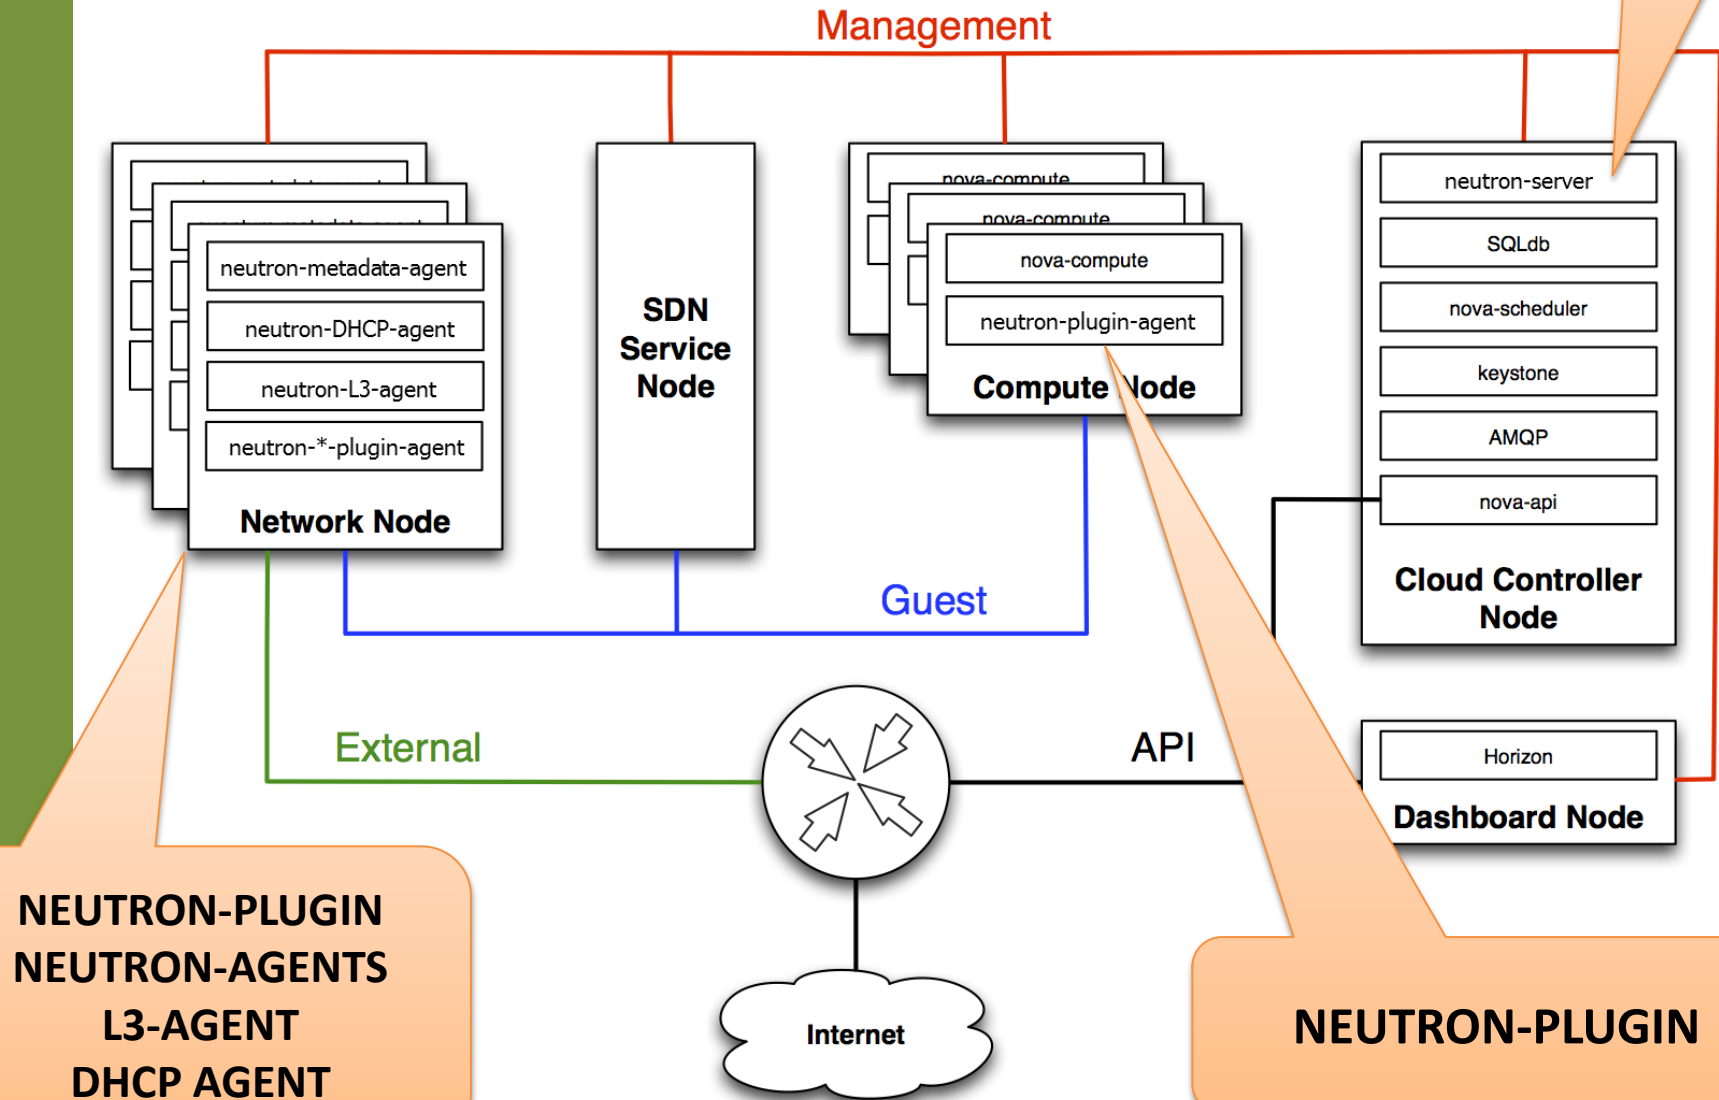


Moving Networking toward the Edge



Component Layout

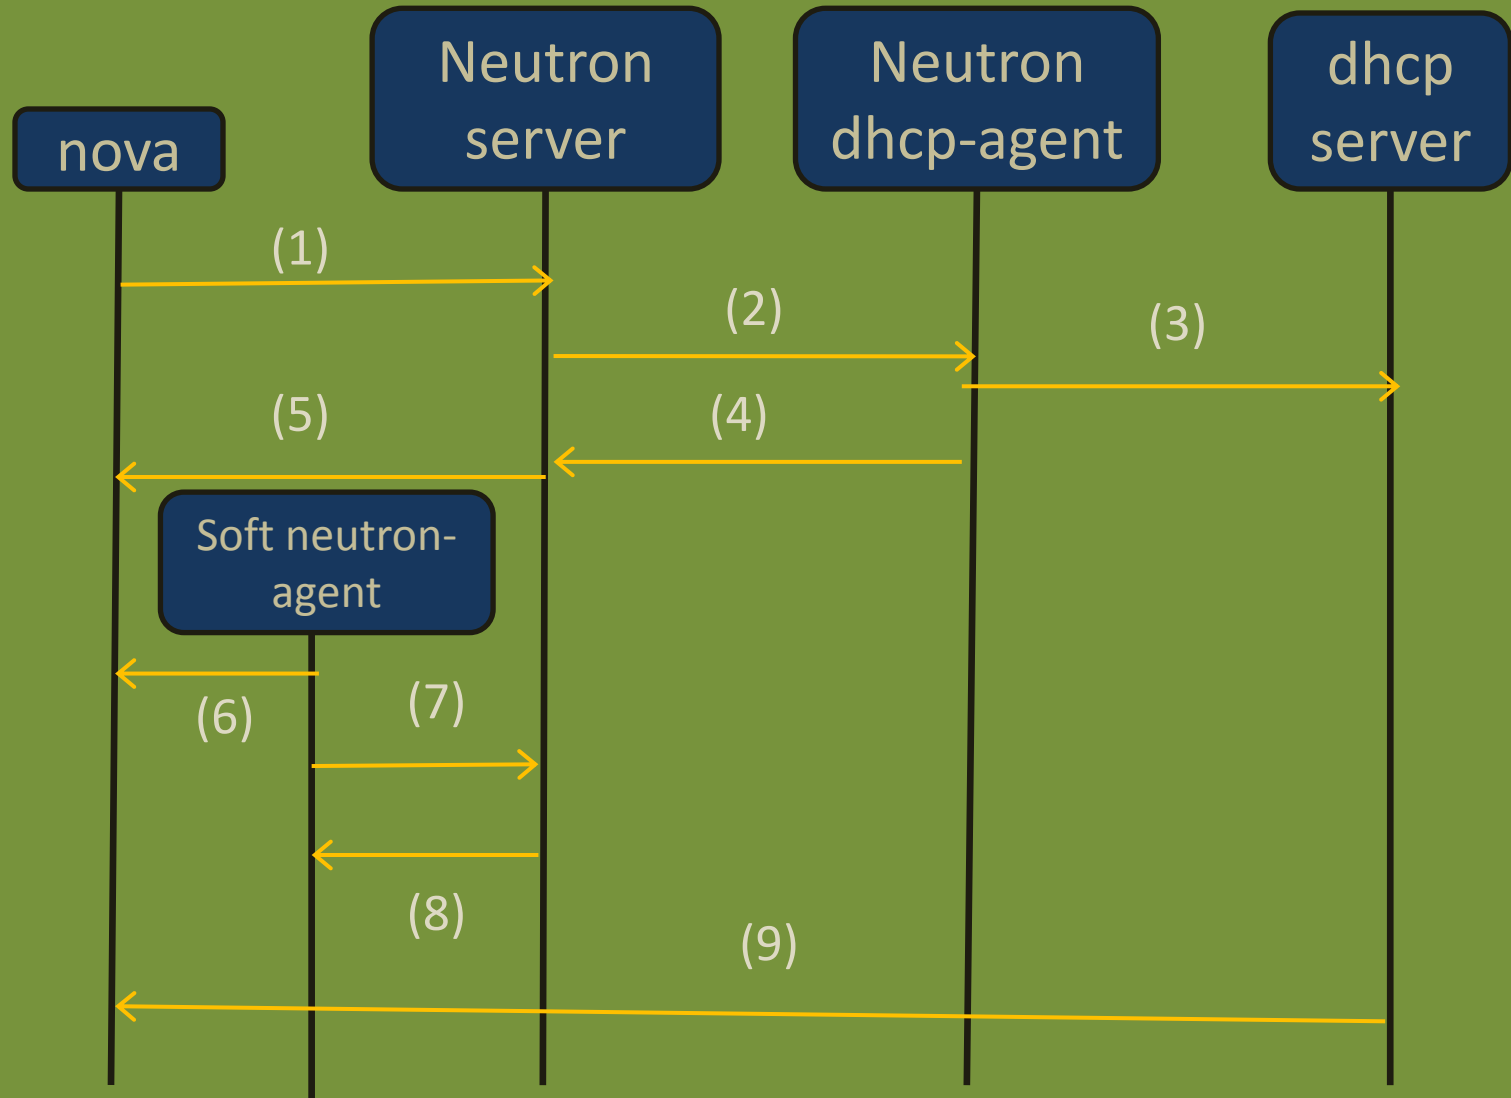
**NEUTRON
SERVER**



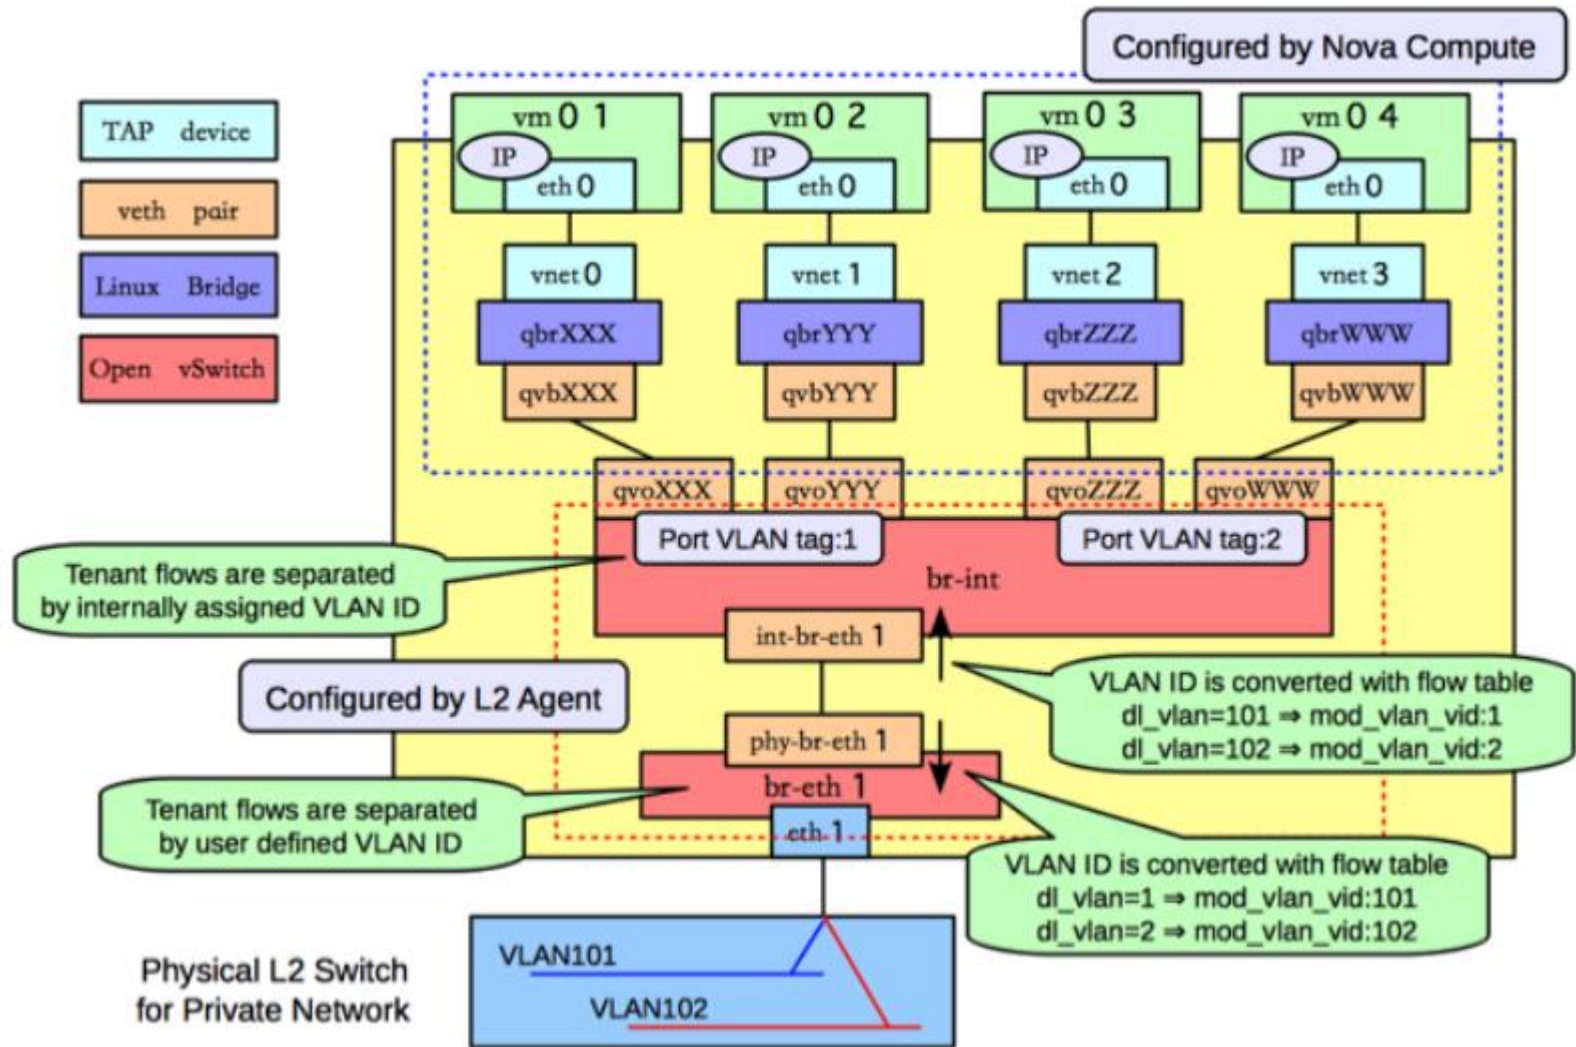
VM booting workflow between nova and neutron

1. **nova** boot will get into compute driver, which will call neutron api to create port
2. **neutron-server** creates the port object and allocates it with ip address from subnets
3. **neutron-server** notifies neutron-dhcp agent with the created port object
4. **neutron-dhcp agent** configs the dhcp server with the port object, such as IP, Mac, gateway and routes
5. **compute-driver gets** the network information, and then create port on br-int soft-switch, and then starts the VM with a tap device attached on the soft-switch port.
6. **soft-neutron-agent** detects and gets to know there is a new soft-switch port created
7. **soft-neutron-agent** asks information from neutron-server
8. **soft-neutron-agent** set up the port, such as the flows and vlan id of the soft-switch port. After this step, the VM's network is connected.
9. **VM** gets the IP address with the dhcp client.

Neutron Control Flow

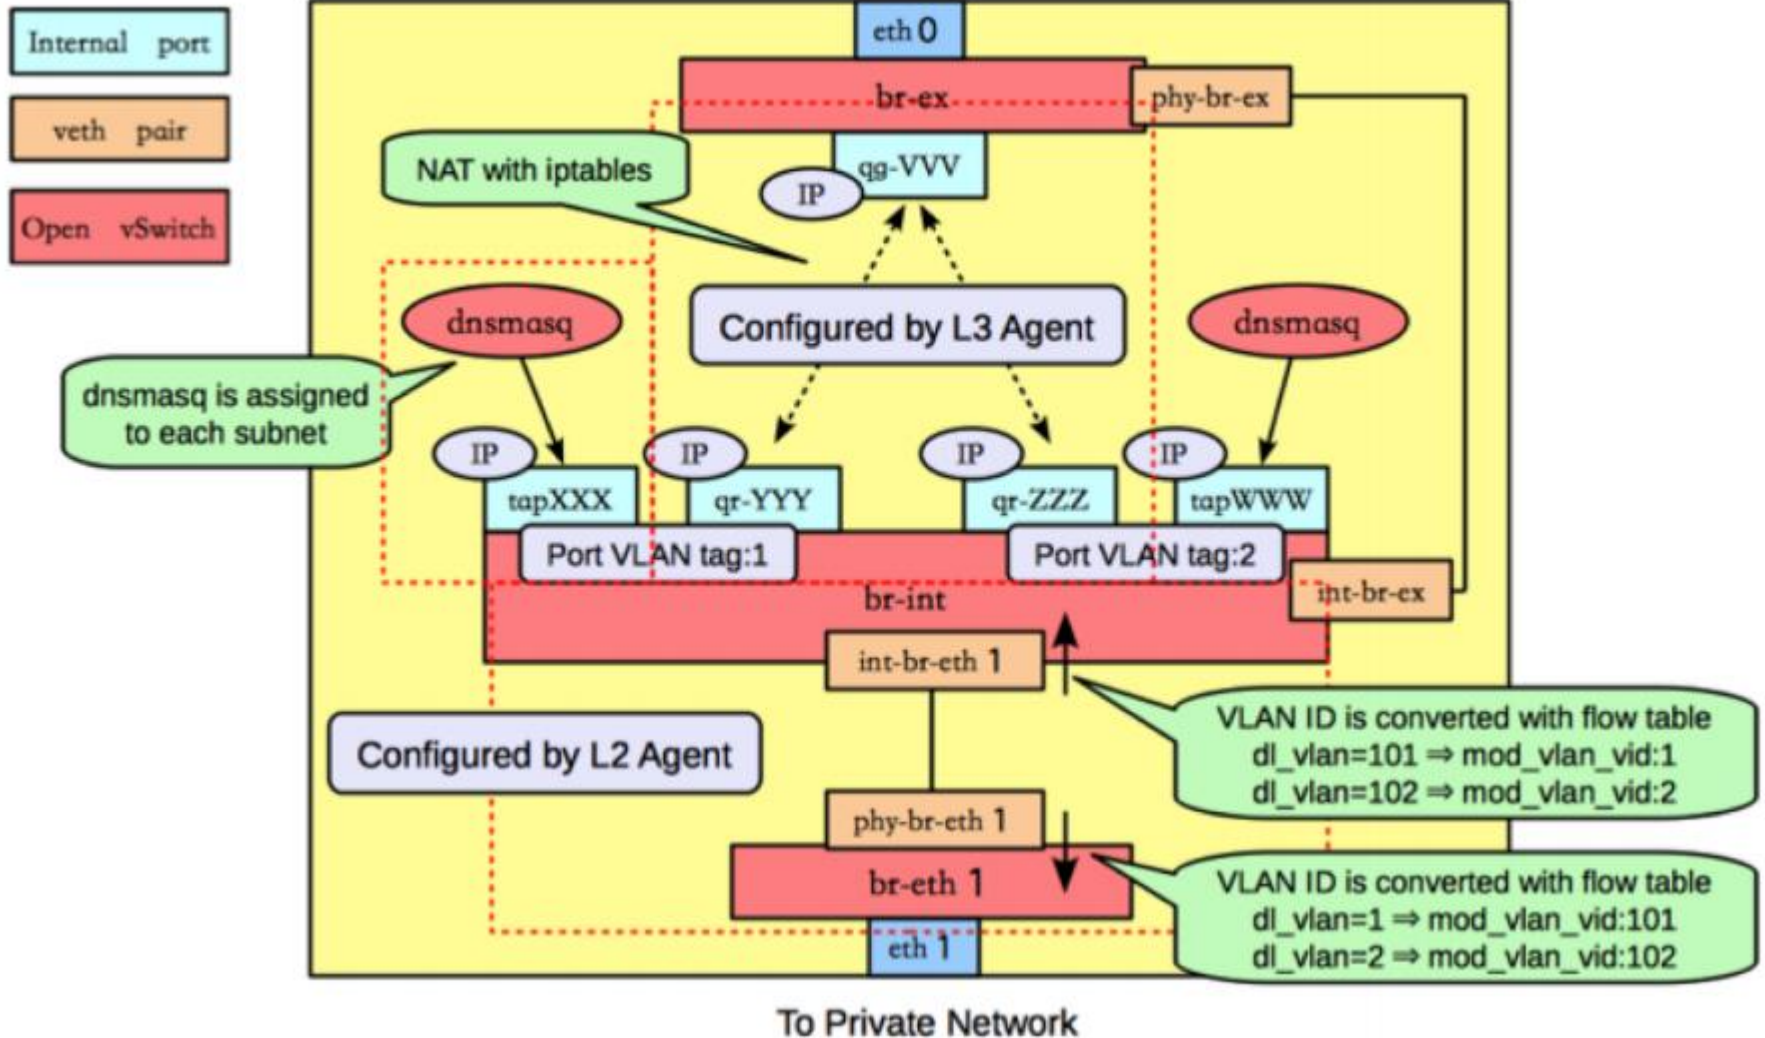


Data flow



The diagram illustrates the Open vSwitch configuration for NAT and L2/L3 agents. It shows the internal network structure, including the vSwitch, bridges, and interfaces, and how they are configured for NAT and L2/L3 agents.

- Legend:**
 - Internal port (light blue box)
 - veth pair (orange box)
 - Open vSwitch (yellow background)
- Network Structure:**
 - Top:** "To Public Network" interface (eth 0) connected to the vSwitch.
 - Bottom:** "To Private Network" interface (eth 1) connected to the vSwitch.
 - Internal Bridges:**
 - br-ex:** External bridge connected to eth 0.
 - br-int:** Internal bridge connected to eth 1.
 - br-eth 1:** Bridge connected to eth 1.
 - Interfaces and Veth Pairs:**
 - qg-VVV:** Interface connected to br-ex.
 - qr-YYY, qr-ZZZ:** Interfaces connected to br-int.
 - tapXXX, tapWWW:** Interfaces connected to br-int.
 - int-br-eth 1:** Veth pair connecting br-int and br-eth 1.
 - phy-br-eth 1:** Veth pair connecting br-eth 1 and br-eth 1.
- Configuration and Agents:**
 - NAT with iptables:** Configured on the vSwitch.
 - Configured by L3 Agent:** Configures the vSwitch and the internal bridges (br-ex, br-int, br-eth 1).
 - Configured by L2 Agent:** Configures the vSwitch and the internal bridges (br-ex, br-int, br-eth 1).
 - dnsmaq:** Assigned to each subnet.
- VLAN ID Conversion:**
 - Top:** VLAN ID is converted with flow table: $dl_vlan=101 \Rightarrow mod_vlan_vid:1$, $dl_vlan=102 \Rightarrow mod_vlan_vid:2$.
 - Bottom:** VLAN ID is converted with flow table: $dl_vlan=1 \Rightarrow mod_vlan_vid:101$, $dl_vlan=2 \Rightarrow mod_vlan_vid:102$.




Heat

What is Heat?

- Orchestration service for OpenStack (officially supports Grizzly release)
- Uses templating mechanism
- Controls complex groups of cloud resources
- Huge potential and multiple use cases
- More than 20 active contributors

Heat Basic WorkFlow



<http://www.slideshare.net/mirantis/an-introduction-to-openstack-heat>

Heat Components

Heat API

- heat-api (OpenStack native REST API) or heat-api-cfn (provides AWS Query API)
- Communicates with Heat Engine and tells it what actions to do

Heat Engine

- Does all the orchestration work
- Layer on which resource integration is implemented
- Contains abstractions to use Auto Scaling and High Availability

<http://www.slideshare.net/mirantis/an-introduction-to-openstack-heat>

HOT format

- *Stack*: In Heat parlance, a stack is the collection of objects—or *resources*—that will be created by Heat. This might include instances (VMs), networks, subnets, routers, ports, router interfaces, security groups, security group rules, auto-scaling rules, etc.
- *Template*: Heat uses the idea of a template to define a stack. If you wanted to have a stack that created two instances connected by a private network, then your template would contain the definitions for two instances, a network, a subnet, and two network ports. Since templates are central to how Heat operates, I'll show you examples of templates in this post.
- *Parameters*: A Heat template has three major sections, and one of those sections defines the template's parameters. These are tidbits of information—like a specific image ID, or a particular network ID—that are passed to the Heat template by the user. This allows users to create more generic templates that could potentially use different resources.
- *Resources*: Resources are the specific objects that Heat will create and/or modify as part of its operation, and the second of the three major sections in a Heat template.

Basic HOT

```
heat_template_version: 2013-05-23

description: Test Template

parameters:
  ImageID:
    type: string
    description: Image use to boot a server
  NetID:
    type: string
    description: Network ID for the server

resources:
  server1:
    type: OS::Nova::Server
    properties:
      name: "Test server"
      image: { get_param: ImageID }
      flavor: "ml.tiny"
      networks:
        - network: { get_param: NetID }

outputs:
  server1_private_ip:
    description: IP address of the server in the private network
    value: { get_attr: [ server1, first_address ] }
```

```
$ NET_ID=$(nova net-list | awk '/ demo-net / { print $2 }')  
$ heat stack-create -f test-stack.yml \  
  -P "ImageID=cirros-0.3.2-x86_64;NetID=$NET_ID" testStack
```

id	stack_name	stack_status	creation_time
477d96b4-d547-4069-938d-32ee990834af	testStack	CREATE_IN_PROGRESS	2014-04-01T14:10:10Z

Example HOT

```
1 heat_template_version: 2013-05-23
2 description: >
3   A simple Heat template that spins up multiple instances and a private
4 resources:
5   heat_network_01:
6     type: OS::Neutron::Net
7     properties:
8       admin_state_up: true
9       name: heat-network-01
10  heat_subnet_01:
11    type: OS::Neutron::Subnet
12    properties:
13      name: heat-subnet-01
14      cidr: 10.10.10.0/24
15      dns_nameservers: [172.16.1.11, 172.16.1.6]
16      enable_dhcp: true
17      gateway_ip: 10.10.10.254
18      network_id: { get_resource: heat_network_01 }
```

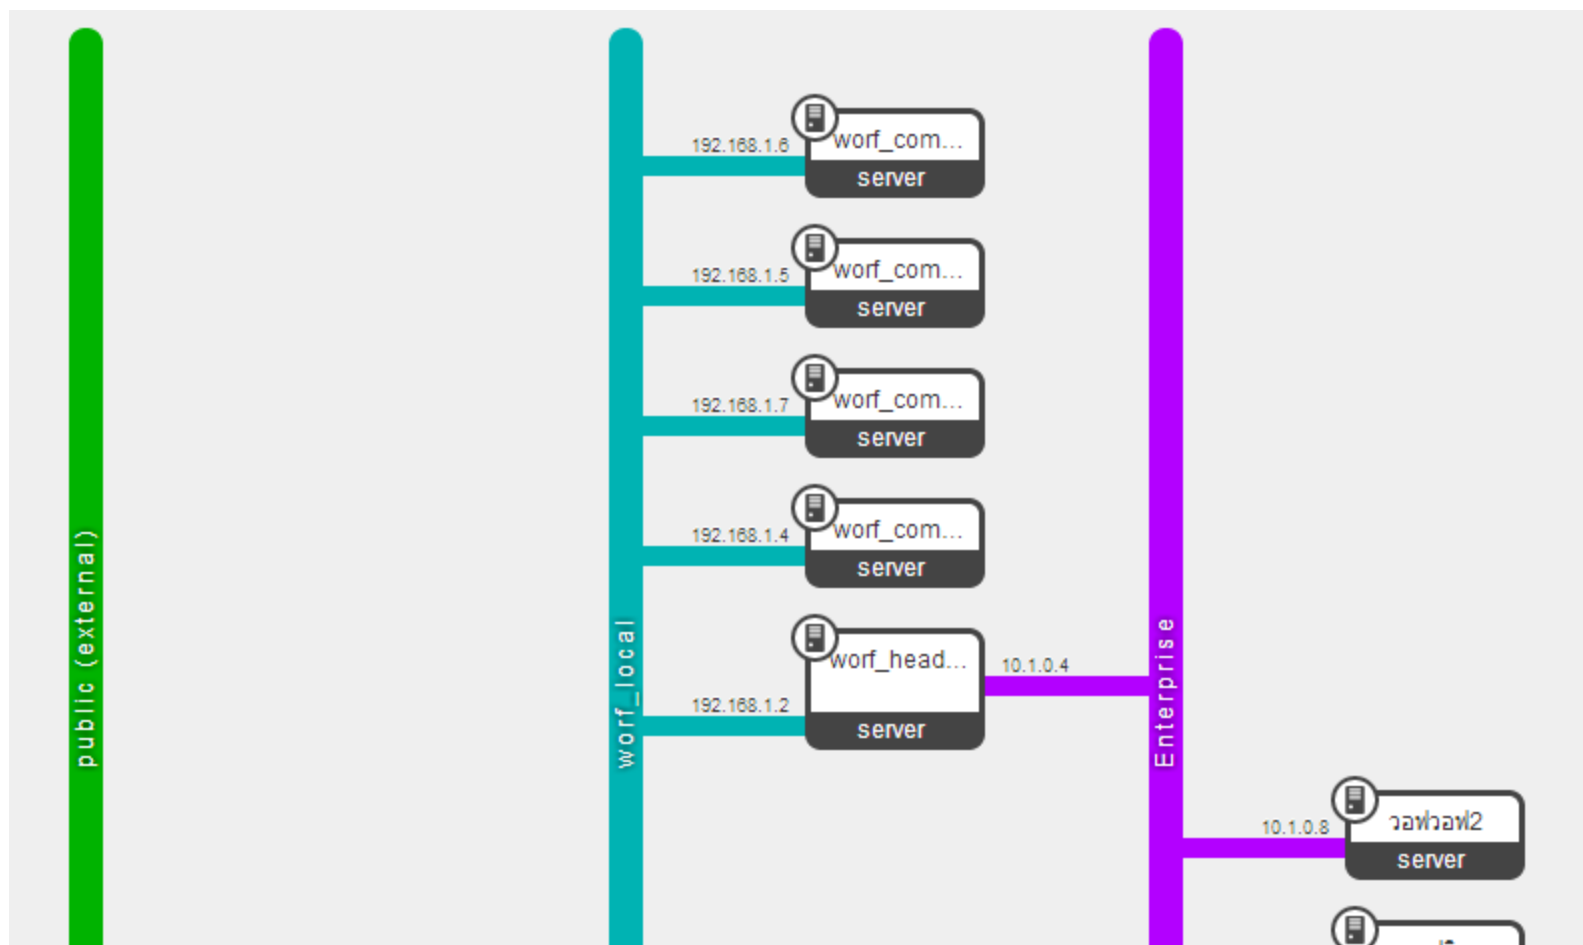


```
19 heat_router_01:
20   type: OS::Neutron::Router
21   properties:
22     admin_state_up: true
23     name: heat-router-01
24 heat_router_01_gw:
25   type: OS::Neutron::RouterGateway
26   properties:
27     network_id: 604146b3-2e0c-4399-826e-a18cbc18362b
28     router_id: { get_resource: heat_router_01 }
29 heat_router_int0:
30   type: OS::Neutron::RouterInterface
31   properties:
32     router_id: { get_resource: heat_router_01 }
33     subnet_id: { get_resource: heat_subnet_01 }
```

```
34 instance0_port0:
35     type: OS::Neutron::Port
36     properties:
37         admin_state_up: true
38         network_id: { get_resource: heat_network_01 }
39         security_groups:
40             - b0ab35c3-63f0-48d2-8a6b-08364a026b9c
41 instance1_port0:
42     type: OS::Neutron::Port
43     properties:
44         admin_state_up: true
45         network_id: { get_resource: heat_network_01 }
46         security_groups:
47             - b0ab35c3-63f0-48d2-8a6b-08364a026b9c
```

```
48 instance0:
49     type: OS::Nova::Server
50     properties:
51         name: heat-instance-01
52         image: 01b0eb5d-14ae-4c9e-8025-a21e6f733034
53         flavor: m1.xsmall
54         networks:
55             - port: { get_resource: instance0_port0 }
56 instance1:
57     type: OS::Nova::Server
58     properties:
59         name: heat-instance-02
60         image: 01b0eb5d-14ae-4c9e-8025-a21e6f733034
61         flavor: m1.xsmall
62         networks:
63             - port: { get_resource: instance1_port0 }
```

Creating a Cluster Computer on **OpenStack**



Create a local network

Science Cloud
Thammasat University

Project

CURRENT PROJECT
Qo'noS

Manage Compute

Networks

Logged in as: worf [Community](#) [Settings](#)

Networks

+ Create Network

Create Netv

<input type="checkbox"/>	Name	Subnets Associated	Shared	Status	Admin State	Actions
<input type="checkbox"/>	Enterprise	Enterprise 10.1.0.0/24	Yes	ACTIVE	UP	

Displaying 1 item

Science Cloud
Thammasat University

Project

CURRENT PROJECT
Qo'noS

Manage Compute

Overview

Networks

Logged in as: worf [Community](#) [Settings](#) [Help](#)

Networks

+ Create Network

Delete N

<input type="checkbox"/>	Name	Subnets Associated	Shared	Status	Admin State	Actions
<input type="checkbox"/>	worf_local	192.168.1.0/24	No	ACTIVE	UP	<div>Edit Network More</div>
<input type="checkbox"/>	Enterprise	Enterprise 10.1.0.0/24	Yes	ACTIVE	UP	

Displaying 2 items

Launch a head node

Launch Instance

DetailsAccess & SecurityNetworkingVolume OptionsPost-Creation

Instance Source

Image

Image

Ubuntu-12.04-Desktop

Instance Name

worf_head_node

Flavor

medium

Instance Count

1

Specify the details for launching an instance.

The chart below shows the resources used by this project in relation to the project's quotas.

Flavor Details

Name	medium
VCPUs	2
Root Disk	40 GB
Ephemeral Disk	40 GB
Total Disk	80 GB
RAM	4,096 MB

Project Quotas

Number of Instances (3) 7 Available

Launch Instance

Details

Access & Security

Networking

Volume Options

Post-Creation

Selected Networks

- nic:1 ↕ Enterprise (f6c0f209-9b04-4acc-843b-262a3a8e38f5) -
- nic:2 ↕ worf_local (8b7b1b52-6718-4592-a8eb-e557669728c3) -

Available networks

Choose network from Available networks to Selected Networks by push button or drag and drop, you may change nic order by drag and drop as well.

Cancel

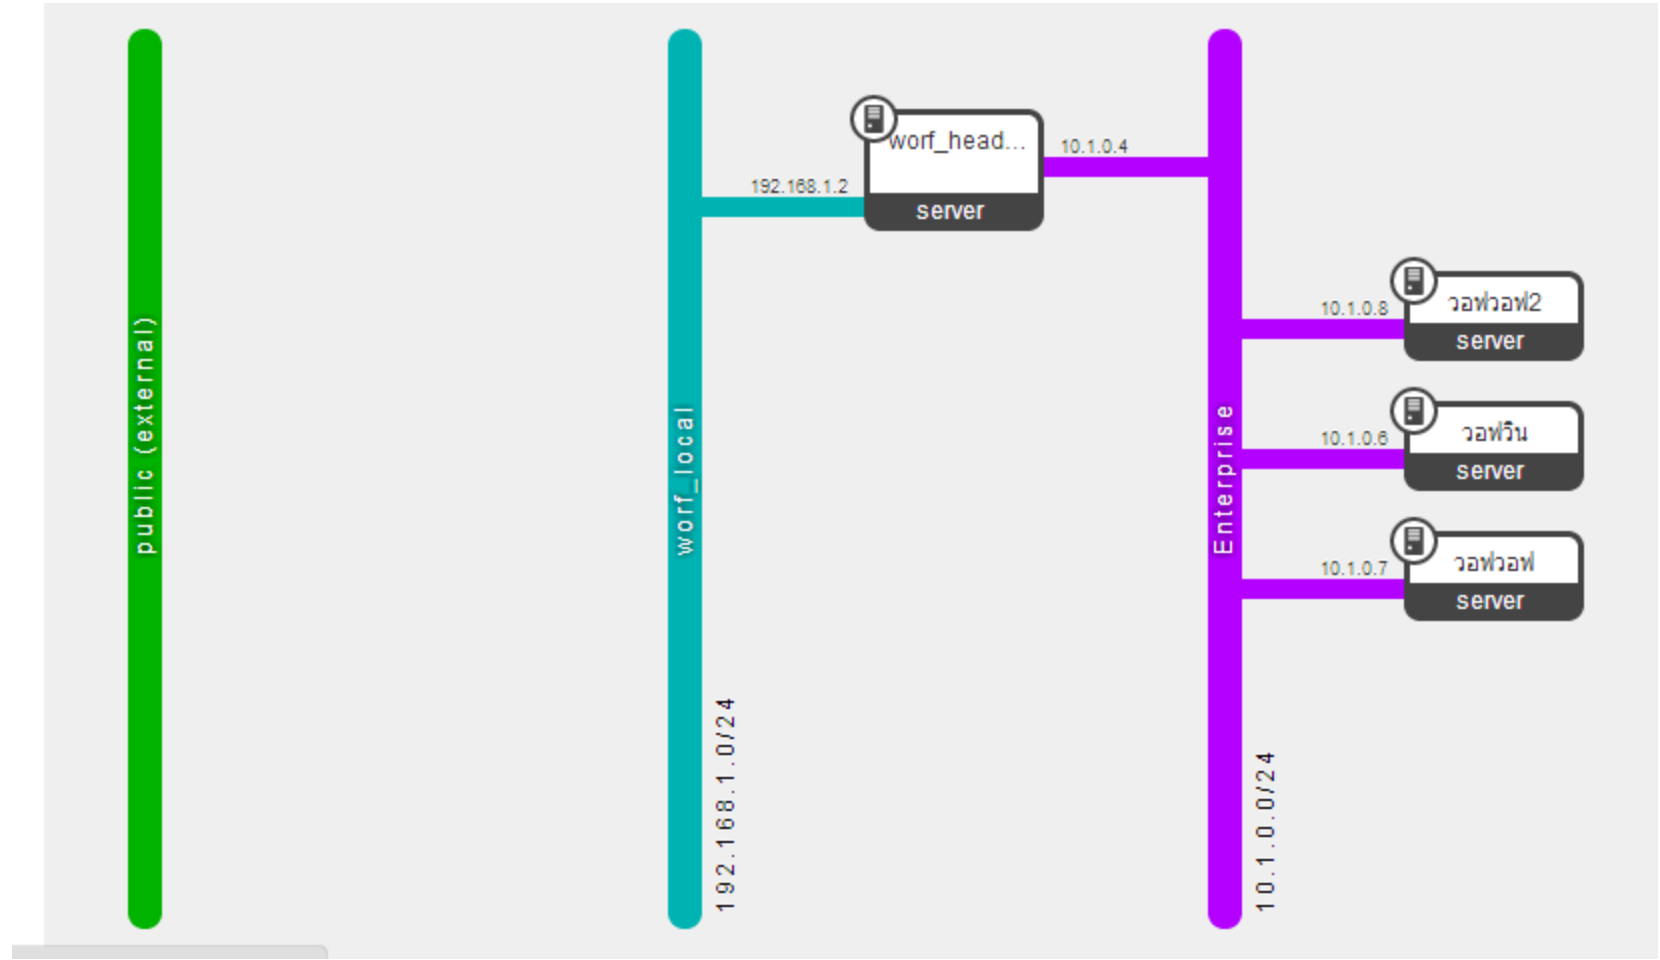
Launch

Instances

[+ Launch Instance](#)[🗑 Terminate Instances](#)

<input type="checkbox"/>	Instance Name	IP Address	Size	Keypair	Status	Task	Power State	Actions
<input type="checkbox"/>	worf_head_node	worf_local 192.168.1.2 Enterprise 10.1.0.4	medium 4GB RAM 2 VCPU 40GB Disk	-	Active	None	Running	Create Snapshot More ▾
<input type="checkbox"/>	วอฟหัวฟ2	10.1.0.8 10.100.20.19	medium 4GB RAM 2 VCPU 40GB Disk	-	Active	None	Running	Create Snapshot More ▾
<input type="checkbox"/>	วอฟวิน	10.1.0.6 10.100.20.21	medium 4GB RAM 2 VCPU 40GB Disk	-	Active	None	Running	Create Snapshot More ▾
<input type="checkbox"/>	วอฟฟ	10.1.0.7 10.100.20.20	m1.medium 4GB RAM 2 VCPU 40GB Disk	-	Suspended	None	Shutdown	Associate Floating IP More ▾

Displaying 4 items



Launch compute nodes

Insta

Launch Instance

×

Details

Access & Security

Networking

Volume Options

Post-Creation

Instance Source

Image

▼

Image

Ubuntu-12.04-Server

▼

Instance Name

worf_compute_node_

Flavor

m1.small

▼

Instance Count

4

Specify the details for launching an instance.

The chart below shows the resources used by this project in relation to the project's quotas.

Flavor Details

Name	m1.small
VCPUs	1
Root Disk	20 GB
Ephemeral Disk	0 GB
Total Disk	20 GB
RAM	2,048 MB

Number of instances to launch.

Number of Instances (4)

6 Available

Launch Instance



Details

Access & Security

Networking

Volume Options

Post-Creation

Selected Networks

nic:1 ↕ worf_local (3b7b1b52-6718-4592-a5eb-e557669728c3) -

Available networks


↕ Enterprise (f5c0f209-9b04-4aca-843b-262a3a8e38f5) +

Choose network from Available networks to Selected Networks by push button or drag and drop, you may change nic order by drag and drop as well.

Cancel

Launch

<input type="checkbox"/>	Instance Name	IP Address	Size	Keypair	Status	Task	Power State	Actions
<input type="checkbox"/>	worf_compute_node_-49e4dc1d-8939-4dc2-9492-fe98abd2ac3d	192.168.1.6	m1.small 2GB RAM 1 VCPU 20GB Disk	-	Active	None	Running	<button>Create Snapshot</button> <button>More ▾</button>
<input type="checkbox"/>	worf_compute_node_-e1a1023c-d1de-4c1f-91fb-5f0605047598	192.168.1.5	m1.small 2GB RAM 1 VCPU 20GB Disk	-	Active	None	Running	<button>Create Snapshot</button> <button>More ▾</button>
<input type="checkbox"/>	worf_compute_node_-673ce66f-09a0-4eb1-aa6a-2be9756d52f0	192.168.1.7	m1.small 2GB RAM 1 VCPU 20GB Disk	-	Active	None	Running	<button>Create Snapshot</button> <button>More ▾</button>
<input type="checkbox"/>	worf_compute_node_-ab4cb991-aeee-4b69-b0fd-9b82c903f592	192.168.1.4	m1.small 2GB RAM 1 VCPU 20GB Disk	-	Active	None	Running	<button>Create Snapshot</button> <button>More ▾</button>
<input type="checkbox"/>	worf_head_node	worf_local 192.168.1.2 Enterprise 10.1.0.4	medium 4GB RAM 2 VCPU 40GB Disk	-	Active	None	Running	<button>Create Snapshot</button> <button>More ▾</button>



ubuntu@ubuntu-virtual-machine: ~



ubuntu@ubuntu-virtual-machine:~\$ ping www.google.com
PING www.google.com (61.19.1.118) 56(84) bytes of data.
64 bytes from 61.19.1.118: icmp_req=1 ttl=51 time=15.5 ms
64 bytes from 61.19.1.118: icmp_req=2 ttl=51 time=4.13 ms
64 bytes from 61.19.1.118: icmp_req=3 ttl=51 time=4.19 ms
^C64 bytes from 61.19.1.118: icmp_req=4 ttl=51 time=4.23 ms

--- www.google.com ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 15302ms
rtt min/avg/max/mdev = 4.138/7.017/15.501/4.898 ms

ubuntu@ubuntu-virtual-machine:~\$ ping 192.168.1.4
PING 192.168.1.4 (192.168.1.4) 56(84) bytes of data.
64 bytes from 192.168.1.4: icmp_req=1 ttl=64 time=30.4 ms
64 bytes from 192.168.1.4: icmp_req=2 ttl=64 time=0.986 ms
^C

--- 192.168.1.4 ping statistics ---
2 packets transmitted, 2 received, 0% packet loss, time 1001ms
rtt min/avg/max/mdev = 0.986/15.711/30.437/14.726 ms

ubuntu@ubuntu-virtual-machine:~\$ ping 192.168.1.5
PING 192.168.1.5 (192.168.1.5) 56(84) bytes of data.
64 bytes from 192.168.1.5: icmp_req=1 ttl=64 time=33.7 ms
64 bytes from 192.168.1.5: icmp_req=2 ttl=64 time=0.862 ms
^C

--- 192.168.1.5 ping statistics ---
2 packets transmitted, 2 received, 0% packet loss, time 1001ms